

CS 330: Network Applications & Protocols

Application Layer: HTTP

Galin Zhelezov
Department of Physical Sciences
York College of Pennsylvania



Overview of Application Layer


- **Network Application Architectures**
- **HyperText Transfer Protocol (HTTP)**
- **File Transfer and Email protocols (FTP, SMTP)**
- **Domain Name System (DNS)**
- **Peer-to-Peer Applications (P2P)**
- **Content Delivery Networks**
- **Socket Programming**

Overview of Application Layer

- **Network Application Architectures**
 - Protocol Layers
 - Client-Server vs. Peer-to-Peer
 - Process Communication
 - Transport Services
- **HyperText Transfer Protocol (HTTP)**
- **File Transfer and Email protocols (FTP, SMTP)**
- **Domain Name System (DNS)**
- **Peer-to-Peer Applications (P2P)**

Protocol Layers

- **Top-Down Approach**




Example Applications

- **E-mail**
- **Web**
- **Text messaging**
- **Remote login**
- **P2P file sharing**
- **Multi-user network games**
- **Streaming stored video**
(YouTube, Hulu, Netflix)
- **Voice over IP (e.g. Skype)**
- **Real-time video conferencing**
- **Social networking**
- **Search**
- **...**

Creating a network app

write programs that:

- run on (different) *end systems*
 - communicate over network
 - e.g., web server software communicates with browser software
- no need to write software for network-core devices**
- network-core devices do not run user applications
 - applications on end systems allows for rapid app development, propagation



Network Application Architectures

- **Client-Server**
- **Peer-to-Peer**


Client-Server Architecture

- **Clients:**

- Request service from server
- May be intermittently connected
- May have dynamic IP addresses
- Do not communicate directly with each other

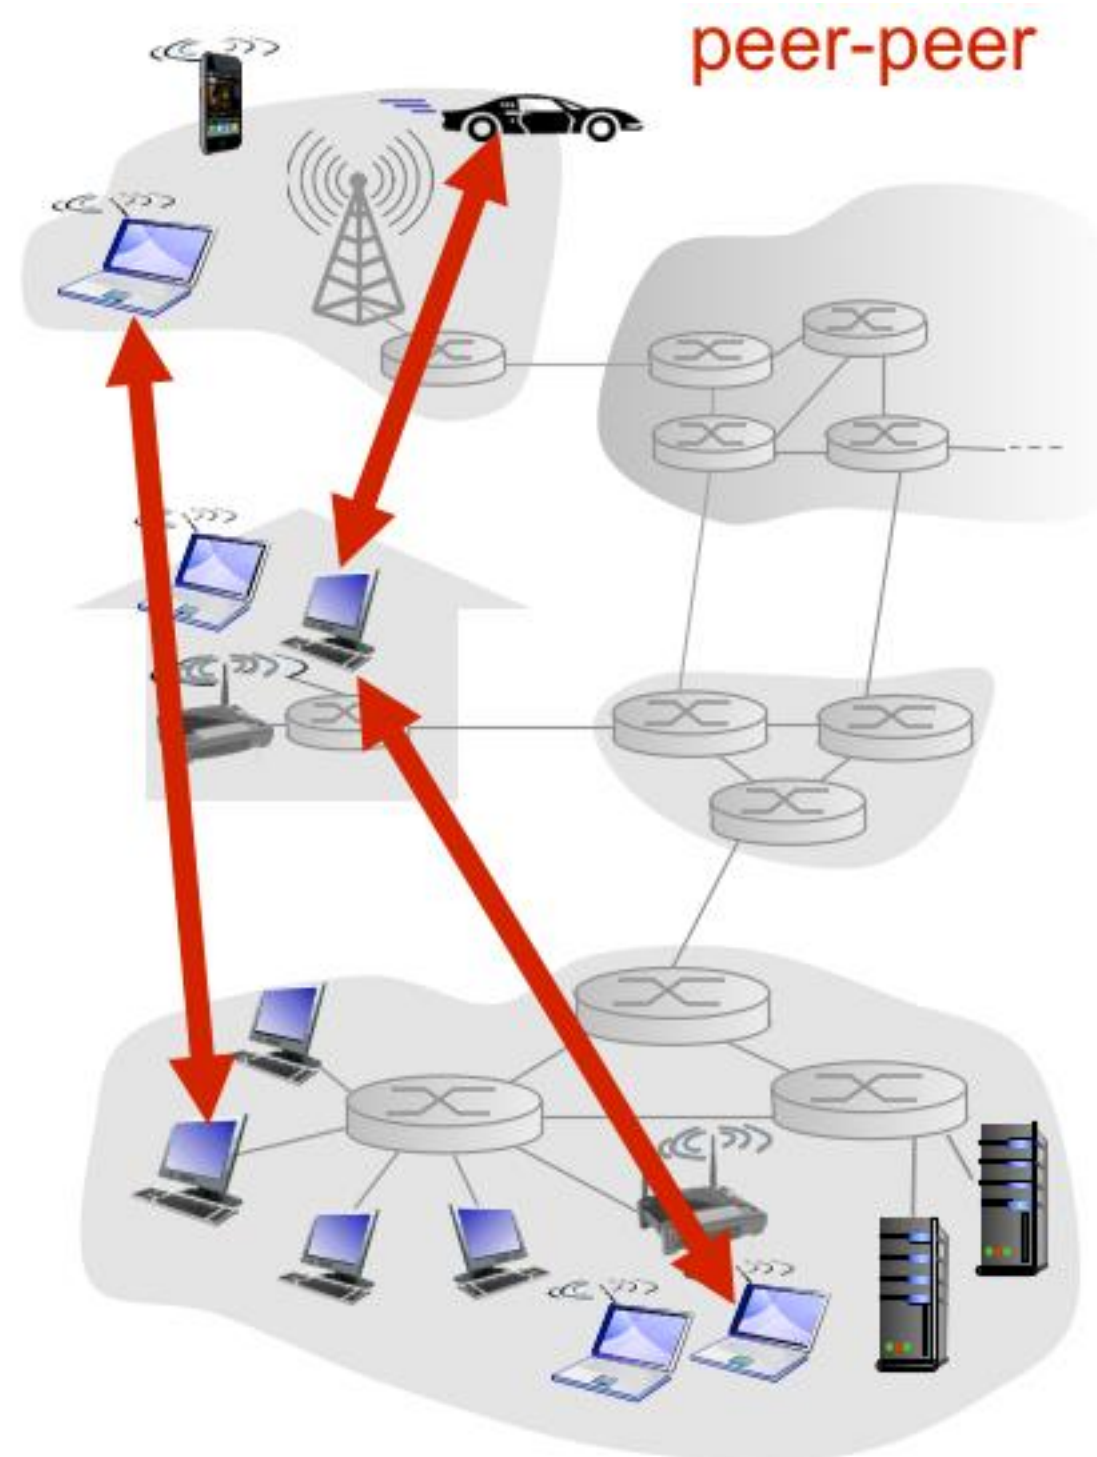
- **Server:**

- Provides a service to clients
- Always-on host
- Permanent IP address
- Data centers for scaling



Peer-to-Peer Architecture

- Does not require always-on servers
- Hosts communicate directly with each other
- Peers request service from other peers, and provide service in return to other peers
- Highly scalable
- Self scalability – new peers bring new service capacity, as well as new service demands
- Hosts are intermittently connected and may change IP addresses
- Difficult to authenticate - possibly insecure
- Hosts need incentive to share data



Process Communications

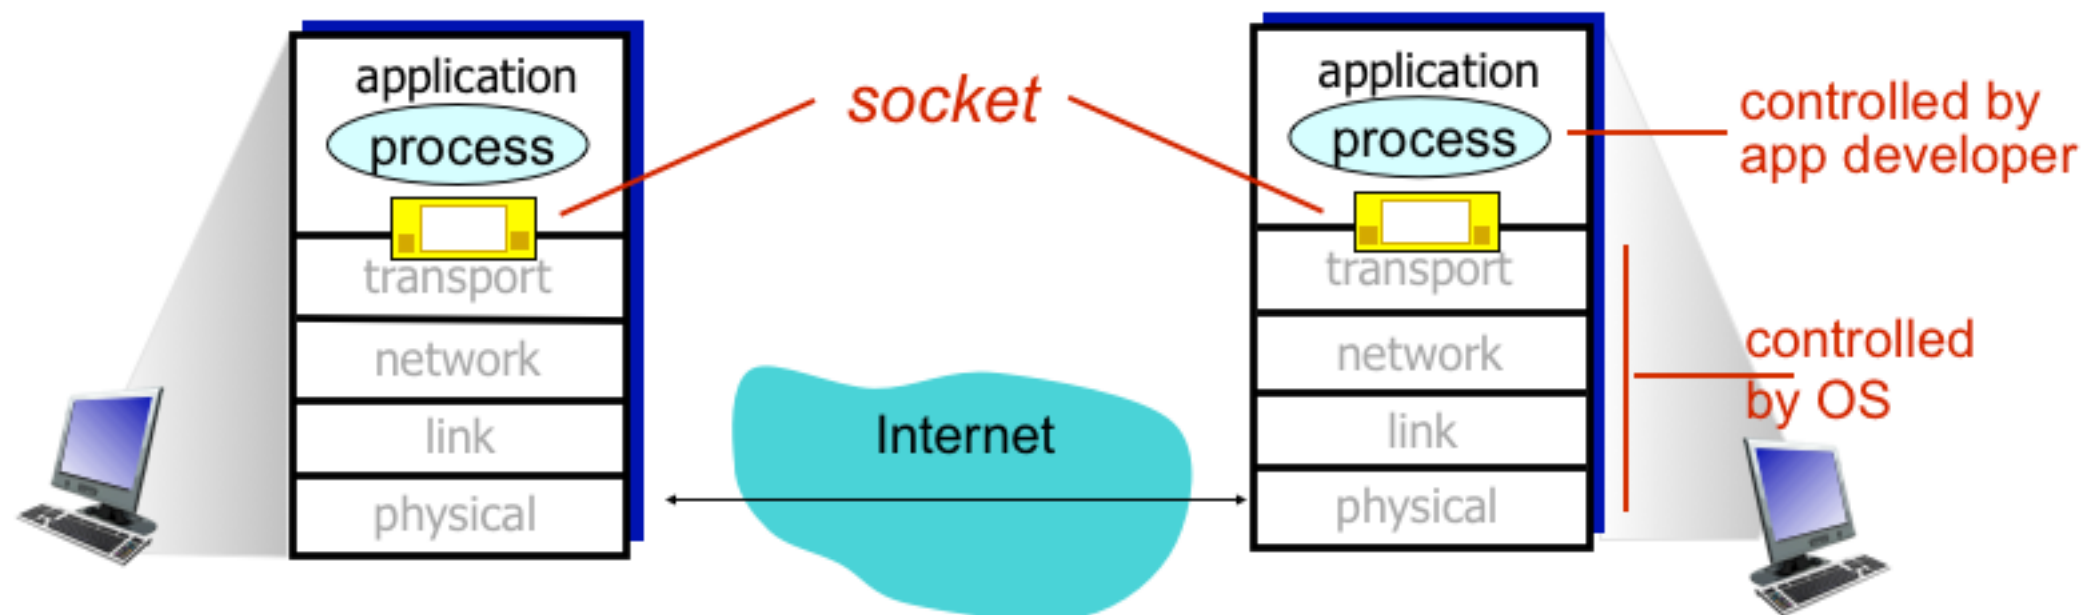
- **A process is program running on a host**
 - **Client process** - process that initiates communication
 - **Server process** - process that waits to be contacted

Note: Applications with P2P architectures have both client processes & server processes

- **Sockets provide a mechanism for inter-process communication (IPC)**
 - Inter-process communication on the same host
 - **Operating system provides message passing**
 - Inter-process communication on different hosts
 - **Network provides message passing**

Sockets

- **Process sends/receives messages to/from its socket**
- **Socket analogous to door**
 - Sending process shoves message out door
 - Sending process relies on transport infrastructure on other side of door to deliver message to socket at receiving process



Addressing Processes

- To receive messages, process must have some **identifier**
- Host device has unique 32-bit IP address (IPv4)
 - IP address alone is insufficient to address a process on a system
 - Many processes may be running on the same system
- Identifier includes both **IP address and port numbers associated with process on host**
 - Example port numbers:
 - HTTP server: port 80
 - SMTP Mail server: port 25
- To send an **HTTP** message to **cs.ycp.edu** web server:
 - IP address: 192.245.87.64
 - Port number: 80

Application Layer Protocol Defines

- **What types of messages are exchanged**
 - e.g. request, response
- **Message syntax:**
 - What fields are in messages
 - How fields are delineated
- **Message semantics:**
 - Meaning of information in fields
 - Rules for when and how processes send & respond to messages

What Transport Service Does an Application Need?

- **Data integrity**

- Some applications require 100% reliable data transfer (e.g. file transfer, web transactions)
- Other applications can tolerate some loss (e.g. audio)

- **Timing**

- Some applications require low delay to be “effective” (e.g. Internet telephony, interactive games)

- **Throughput**

- Some applications require minimum amount of throughput to be “effective” (e.g., multimedia)
- Other “elastic” applications make use of whatever throughput they get

- **Security**

- Encryption, data integrity, etc.

Transport service requirements: common apps

application	data loss	throughput	time sensitive
file transfer	no loss	elastic	no
e-mail	no loss	elastic	no
Web documents	no loss	elastic	no
real-time audio/video	loss-tolerant	audio: 5kbps-1Mbps video: 10kbps-5Mbps	yes, 100' s msec
stored audio/video	loss-tolerant	same as above	yes, few secs
interactive games	loss-tolerant	few kbps up	yes, 100' s msec
text messaging	no loss	elastic	yes and no

Internet Transport Protocol Services

TCP (Transmission Control Protocol)	UDP (User-Datagram Protocol)
Reliable data transfer	Unreliable data transfer
Packet sequence # required	Sequence # is optional
Every packet is acked	Not acked
Lost packets are retransmitted	No retransmission
May cause long delay	Quick and lossy
Connection-oriented service	Connection-less service
Good for reliable and delay-insensitive applications	Good for loss-tolerant and delay sensitive applications
Application examples: email, HTTP, FTP, remote terminal access	Application examples: Telephony, streaming multimedia

Overview of Application Layer

- **Network Application Architectures**
- **HyperText Transfer Protocol (HTTP)**
 - HTTP Overview
 - HTTP Communication
 - HTTP Request Message
 - HTTP Response Message
 - Cookies
 - Web Caching
- **File Transfer and Email protocols (FTP, SMTP)**
- **Domain Name System (DNS)**
- **Peer-to-Peer Applications (P2P)**

Internet apps: application, transport protocols

application	application layer protocol	underlying transport protocol
e-mail	SMTP [RFC 2821]	TCP
remote terminal access	Telnet [RFC 854]	TCP
Web	HTTP [RFC 2616]	TCP
file transfer	FTP [RFC 959]	TCP
streaming multimedia	HTTP (e.g., YouTube), RTP [RFC 1889]	TCP or UDP
Internet telephony	SIP, RTP, proprietary (e.g., Skype)	TCP or UDP

Securing TCP

TCP & UDP

- no encryption
- cleartext passwds sent into socket traverse Internet in cleartext

SSL

- provides encrypted TCP connection
- data integrity
- end-point authentication

SSL is at app layer

- apps use SSL libraries, that “talk” to TCP

SSL socket API

- cleartext passwords sent into socket traverse Internet encrypted


Overview of Application Layer

- Network Application Architectures
- **HyperText Transfer Protocol (HTTP)**
- **File Transfer and Email protocols (FTP, SMTP)**
- **Domain Name System (DNS)**
- **Peer-to-Peer Applications (P2P)**
- **Content Delivery Networks**
- **Socket Programming**

HTTP Overview


- **HTTP client - web browser (e.g. Chrome, Firefox, Safari)**
- **HTTP server - web server (Apache, Microsoft Internet Information Service (IIS))**
- **Web page consists of:**
 - A group of objects
 - HTML files, images, Java applets, audio files, etc.
 - A base HTML file which references objects
 - Each object is addressable by a URL

http://faculty.ycp.edu/~jmoscola/classes/cs330/docs/cs330_syllabus.pdf



HTTP Overview (Cont.)

- **HTTP: hypertext transfer protocol**
- **Web's application layer protocol**
- **Uses client/server model**
 - **Client**: browser that requests, receives, (using HTTP protocol) and “displays” web objects
 - **Server**: web server sends (using HTTP protocol) objects in response to requests



HTTP Communication

- **Uses TCP**

- Client initiates TCP connection (creates socket) to server, port 80
- Server accepts TCP connection from client
- HTTP messages exchanged between web browser (HTTP client) and web server (HTTP server)
- TCP connection closed

- **HTTP is “stateless”**

- Server maintains no information about past client requests

HTTP Connections

- **Non-persistent HTTP**

- At most one object is sent over a TCP connection
 - Open connection, get one object, close connection
- Downloading multiple objects requires multiple connections

- **Persistent HTTP**

- Multiple objects can be sent over single TCP connection between client and server
 - Server leaves the connection open after sending an object and closes on timeout

HTTP Request Message

- Two types of HTTP messages: request, response
- HTTP Request Message
 - ASCII (human-readable format)

Request line
(GET, POST,
HEAD commands)

GET /~jmoscola/index.html HTTP/1.1\r\n

Header Lines

Host: faculty.ycp.edu\r\n

User-Agent: Firefox/3.6.10\r\n

Accept: text/html,application/xhtml+xml\r\n

Accept-Language: en-us,en;q=0.5\r\n

Accept-Encoding: gzip,deflate\r\n

Accept-Charset: ISO-8859-1,utf-8;q=0.7\r\n

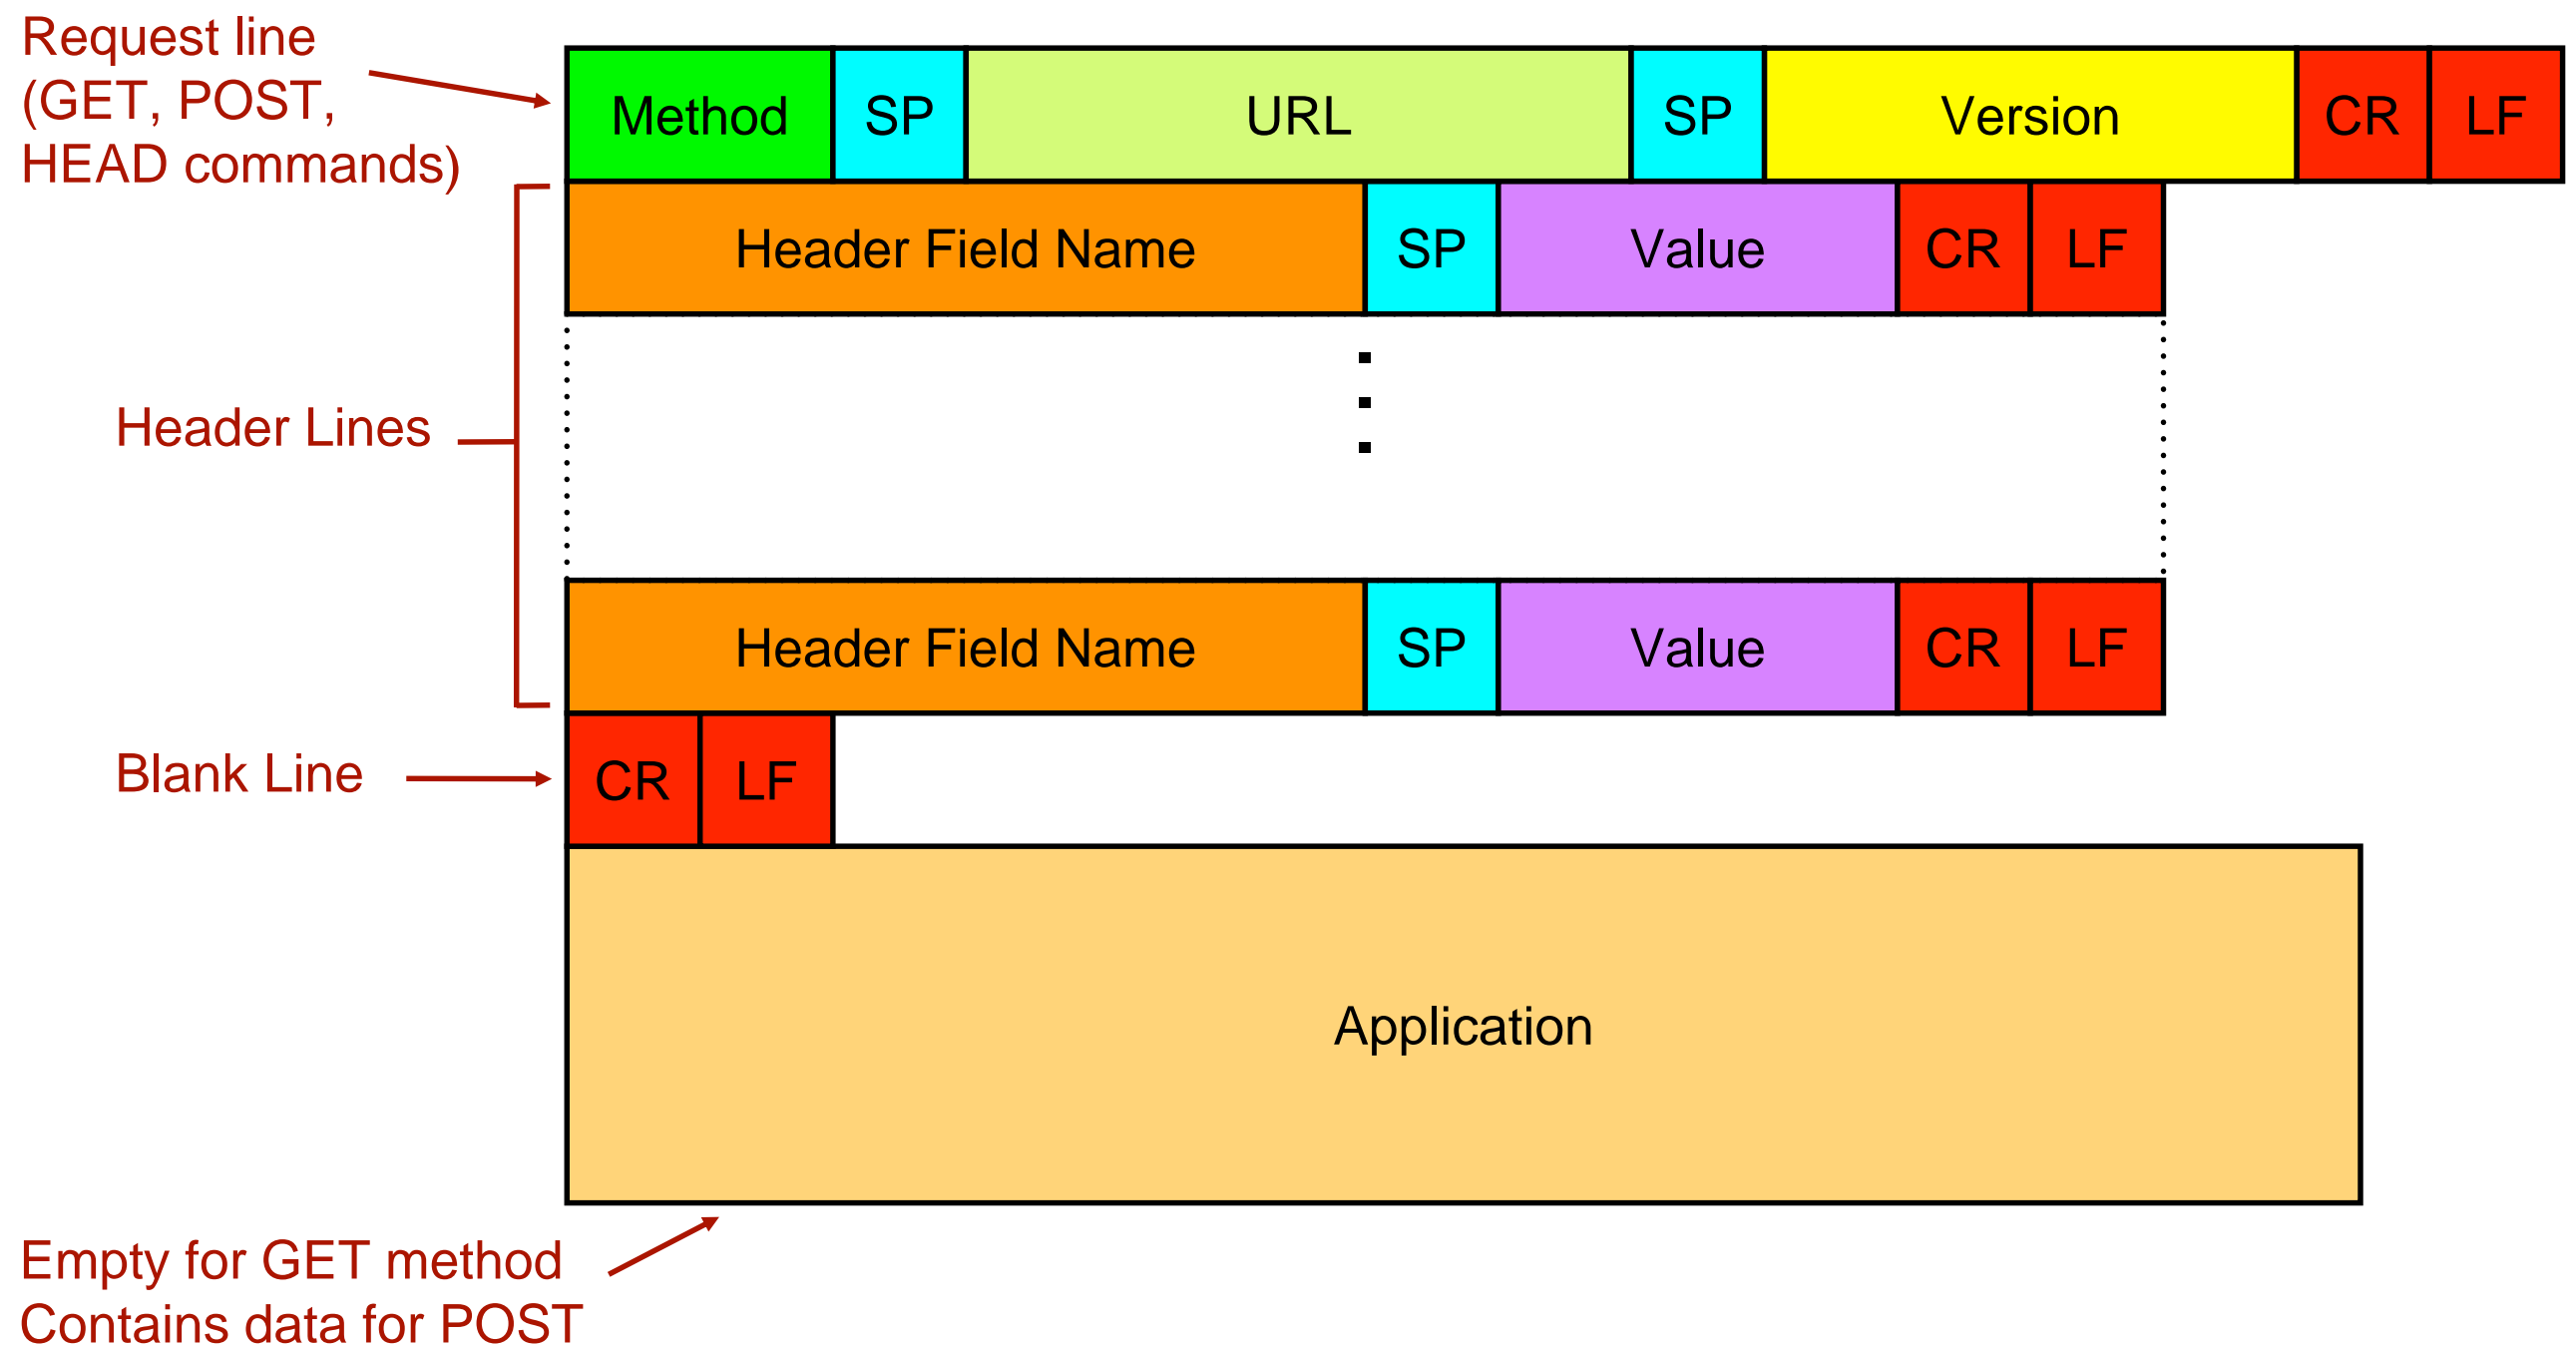
Keep-Alive: 115\r\n

Connection: keep-alive\r\n

\r\n

Carriage return,
line feed at start
of line indicates
end of header lines

HTTP Request Message: General Format



HTTP Methods

- **GET** - used to request an object from a server
 - Requested object is in URL field of HTTP request message
- **HEAD** - same as GET, but only sends header
 - Doesn't actually send requested object
 - Useful for testing/debugging
- **POST** - used to send information to a server when requesting an object
 - The object returned may depend on the information posted
 - Often used when filling out web forms
- **PUT** - uploads file in entity body to path specified in URL field
- **DELETE** - deletes file specified in the URL field

Method types

HTTP/1.0:

- **GET**
- **POST**
- **HEAD**
 - asks server to leave requested object out of response

HTTP/1.1:

- **GET, POST, HEAD**
- **PUT**
 - uploads file in entity body to path specified in URL field
- **DELETE**
 - deletes file specified in the URL field

Uploading Form Input

- **Two methods:**

- POST method:

- Web page often includes form input
 - Input is uploaded to server in the Entity portion of the HTTP request message

- URL method:

- Uses HTTP GET method
 - Input is uploaded in URL field of HTTP request message

`www.somesite.com/animalsearch?monkeys&banana`

HTTP Response Message

status line
(protocol status
code status
phrase)

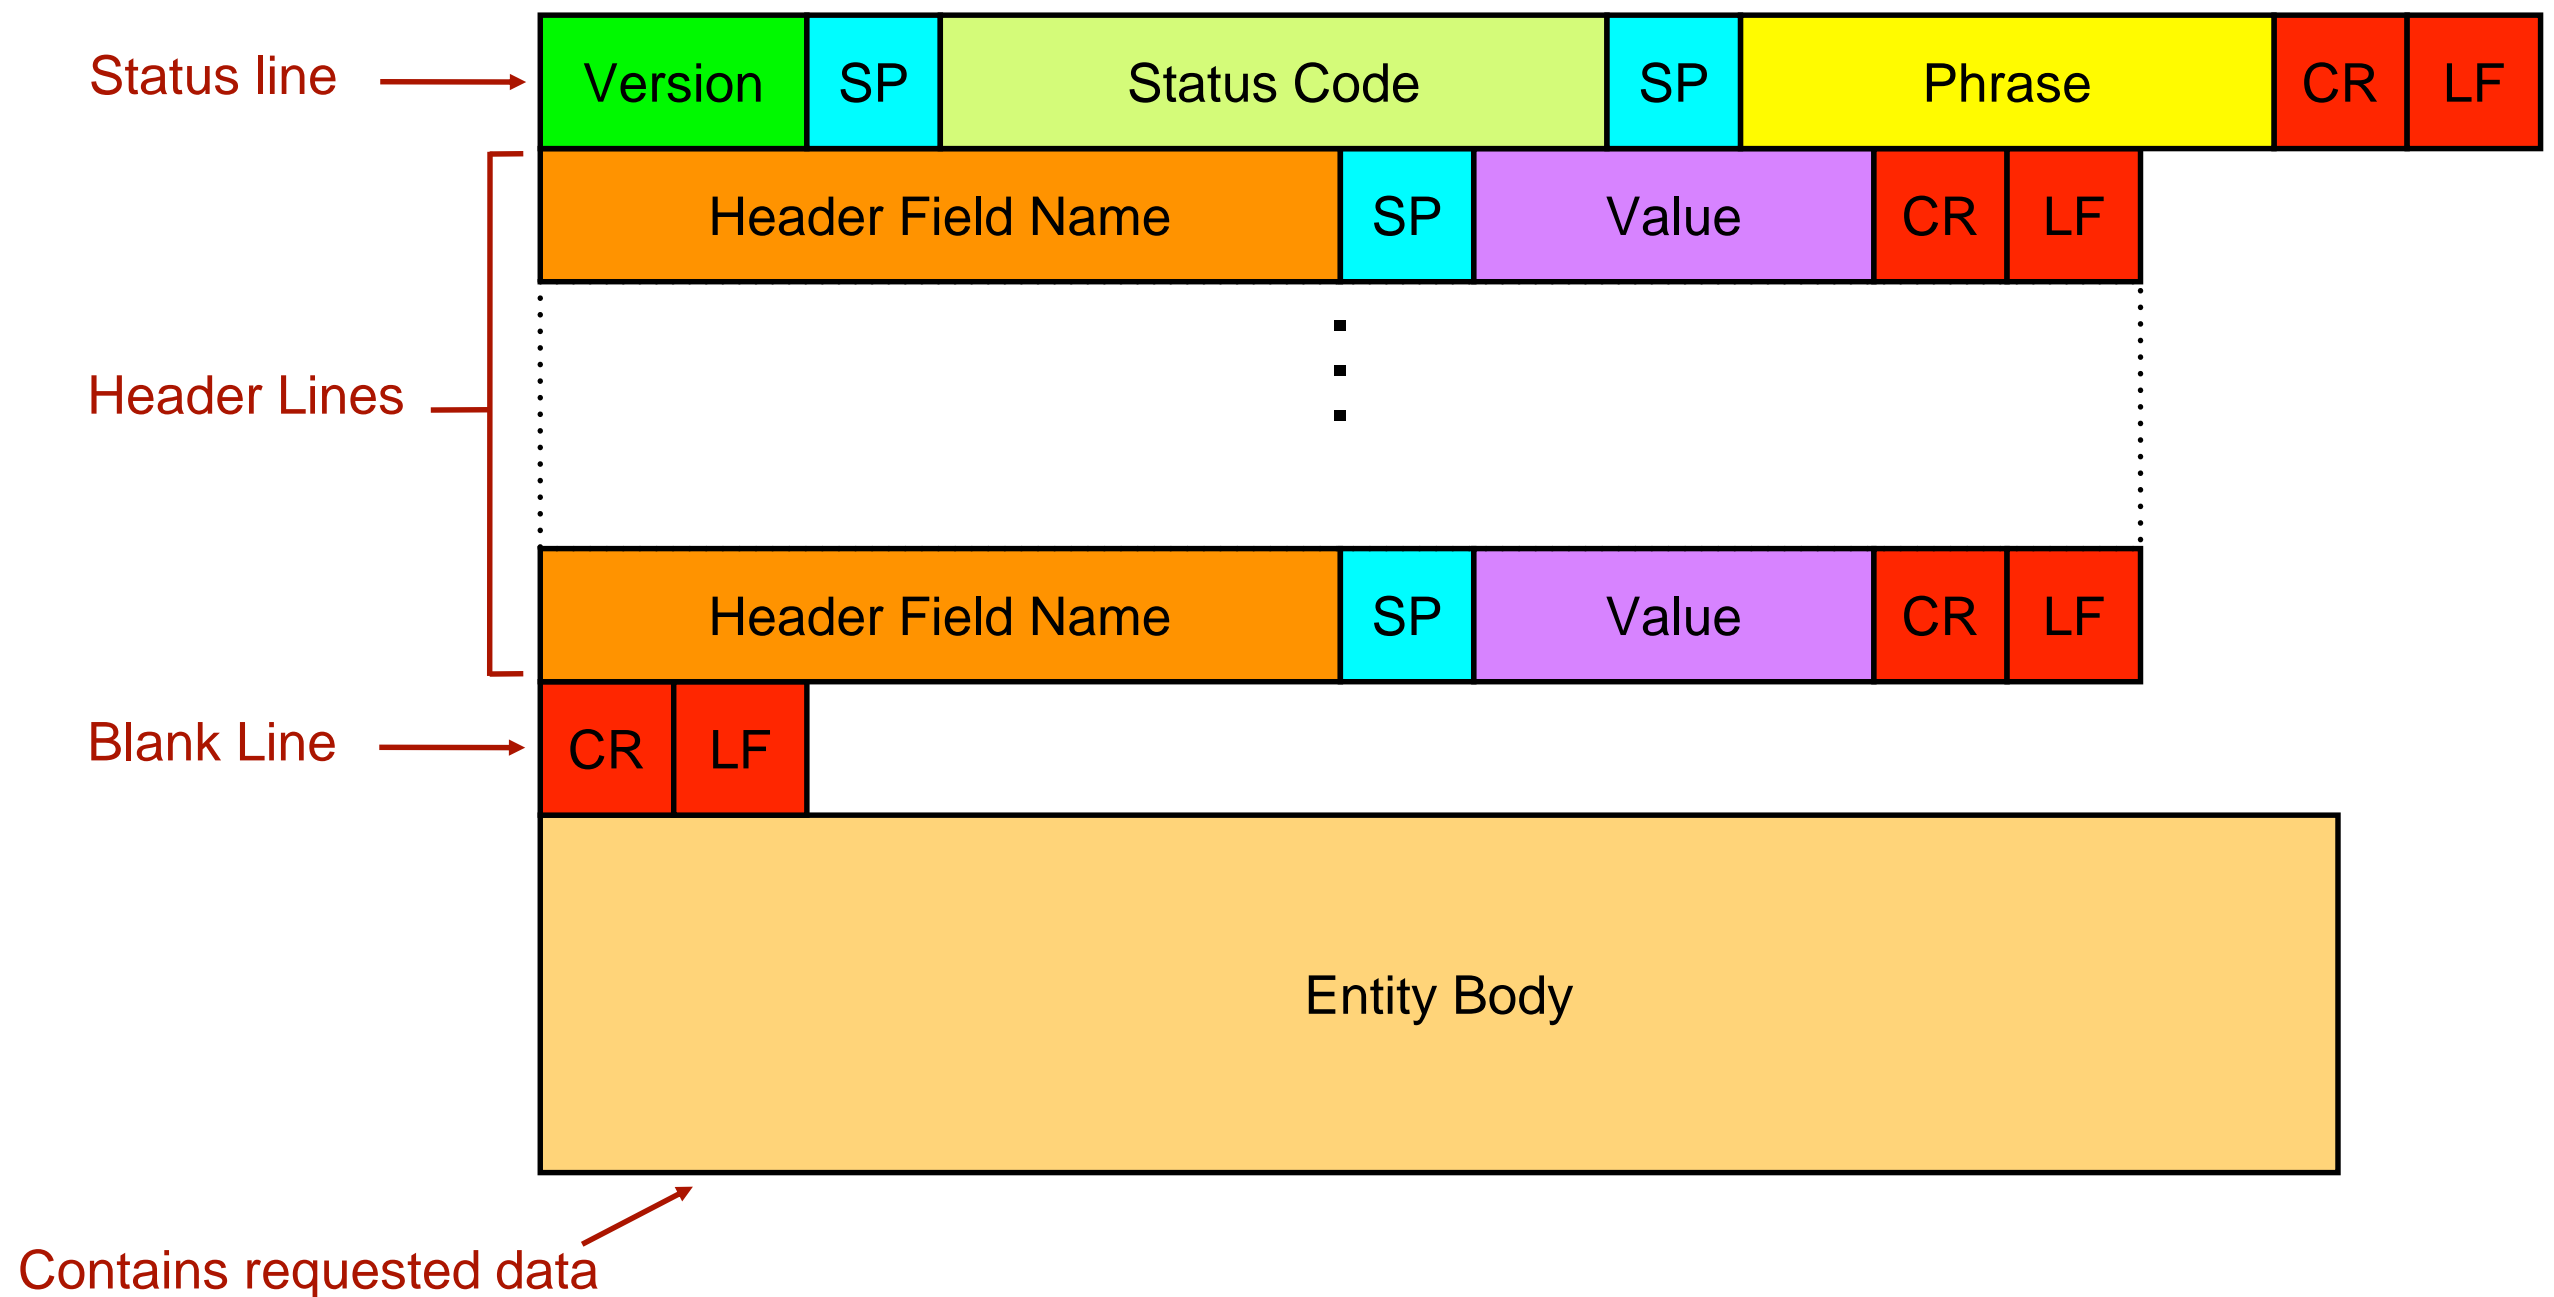
Header Lines

Carriage return,
line feed at start
of line indicates
end of header lines

data (e.g.
requested
HTML file)

```
HTTP/1.1 200 OK\r\n
Date: Sun, 26 Sep 2010 20:09:20 GMT\r\n
Server: Apache/2.0.52 (CentOS)\r\n
Last-Modified: Tue, 30 Oct 2007 17:00:02 GMT\r\n
ETag: "17dc6-a5c-bf716880"\r\n
Accept-Ranges: bytes\r\n
Content-Length: 2652\r\n
Keep-Alive: timeout=10, max=100\r\n
Connection: Keep-Alive\r\n
Content-Type: text/html; charset=ISO-8859-1\r\n
\r\n
data data data data data ...
```

HTTP Response Message: General Format



HTTP Response Status Codes

- Status code appears in first line in server-to-client response message

- Some sample codes:

200 OK

- Request succeeded, requested object later in this message

301 Moved Permanently

- Requested object moved, new location specified later in this msg (Location:)

400 Bad Request

- Request message not understood by server

404 Not Found

- Requested document not found on this server

505 HTTP Version Not Supported

Trying Out HTTP For Yourself (Client Side)

1. Telnet to your favorite Web server:

```
telnet faculty.ycp.edu 80
```

Opens a TCP connection to port 80 (default HTTP server port) at faculty.ycp.edu
Anything typed in sent to port 80 at faculty.ycp.edu

2. Type in a GET HTTP request:

```
GET /~jmoscola/test.html HTTP/1.1  
Host: faculty.ycp.edu
```

By typing this in (hit carriage return twice at the end), you send this minimal (but complete) GET request to the HTTP server

3. Look at response message sent by HTTP server!

Trying Out HTTP For Yourself (Another Example)

1. Telnet to your favorite Web server:

```
telnet faculty.ycp.edu 80
```

2. Type in a GET HTTP request:

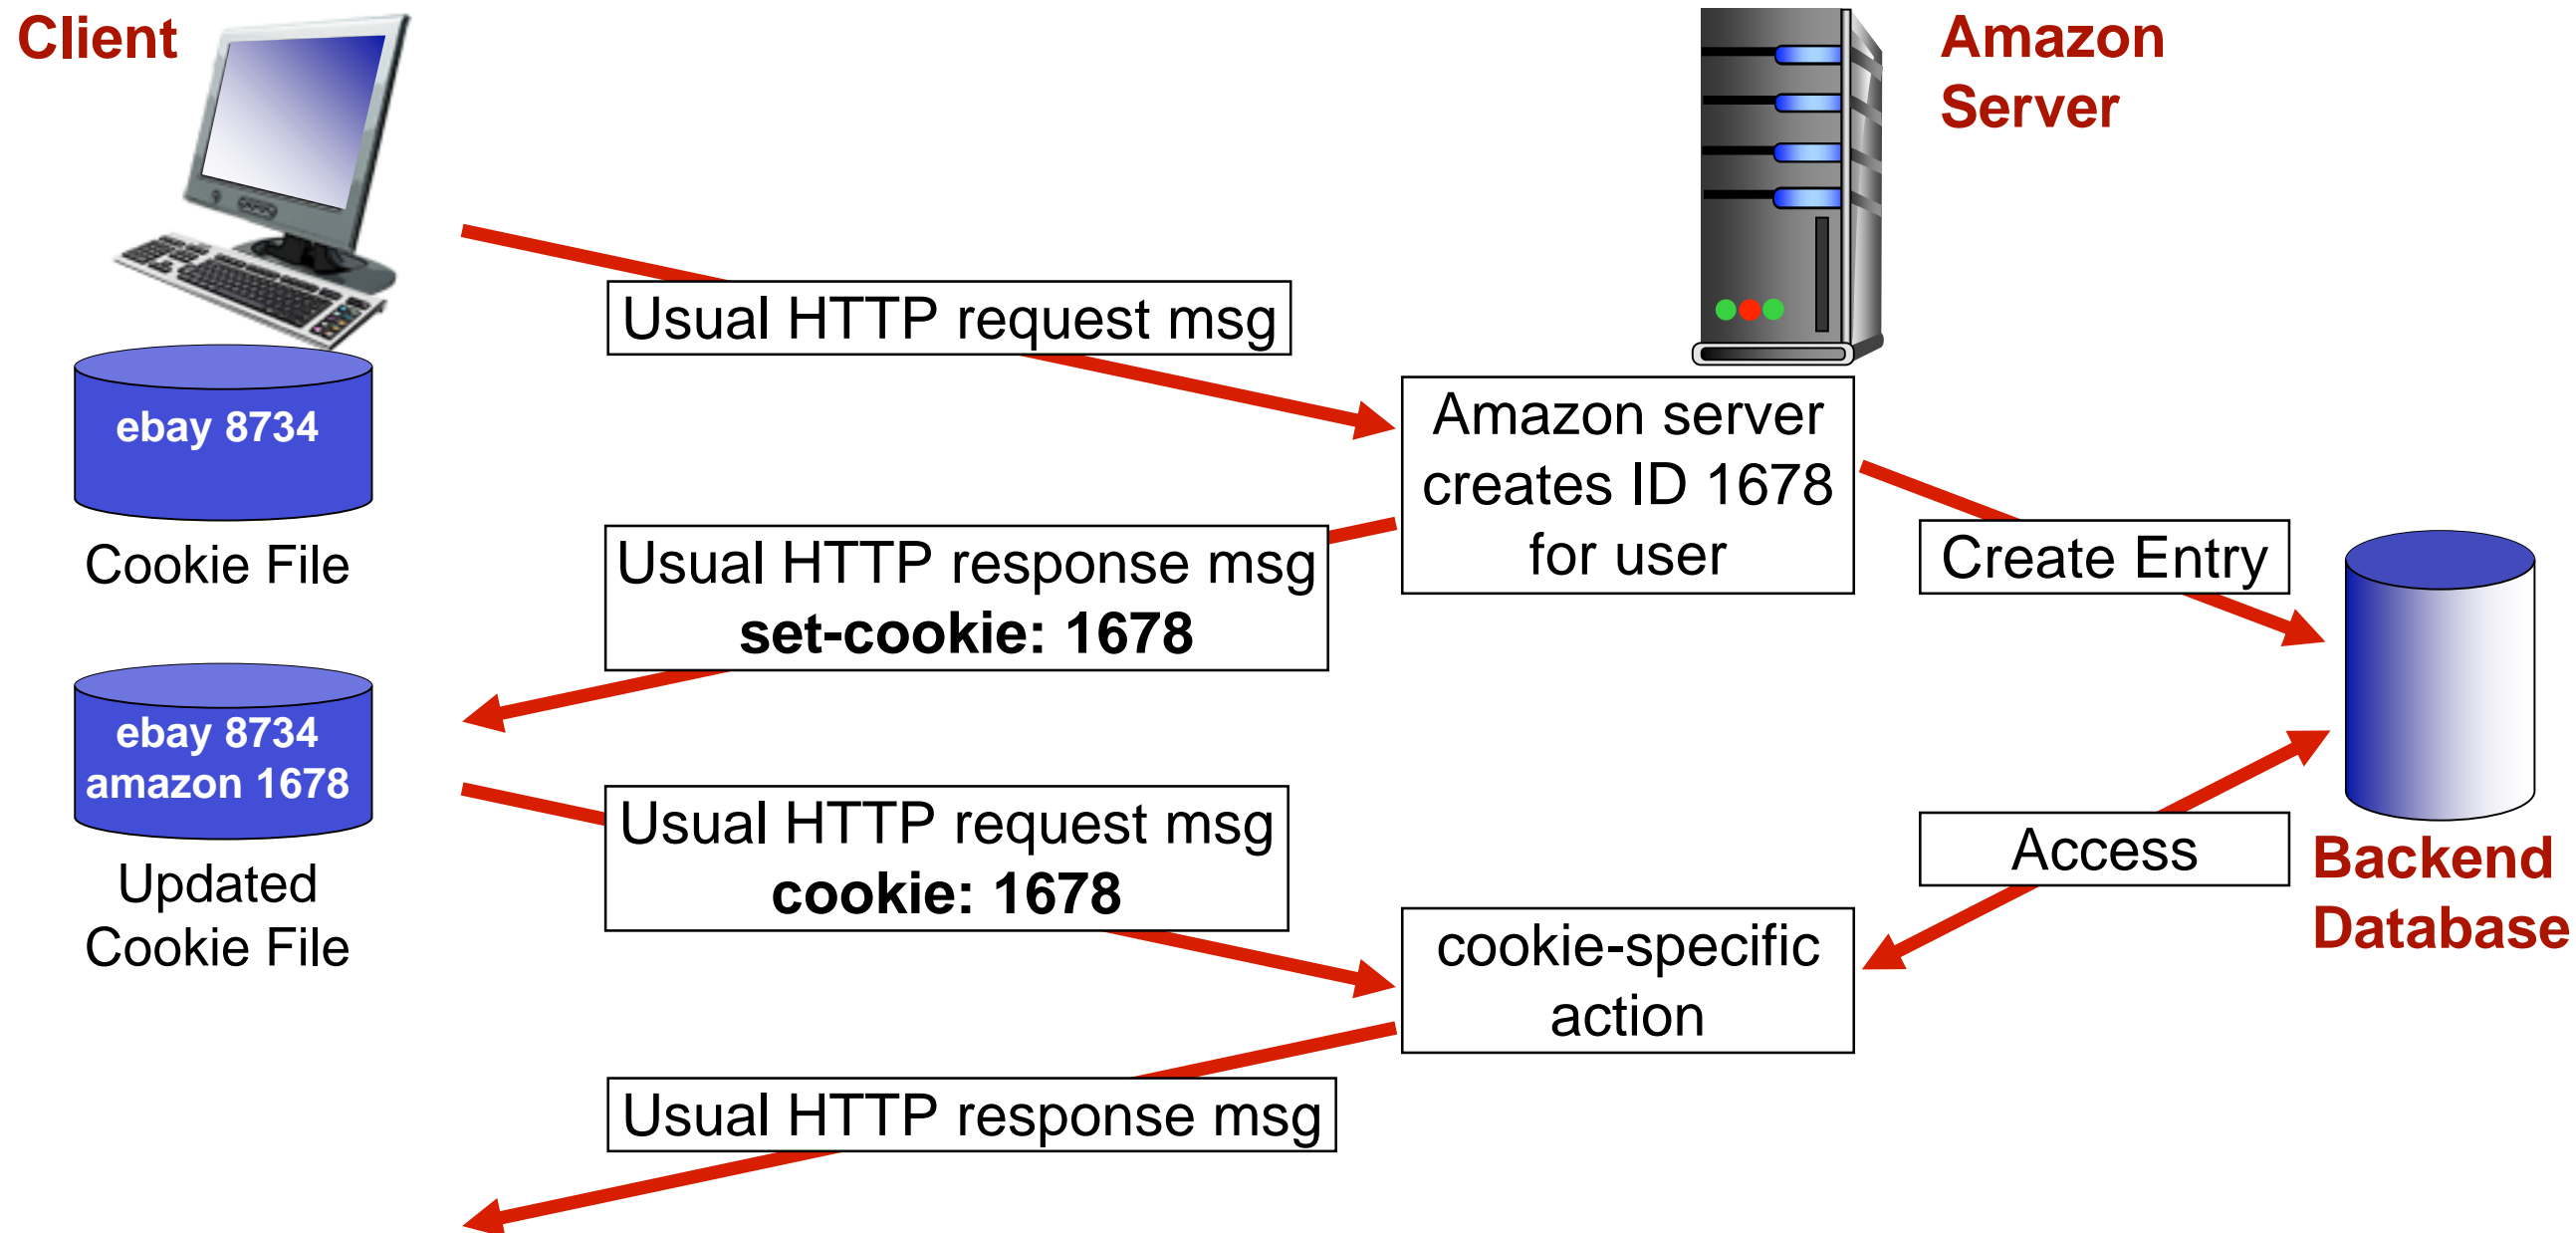
```
GET /~jmoscola/exam_solutions.html HTTP/1.1  
Host: faculty.ycp.edu
```

3. Look at response message sent by HTTP server!

User-Server State: Cookies

- **Cookies allow servers to remember previous information about user**
 - Stored in file on end user system since HTTP server is stateless
 - Many companies use cookies to identify a user
 - Content is dependent on the identity of the user
- **Four components to cookies:**
 - A cookie header line in the HTTP response message
 - A cookie header line in the HTTP request message
 - A cookie file kept on user's end system & managed by user's browser
 - A back-end database at web site

Cookies: Keeping “state”



Cookies

- **What cookies can be used for:**

- Authorization
- Shopping carts
- Recommendations
- User session state (web-based e-mail)

cookies and privacy:


- cookies permit sites to learn a lot about you
- you may supply name and e-mail to sites

- **How to keep “state”:**

- Protocol endpoints: maintain state at sender/receiver over multiple transactions
- Cookies: http messages carry state

Web Caches (proxy server)

- **A network entity that satisfies requests on behalf of an origin web server**
 - All requests are sent to proxy server
 - Proxy server caches objects
 - Only new objects are requested from origin server



More About Web Caching

- **Cache acts as both client and server**
 - Server for requesting clients
 - Client to the origin server
- **Typically, cache is installed by ISP (university, company, residential ISP)**
- **Why Web caching?**
 - Reduce response time for client request
 - Reduce traffic on an institution's access link


Caching example:

assumptions:

- avg object size: 100K bits
- avg request rate from browsers to origin servers: 15/sec
- avg data rate to browsers: 1.50 Mbps
- RTT from institutional router to any origin server: 2 sec
- access link rate: 1.54 Mbps

consequences:

- LAN utilization: 15%
- access link utilization = 99% *problem!*
- total delay = Internet delay + access delay + LAN delay
= 2 sec + minutes + usecs



Caching example: fatter access link

assumptions:


- avg object size: 100K bits
- avg request rate from browsers to origin servers: 15/sec
- avg data rate to browsers: 1.50 Mbps
- RTT from institutional router to any origin server: 2 sec
- access link rate: 1.54 Mbps

→ 154 Mbps

consequences:

- LAN utilization: 15%
- access link utilization = 99% → 9.9%
- total delay = Internet delay + access delay + LAN delay
= 2 sec + minutes + usecs

→ msec



Cost: increased access link speed (not cheap!)

Caching example: install local cache

assumptions:


- avg object size: 100K bits
- avg request rate from browsers to origin servers: 15/sec
- avg data rate to browsers: 1.50 Mbps
- RTT from institutional router to any origin server: 2 sec
- access link rate: 1.54 Mbps

consequences:

- LAN utilization: 15%
- access link utilization = ?
- total delay = ?

How to compute link utilization, delay?

Cost: web cache (cheap!)



Caching example: install local cache

**Calculating access link utilization,
delay with cache:**

suppose cache hit rate is 0.4

- 40% requests satisfied at cache - --- 60% requests satisfied at origin

access link utilization:

- 60% of requests use access link


data rate to browsers over access link

- = $0.6 * 1.50 \text{ Mbps} = .9 \text{ Mbps}$
- utilization = $0.9 / 1.54 = .58$

total delay

- = $0.6 * (\text{delay from origin servers}) + 0.4 * (\text{delay when satisfied at cache})$
- = $0.6 (2.01) + 0.4 (\sim \text{msecs}) = \sim 1.2 \text{ secs}$

less than with 154 Mbps link (and cheaper too!)



Conditional GET

- **Goal: don't send object if cache has up-to-date cached version**
 - No object transmission delay
 - lower link utilization
- **Cache: specify date of cached copy in HTTP request**
 - If-modified-since: <date>
- **Server: response contains no object if cached copy is up-to-date:**
 - HTTP/1.0 304 Not Modified

Client



Server

