

# CS 330: Network Applications & Protocols

Application Layer: P2P, Sockets

---

Galin Zhelezov  
Department of Physical Sciences  
York College of Pennsylvania




# Overview of Application Layer

---

- **Network Application Architectures**
- **HyperText Transfer Protocol (HTTP)**
- **File Transfer and Email protocols (FTP, SMTP)**
- **Domain Name System (DNS)**
- **Peer-to-Peer Applications (P2P)**
  - File Distribution
- **Video Streaming and Content Distribution Networks**
- **Socket Programming**

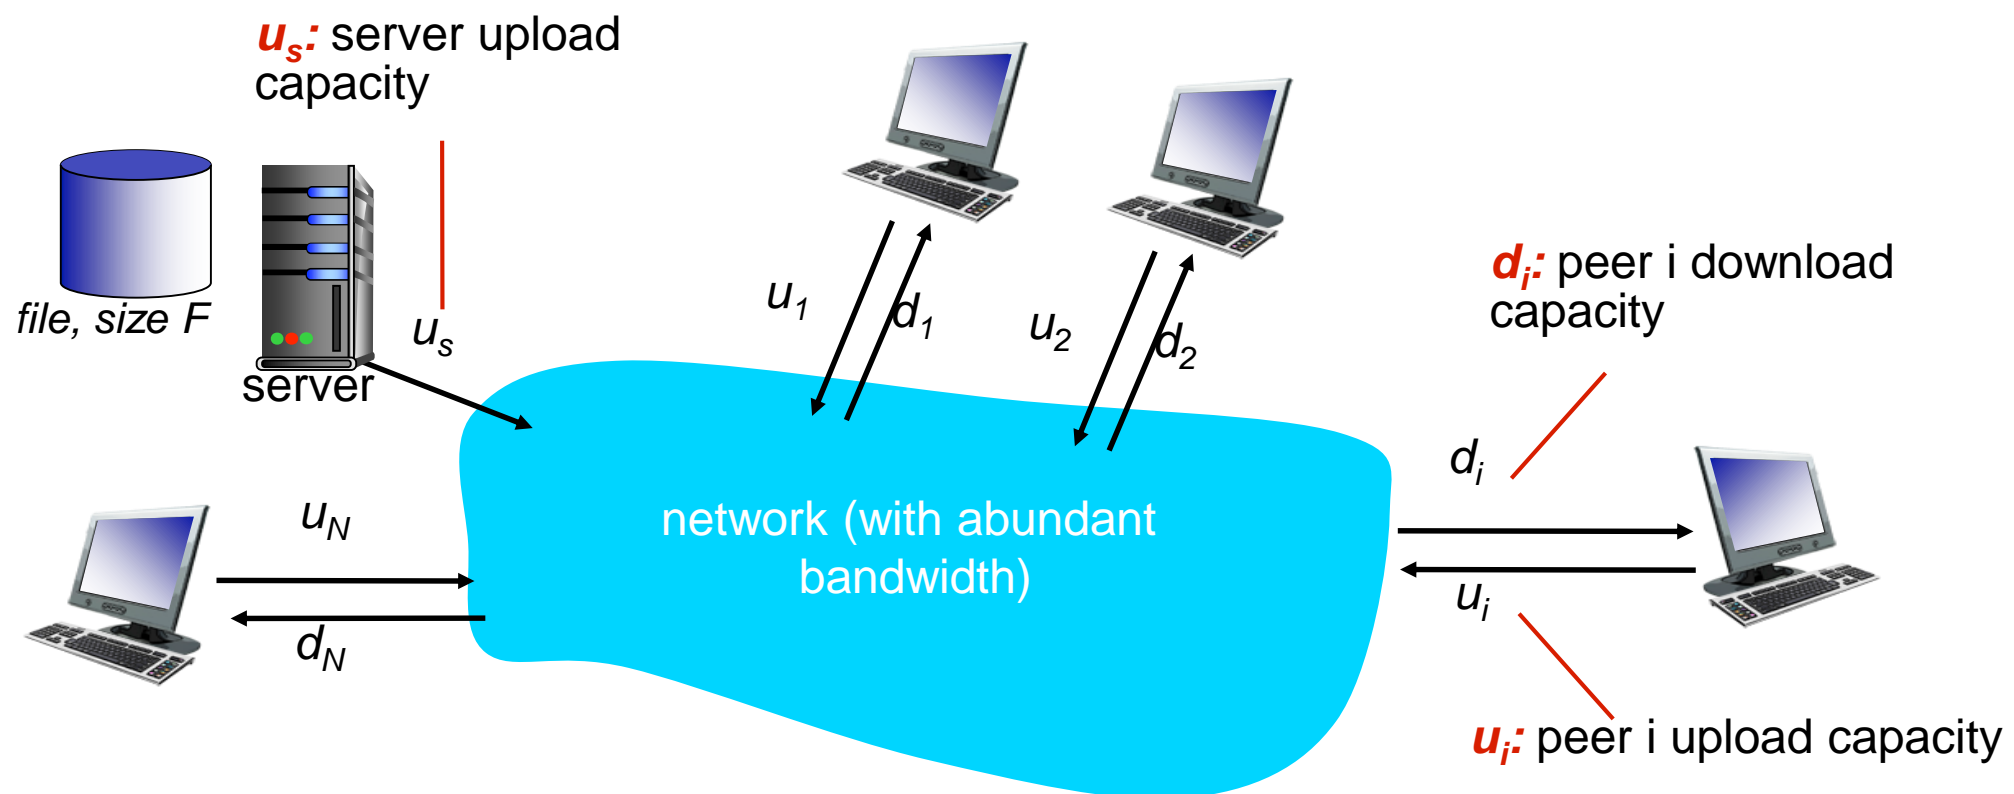
# Peer-to-Peer Architecture

- Does not require always-on servers
- Hosts communicate directly with each other
- Peers request service from other peers, and provide service in return to other peers
- Highly scalable
- Self scalability – new peers bring new service capacity, as well as new service demands
- Hosts are intermittently connected and may change IP addresses
- Difficult to authenticate - possibly insecure
- Hosts need incentive to share data



# File Distribution: Client-Server vs P2P


- **Question:** How much time does it take to distribute a file of size  $F$  from one server to  $N$  peers?
  - Assume the core network has abundant bandwidth
  - Assume clients are not using bandwidth for other tasks



# File Distribution: Client-Server

- **Server transmission: must sequentially upload  $N$  copies of the file:**

- Time to send one copy:  $F/u_s$
- Time to send  $N$  copies:  $(N * F) / u_s$



- **Client: each client must download copy of the file**

- $d_{min}$  = download rate of slowest client
- Minimum download time for clients:  $F/d_{min}$

Time to distribute file of size  $F$  to  $N$  clients using client-server approach



$$D_{cs} \geq \max \left\{ \frac{NF}{u_s}, \frac{F}{d_{min}} \right\}$$

**\*\*Distribution time increases linearly with the number of clients!**

# File Distribution: P2P

- **Server transmission: must upload at least one copy of the file**


- Time to send one copy:  $F/u_s$

- **Client: each client must download a copy of the file**

- Minimum download time for clients:  $F/d_{min}$

- **Clients: as an aggregate must download  $NF$  bits**

- Total upload capacity is upload rate of server plus upload rate of each of the individual peers:  $u_s + \sum u_i$



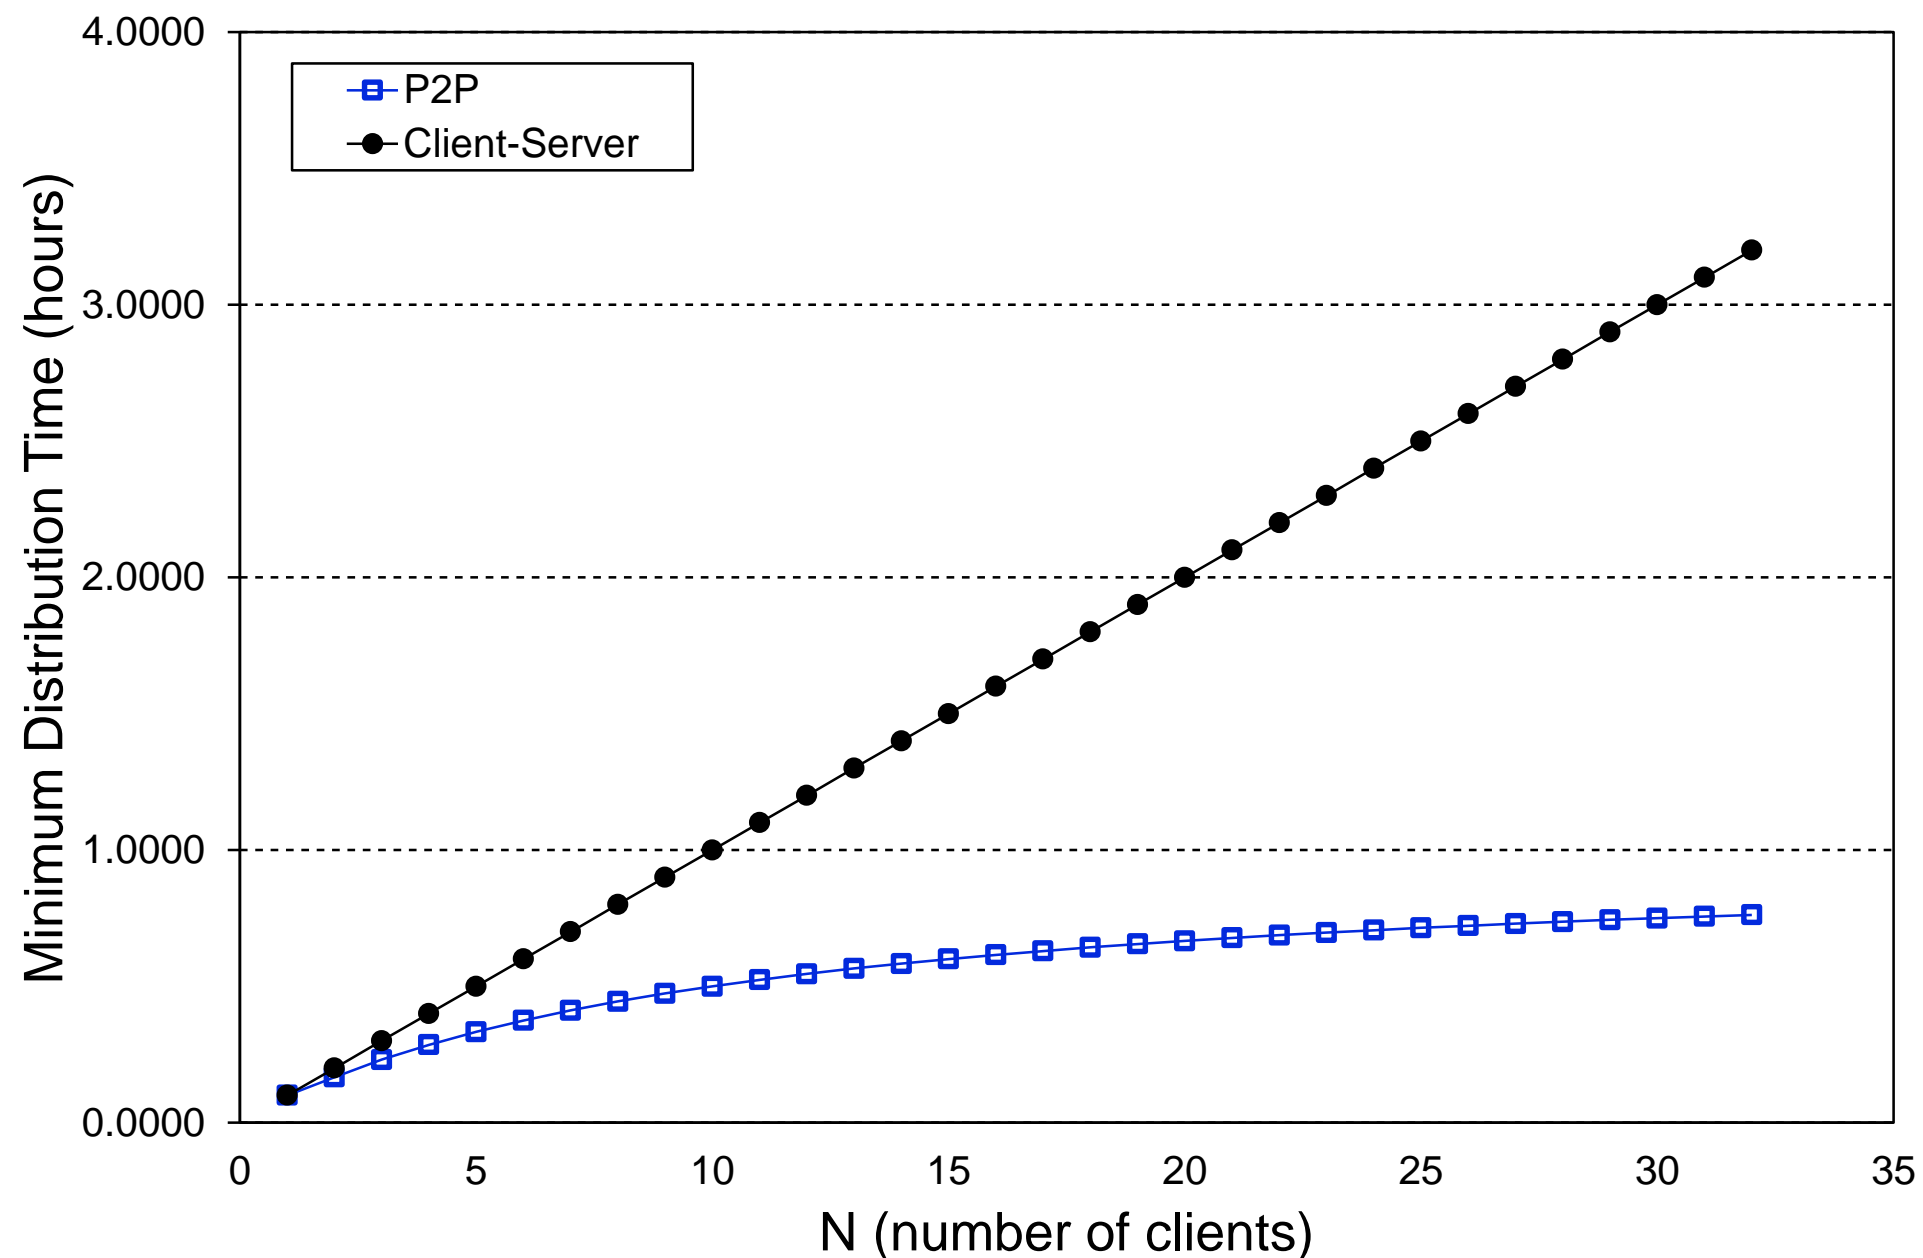
Time to distribute file of size  $F$  to  $N$  clients using P2P approach

$$\Rightarrow D_{P2P} \geq \max \left\{ \frac{F}{u_s}, \frac{F}{d_{min}}, \frac{NF}{u_s + \sum_{i=1}^N u_i} \right\}$$

**\*\*As  $N$  increases, so does service capacity**

# Client-Server vs. P2P

Client upload rate =  $u$ ,  $F/u = 1$  hour,  $u_s = 10u$ ,  $d_{min} \geq u_s$




# P2P File Distribution: BitTorrent

- File divided into 256 KByte chunks
- Peers in torrent send/receive file chunks

*tracker*: tracks peers participating in torrent

*torrent*: group of peers exchanging chunks of a file

Alice arrives ...  
... obtains list of peers from tracker  
... and begins exchanging file chunks with peers in torrent





# P2P File Distribution: BitTorrent

- **Peer joining torrent:**


- Has no chunks, but will accumulate them over time from other peers
- Registers with tracker to get list of peers, connects to subset of peers (“neighbors”)

- **While downloading, peer uploads chunks to other peers**

- **Peer may change peers with whom it exchanges chunks**

- **Churn:** peers may come and go

- **Once peer has entire file, it may (selfishly) leave or (altruistically) remain in torrent**



# BitTorrent: Requesting/Sending File Chunks

---

- **Requesting chunks:**

- At any given time, different peers have different subsets of file chunks
- Periodically, Alice asks each peer for a list of chunks that they have
- Alice requests missing chunks from peers, rarest first


- **Sending chunks: tit-for-tat**

- Alice sends chunks to the four peers currently sending her chunks at highest rate
  - Other peers are choked by Alice (do not receive chunks from her)
  - Re-evaluate top 4 every 10 secs
- Every 30 secs, randomly select another peer, start sending chunks
  - Optimistically unchoke this peer
  - Newly chosen peer may join top four

- **Sending chunks:**

# BitTorrent: tit-for-tat

- (1) Alice “optimistically unchokes” Bob
- (2) Alice becomes one of Bob’s top-four providers; Bob reciprocates
- (3) Bob becomes one of Alice’s top-four providers



# Overview of Application Layer

---

- **Network Application Architectures**
- **HyperText Transfer Protocol (HTTP)**
- **File Transfer and Email protocols (FTP, SMTP)**
- **Domain Name System (DNS)**
- **Peer-to-Peer Applications (P2P)**
- **Video Streaming and Content Distribution Networks**
- **Socket Programming**

# Video Streaming and CDNs: context

---


- video traffic: major consumer of Internet bandwidth
  - Netflix, YouTube: 37%, 16% of downstream residential ISP traffic
    - ~1B YouTube users, ~75M Netflix users
- challenge: scale - how to reach ~1B users?
  - single mega-video server won't work (why?)
    - challenge: heterogeneity
- different users have different capabilities (e.g., wired versus mobile; bandwidth rich versus bandwidth poor)
- **solution:** distributed, application-level infrastructure



# Multimedia: video

- **video: sequence of images displayed at constant rate**
  - e.g., 24 images/sec
- **digital image: array of pixels**
  - each pixel represented by bits
- **coding: use redundancy *within* and *between* images to decrease # bits used to encode image**
  - spatial (within image)
  - temporal (from one image to next)

*spatial coding example:* instead of sending  $N$  values of same color (all purple), send only two values: color value (*purple*) and number of repeated values ( $N$ )



frame  $i$

*temporal coding example:*  
instead of sending complete frame at  $i+1$ ,  
send only differences from frame  $i$




frame  $i+1$

# Multimedia: video

- **CBR: (constant bit rate):** video encoding rate fixed
- **VBR: (variable bit rate):** video encoding rate changes as amount of spatial, temporal coding changes
- **examples:**
  - MPEG I (CD-ROM) 1.5 Mbps
  - MPEG2 (DVD) 3-6 Mbps
  - MPEG4 (often used in Internet, < 1 Mbps)

*spatial coding example:* instead of sending  $N$  values of same color (all purple), send only two values: color value (*purple*) and number of repeated values ( $N$ )



frame  $i$

*temporal coding example:* instead of sending complete frame at  $i+1$ , send only differences from frame  $i$




frame  $i+1$

# Streaming stored video:

---

simple scenario:





# Streaming multimedia: DASH

---

- **DASH: Dynamic, Adaptive Streaming over HTTP**
- **server:**
  - divides video file into multiple chunks
  - each chunk stored, encoded at different rates
  - *manifest file*: provides URLs for different chunks
- **client:**
  - periodically measures server-to-client bandwidth
  - consulting manifest, requests one chunk at a time
    - chooses maximum coding rate sustainable given current bandwidth
    - can choose different coding rates at different points in time (depending on available bandwidth at time)

# Streaming multimedia: DASH

---

- ***DASH: Dynamic, Adaptive Streaming over HTTP***
- ***“intelligence” at client: client determines***
  - *when* to request chunk (so that buffer starvation, or overflow does not occur)
  - *what encoding rate* to request (higher quality when more bandwidth available)
  - *where* to request chunk (can request from URL server that is “close” to client or has high available bandwidth)

# Content distribution networks

---

- **challenge:** how to stream content (selected from millions of videos) to hundreds of thousands of *simultaneous* users?
- **option 1:** single, large “mega-server”
  - single point of failure
  - point of network congestion
  - long path to distant clients
  - multiple copies of video sent over outgoing link

....quite simply: this solution ***doesn't scale***


# Content distribution networks

---

- ***challenge:*** how to stream content (selected from millions of videos) to hundreds of thousands of simultaneous users?
- ***option 2:*** store/serve multiple copies of videos at multiple geographically distributed sites (***CDN***)
  - ***enter deep:*** push CDN servers deep into many access networks
    - close to users
    - used by Akamai, 1700 locations
  - ***bring home:*** smaller number (10's) of larger clusters in POPs near (but not within) access networks
    - used by Limelight

# Content Distribution Networks (CDNs)

- CDN: stores copies of content at CDN nodes
  - e.g. Netflix stores copies of MadMen
- subscriber requests content from CDN
  - directed to nearby copy, retrieves content
  - may choose different copy if network path congested



# Content Distribution Networks (CDNs)



## **OTT challenges:** coping with a congested Internet

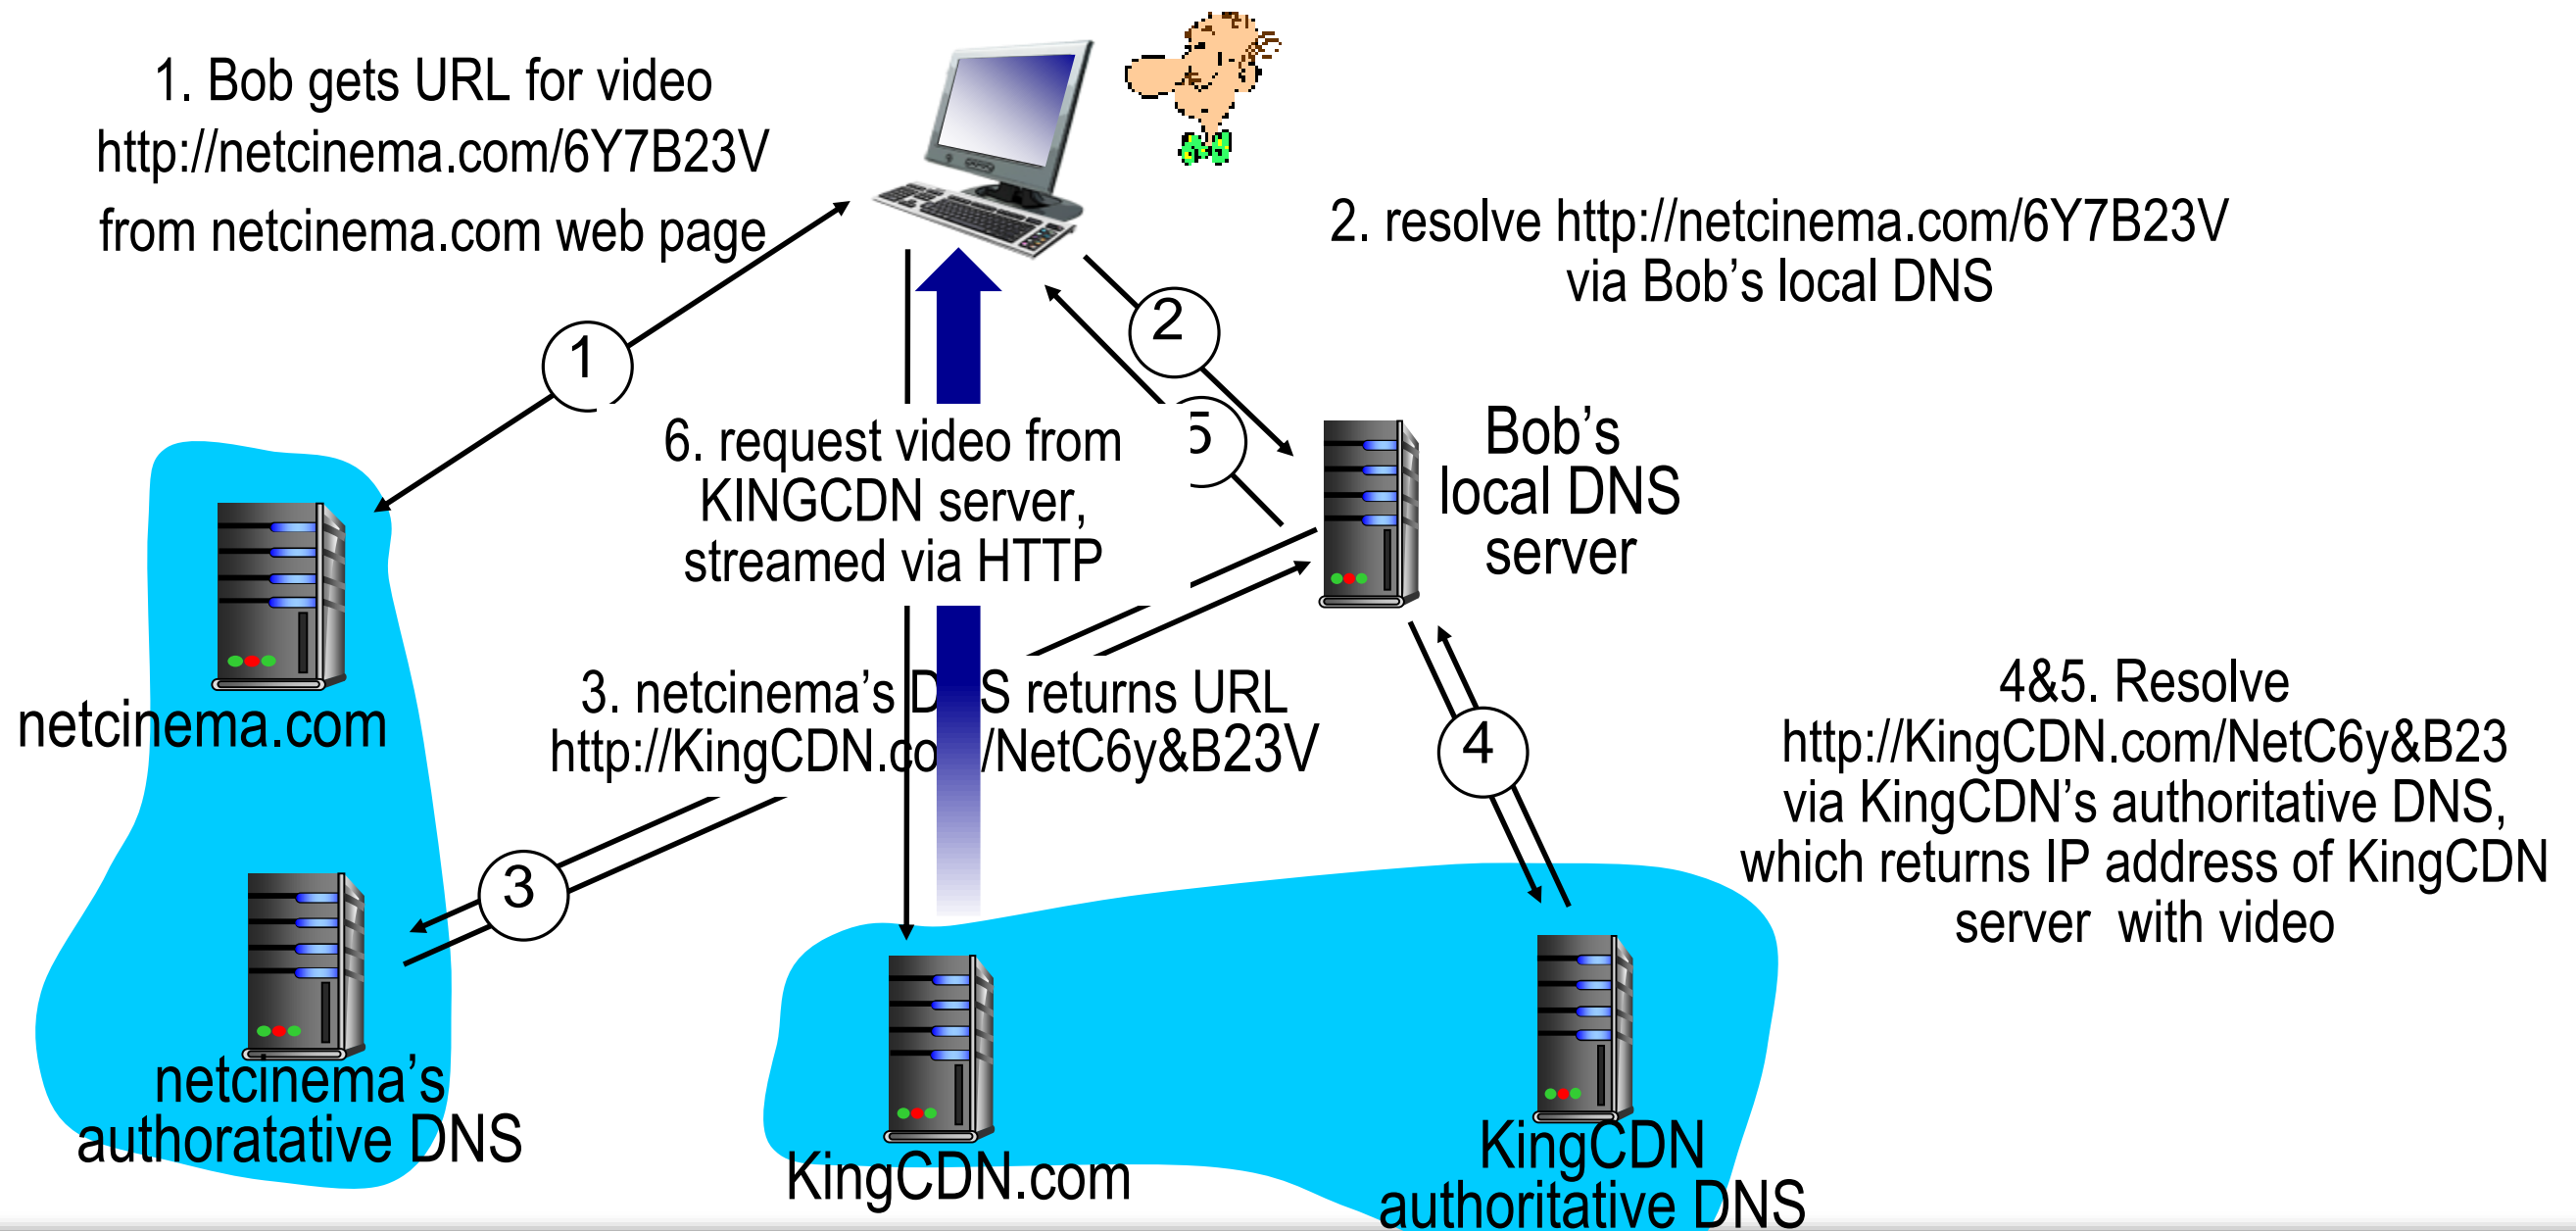
- from which CDN node to retrieve content?
- viewer behavior in presence of congestion?
- what content to place in which CDN node?

*more .. in chapter 7*

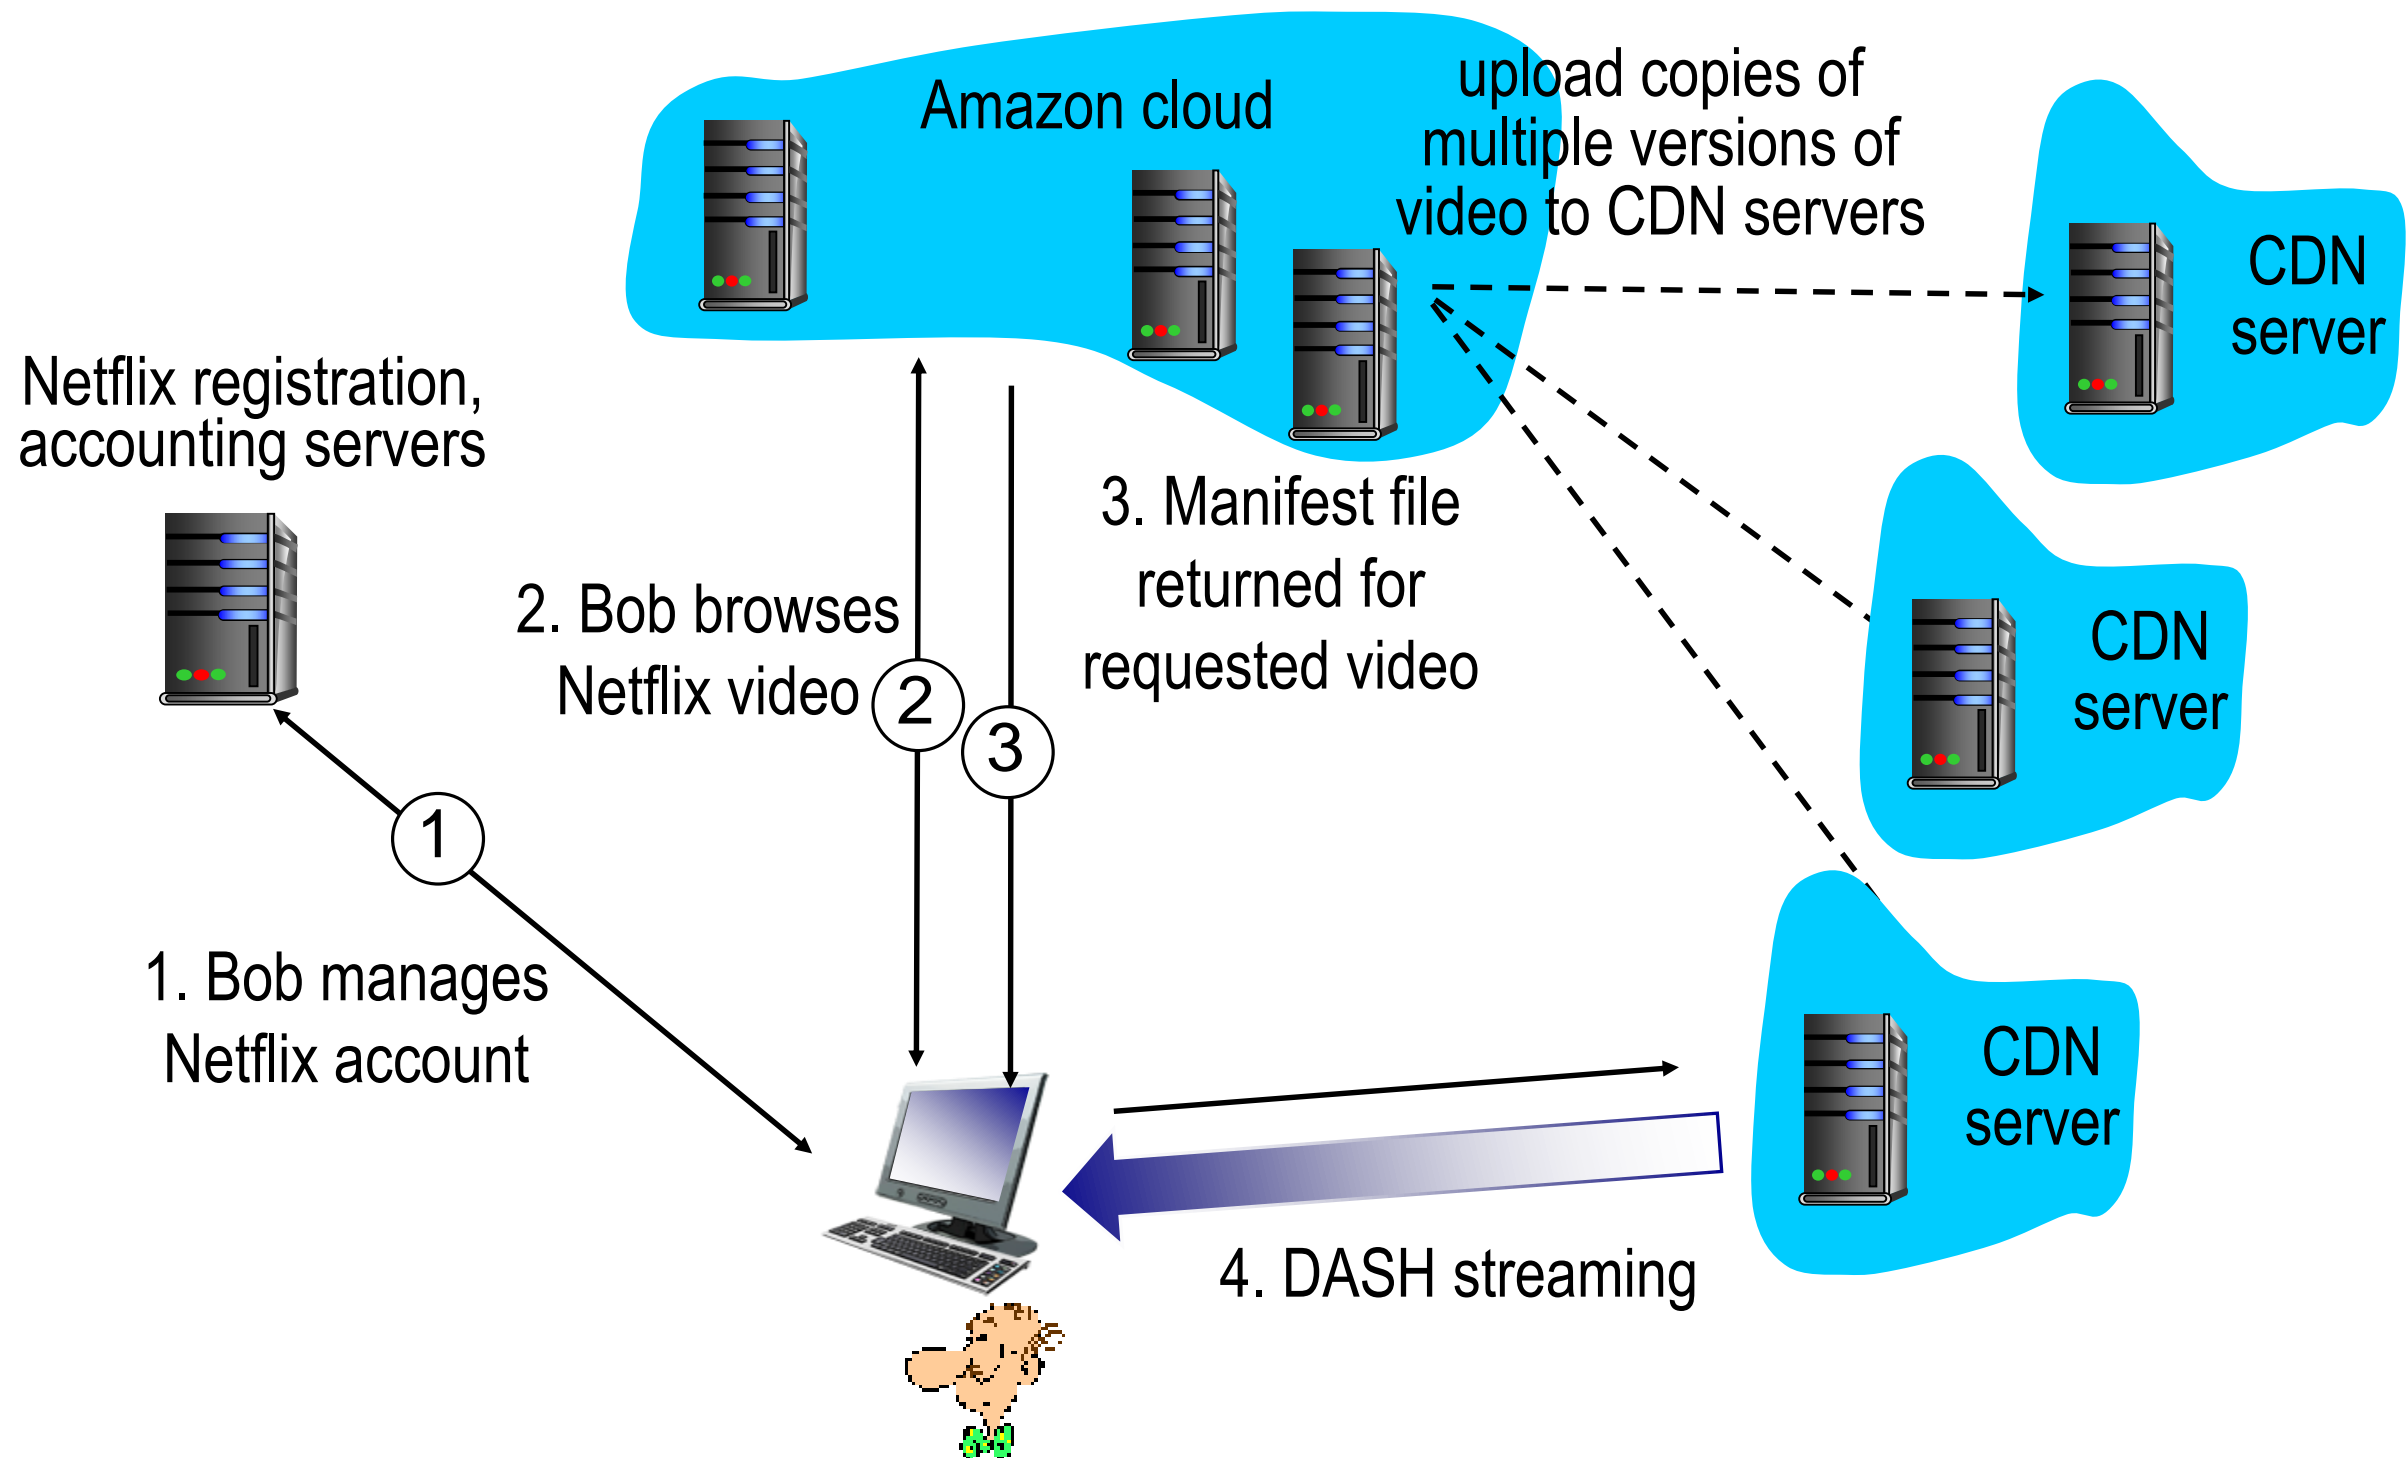
# CDN content access: a closer look

Bob (client) requests video `http://netcinema.com/6Y7B23V`

- video stored in CDN at `http://KingCDN.com/NetC6y&B23V`



# Case study: Netflix





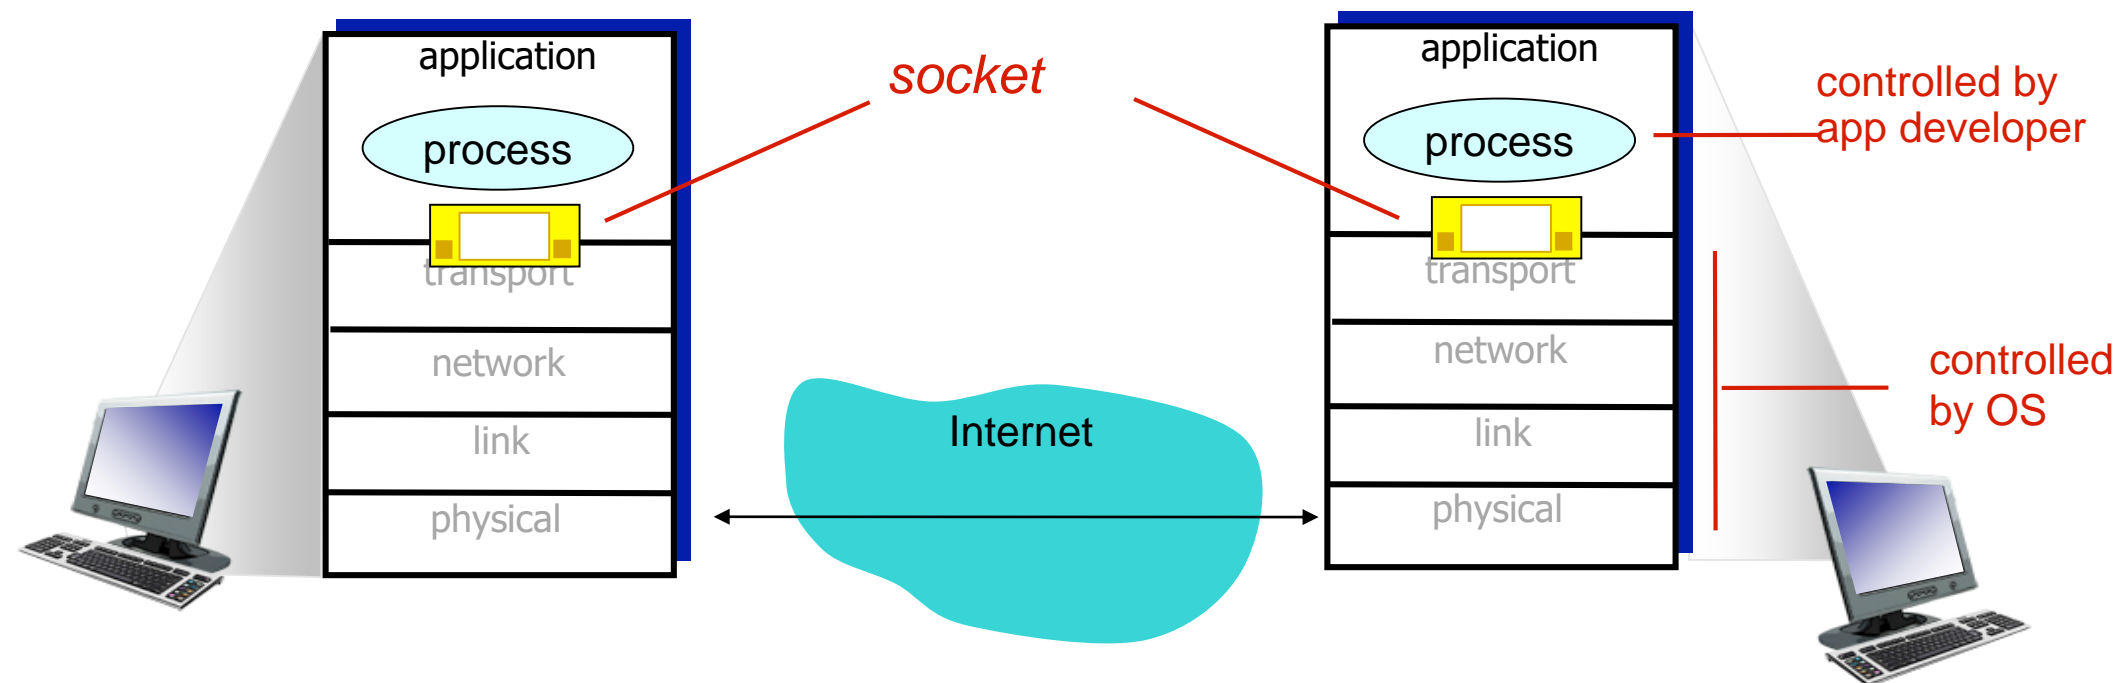
# Overview of Application Layer

---

- **Network Application Architectures**
- **HyperText Transfer Protocol (HTTP)**
- **File Transfer and Email protocols (FTP, SMTP)**
- **Domain Name System (DNS)**
- **Peer-to-Peer Applications (P2P)**
- **Video Streaming and Content Distribution Networks**
- **Socket Programming**

# Socket Programming

- **Goal:** learn how to build client/server applications that communicate using sockets
- **Socket:** door between application process and end-end-transport protocol



# Socket Programming

---

- **Two socket types for two transport services:**
  - **UDP**: unreliable datagram
  - **TCP**: reliable, byte stream-oriented
- **Application Example:**
  - (1) Client reads a line of characters (data) from its keyboard and sends the data to a server.
  - (2) The server receives the data and converts characters to uppercase.
  - (3) The server sends the modified data to the client.
  - (4) The client receives the modified data and displays the line on its screen.

# Socket Programming with TCP

---

- **Client must contact server**
  - Server process must first be running
  - Server must have created socket that welcomes client's connection
- **Client contacts server by:**
  - Creating TCP socket, specifying IP address, port number of server process
  - When client creates socket, client TCP establishes connection to server TCP
- **When contacted by client, server TCP creates new socket for server process to communicate with that particular client**
  - Allows server to talk with multiple clients
  - Source port numbers used to distinguish clients
- **Application viewpoint:**
  - TCP provides reliable, in-order byte-stream transfer (a "pipe") between client and server

# Socket Programming with UDP

---

- **UDP: no “connection” between client & server**
  - No handshaking before sending data
  - Sender explicitly attaches IP destination address and port number to each packet
  - Receiver extracts sender IP address and port number from received packet
- **Transmitted data may be lost or received out-of-order**
- **Application viewpoint:**
  - UDP provides unreliable transfer of datagrams between client and server

# Chapter 2: summary

---

*most importantly: learned about protocols!*

- **typical request/reply message exchange:**

- client requests info or service
- server responds with data, status code

- **message formats:**

- *headers*: fields giving info about data
- *data*: info(payload) being communicated

*important themes:*

- control vs. messages
  - in-band, out-of-band
- centralized vs. decentralized
- stateless vs. stateful
- reliable vs. unreliable message transfer
- “complexity at network edge”