

CS 330: Network Applications & Protocols

Application Layer: P2P, Sockets

Galin Zhelezov
Department of Physical Sciences
York College of Pennsylvania




Overview of Application Layer

- **Network Application Architectures**
- **HyperText Transfer Protocol (HTTP)**
- **File Transfer and Email protocols (FTP, SMTP)**
- **Domain Name System (DNS)**
- **Peer-to-Peer Applications (P2P)**
 - File Distribution
- **Video Streaming and Content Distribution Networks**
- **Socket Programming**

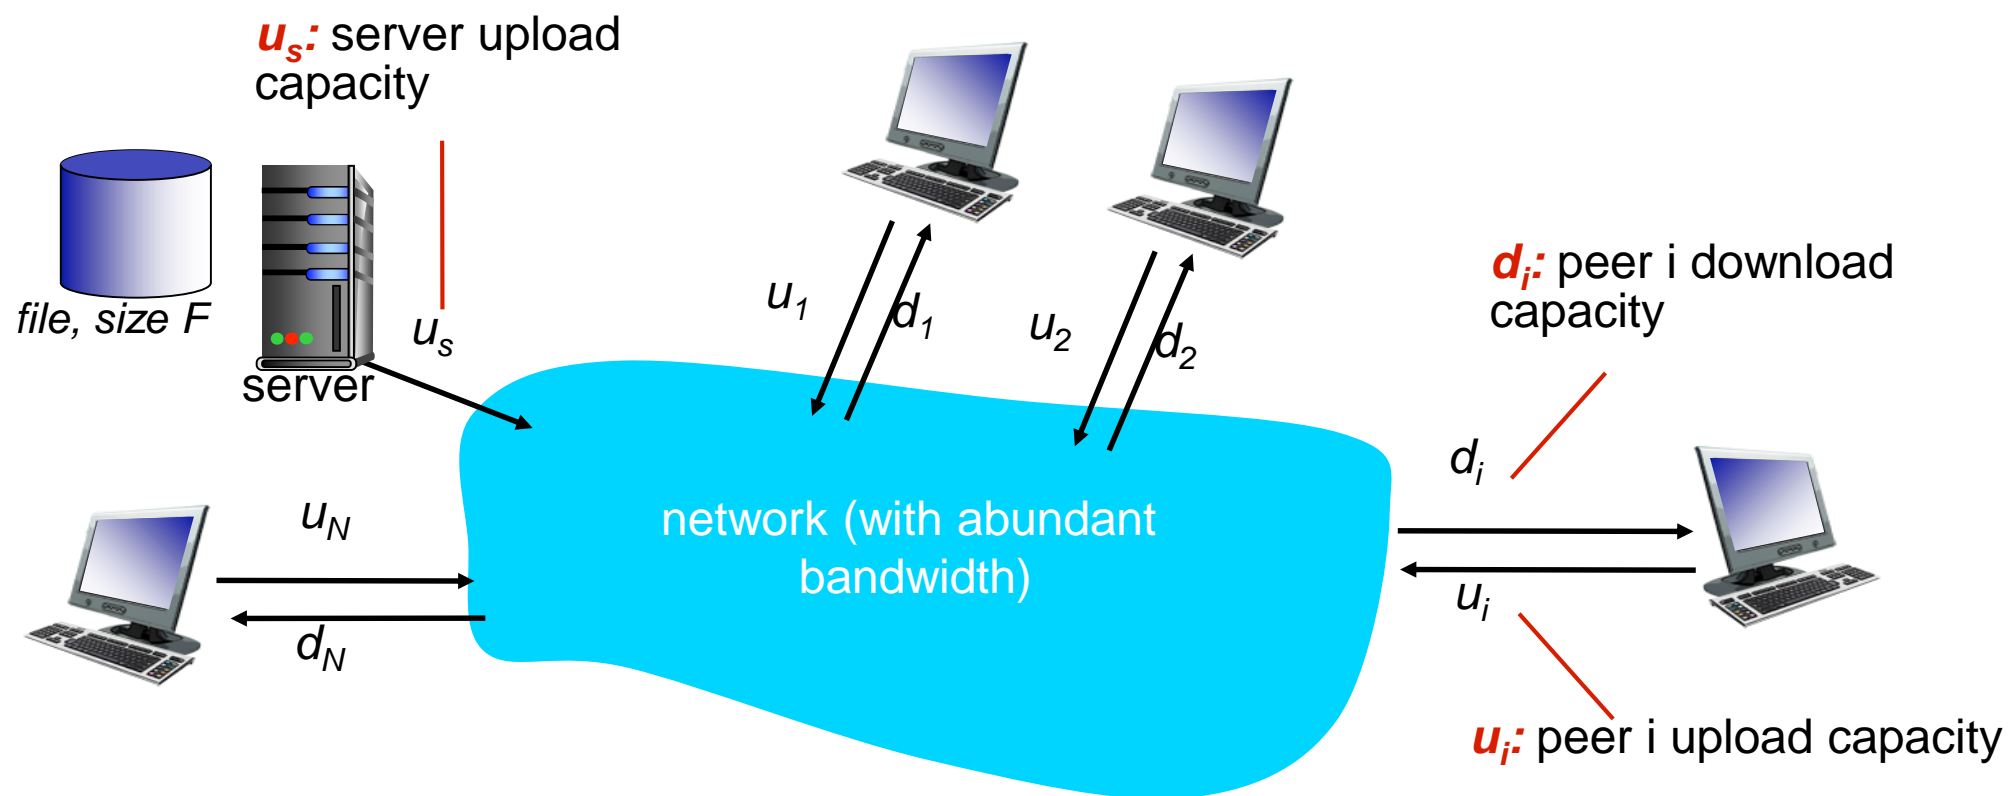
Peer-to-Peer Architecture

- Does not require always-on servers
- Hosts communicate directly with each other
- Peers request service from other peers, and provide service in return to other peers
- Highly scalable
- Self scalability – new peers bring new service capacity, as well as new service demands
- Hosts are intermittently connected and may change IP addresses
- Difficult to authenticate - possibly insecure
- Hosts need incentive to share data



File Distribution: Client-Server vs P2P


- **Question:** How much time does it take to distribute a file of size F from one server to N peers?
 - Assume the core network has abundant bandwidth
 - Assume clients are not using bandwidth for other tasks



File Distribution: Client-Server

- **Server transmission: must sequentially upload N copies of the file:**

- Time to send one copy: F/u_s
- Time to send N copies: $(N * F) / u_s$



- **Client: each client must download copy of the file**

- d_{min} = download rate of slowest client
- Minimum download time for clients: F/d_{min}

Time to distribute file of size F to N clients using client-server approach



$$D_{cs} \geq \max \left\{ \frac{NF}{u_s}, \frac{F}{d_{min}} \right\}$$

****Distribution time increases linearly with the number of clients!**

File Distribution: P2P

- **Server transmission: must upload at least one copy of the file**


- Time to send one copy: F/u_s

- **Client: each client must download a copy of the file**

- Minimum download time for clients: F/d_{min}

- **Clients: as an aggregate must download NF bits**

- Total upload capacity is upload rate of server plus upload rate of each of the individual peers: $u_s + \sum u_i$



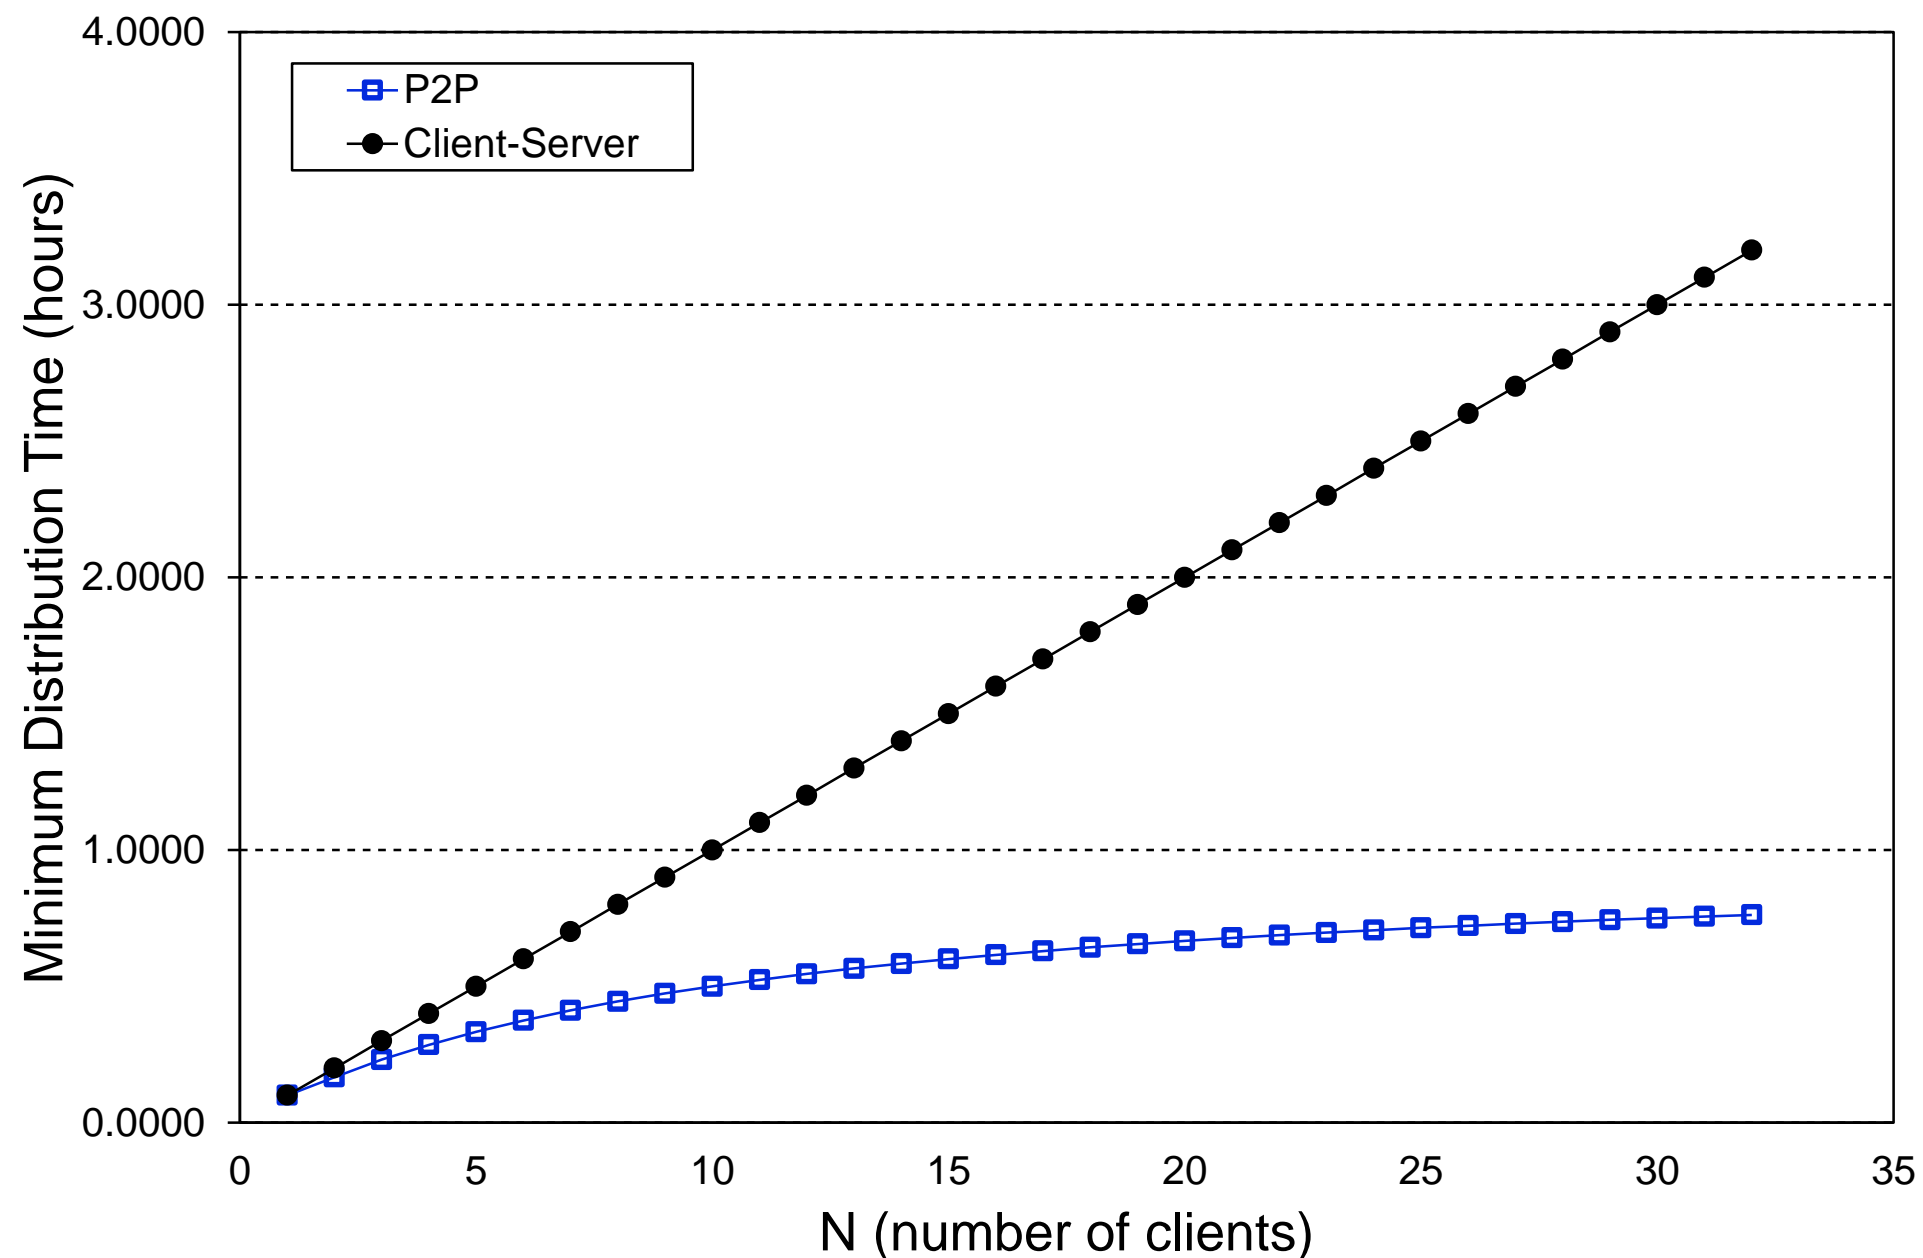
Time to distribute file of size F to N clients using P2P approach

$$\Rightarrow D_{P2P} \geq \max \left\{ \frac{F}{u_s}, \frac{F}{d_{min}}, \frac{NF}{u_s + \sum_{i=1}^N u_i} \right\}$$

****As N increases, so does service capacity**

Client-Server vs. P2P

Client upload rate = u , $F/u = 1$ hour, $u_s = 10u$, $d_{min} \geq u_s$




P2P File Distribution: BitTorrent

- File divided into 256 KByte chunks
- Peers in torrent send/receive file chunks

tracker: tracks peers participating in torrent

torrent: group of peers exchanging chunks of a file

Alice arrives ...
... obtains list of peers from tracker
... and begins exchanging file chunks with peers in torrent



P2P File Distribution: BitTorrent

- **Peer joining torrent:**


- Has no chunks, but will accumulate them over time from other peers
- Registers with tracker to get list of peers, connects to subset of peers (“neighbors”)

- **While downloading, peer uploads chunks to other peers**

- **Peer may change peers with whom it exchanges chunks**

- **Churn:** peers may come and go

- **Once peer has entire file, it may (selfishly) leave or (altruistically) remain in torrent**



BitTorrent: Requesting/Sending File Chunks

- **Requesting chunks:**

- At any given time, different peers have different subsets of file chunks
- Periodically, Alice asks each peer for a list of chunks that they have
- Alice requests missing chunks from peers, rarest first


- **Sending chunks: tit-for-tat**

- Alice sends chunks to the four peers currently sending her chunks at highest rate
 - Other peers are choked by Alice (do not receive chunks from her)
 - Re-evaluate top 4 every 10 secs
- Every 30 secs, randomly select another peer, start sending chunks
 - Optimistically unchoke this peer
 - Newly chosen peer may join top four

- **Sending chunks:**

BitTorrent: tit-for-tat

- (1) Alice “optimistically unchokes” Bob
- (2) Alice becomes one of Bob’s top-four providers; Bob reciprocates
- (3) Bob becomes one of Alice’s top-four providers



Overview of Application Layer

- **Network Application Architectures**
- **HyperText Transfer Protocol (HTTP)**
- **File Transfer and Email protocols (FTP, SMTP)**
- **Domain Name System (DNS)**
- **Peer-to-Peer Applications (P2P)**
- **Video Streaming and Content Distribution Networks**
- **Socket Programming**

Video Streaming and CDNs: context


- video traffic: major consumer of Internet bandwidth
 - Netflix, YouTube: 37%, 16% of downstream residential ISP traffic
 - ~1B YouTube users, ~75M Netflix users
- challenge: scale - how to reach ~1B users?
 - single mega-video server won't work (why?)
 - challenge: heterogeneity
- different users have different capabilities (e.g., wired versus mobile; bandwidth rich versus bandwidth poor)
- **solution:** distributed, application-level infrastructure



Multimedia: video

- **video: sequence of images displayed at constant rate**
 - e.g., 24 images/sec
- **digital image: array of pixels**
 - each pixel represented by bits
- **coding: use redundancy *within* and *between* images to decrease # bits used to encode image**
 - spatial (within image)
 - temporal (from one image to next)

spatial coding example: instead of sending N values of same color (all purple), send only two values: color value (*purple*) and number of repeated values (N)



frame i

temporal coding example:
instead of sending complete frame at $i+1$,
send only differences from frame i




frame $i+1$

Multimedia: video

- **CBR: (constant bit rate):** video encoding rate fixed
- **VBR: (variable bit rate):** video encoding rate changes as amount of spatial, temporal coding changes
- **examples:**
 - MPEG I (CD-ROM) 1.5 Mbps
 - MPEG2 (DVD) 3-6 Mbps
 - MPEG4 (often used in Internet, < 1 Mbps)

spatial coding example: instead of sending N values of same color (all purple), send only two values: color value (*purple*) and number of repeated values (N)



frame i


temporal coding example: instead of sending complete frame at $i+1$, send only differences from frame i



frame $i+1$

Streaming stored video:

simple scenario:



Streaming multimedia: DASH

- **DASH: Dynamic, Adaptive Streaming over HTTP**
- **server:**
 - divides video file into multiple chunks
 - each chunk stored, encoded at different rates
 - *manifest file*: provides URLs for different chunks
- **client:**
 - periodically measures server-to-client bandwidth
 - consulting manifest, requests one chunk at a time
 - chooses maximum coding rate sustainable given current bandwidth
 - can choose different coding rates at different points in time (depending on available bandwidth at time)

Streaming multimedia: DASH

- **DASH: Dynamic, Adaptive Streaming over HTTP**
- “intelligence” at client: **client determines**
 - *when* to request chunk (so that buffer starvation, or overflow does not occur)
 - *what encoding rate* to request (higher quality when more bandwidth available)
 - *where* to request chunk (can request from URL server that is “close” to client or has high available bandwidth)

Content distribution networks

- **challenge:** how to stream content (selected from millions of videos) to hundreds of thousands of *simultaneous* users?
- **option 1:** single, large “mega-server”
 - single point of failure
 - point of network congestion
 - long path to distant clients
 - multiple copies of video sent over outgoing link


....quite simply: this solution ***doesn't scale***

Content distribution networks

- **challenge:** how to stream content (selected from millions of videos) to hundreds of thousands of simultaneous users?
- **option 2:** store/serve multiple copies of videos at multiple geographically distributed sites (**CDN**)
 - **enter deep:** push CDN servers deep into many access networks
 - close to users
 - used by Akamai, 1700 locations
 - **bring home:** smaller number (10's) of larger clusters in POPs near (but not within) access networks
 - used by Limelight

Content Distribution Networks (CDNs)

- CDN: stores copies of content at CDN nodes
 - e.g. Netflix stores copies of MadMen
- subscriber requests content from CDN
 - directed to nearby copy, retrieves content
 - may choose different copy if network path congested



Content Distribution Networks (CDNs)



OTT challenges: coping with a congested Internet

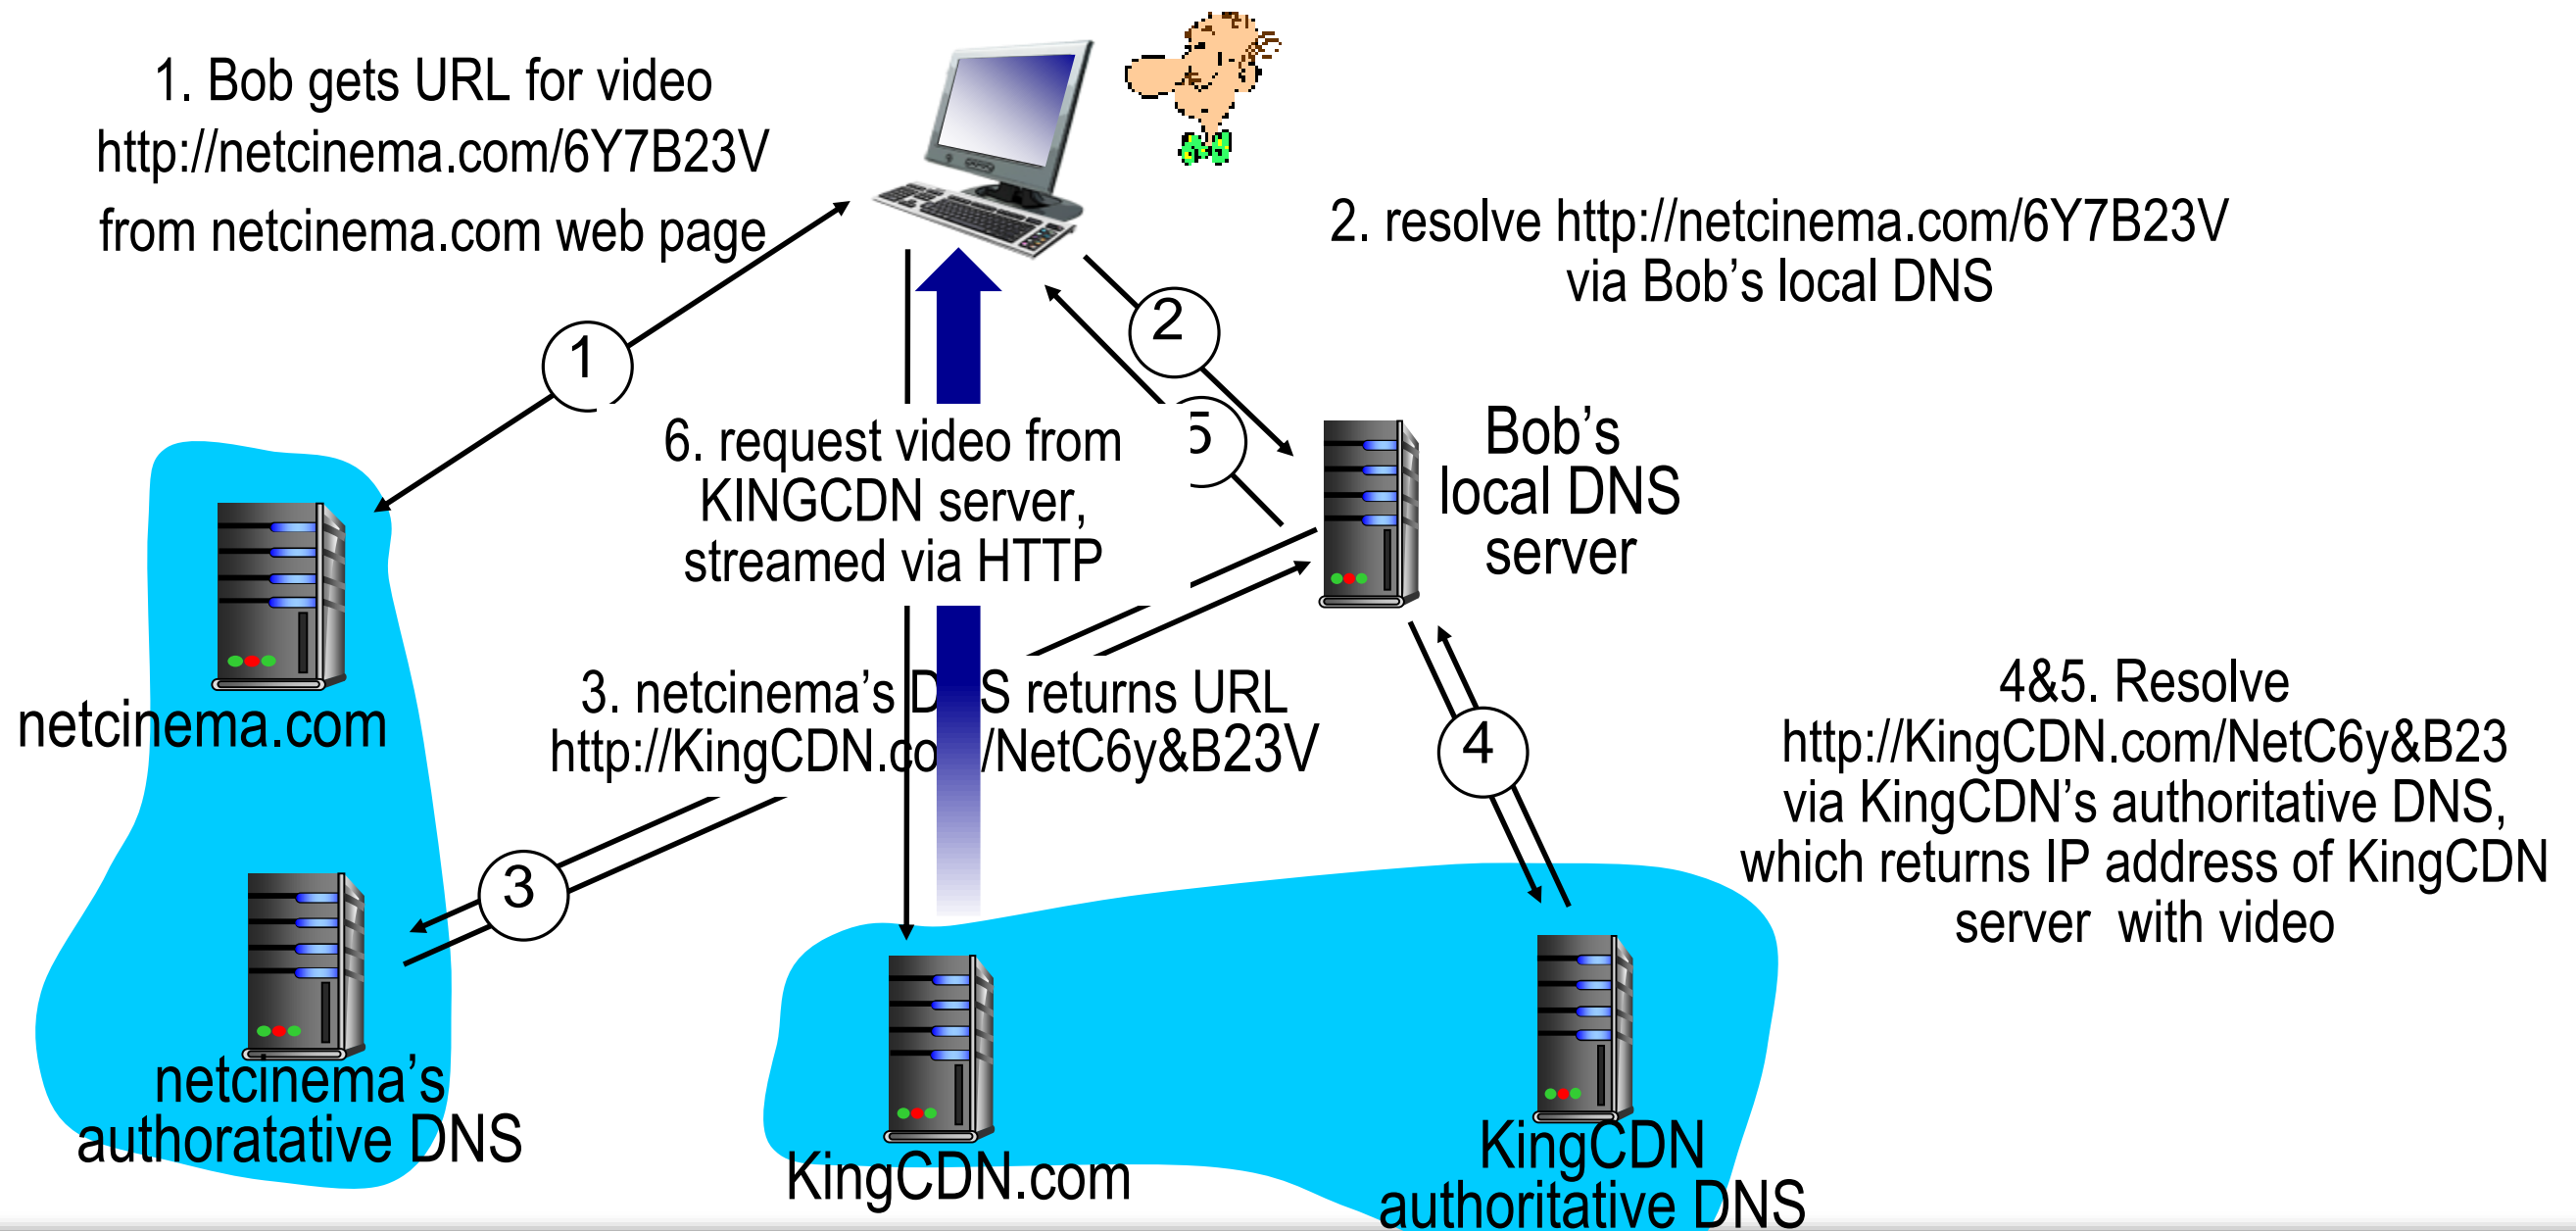
- from which CDN node to retrieve content?
- viewer behavior in presence of congestion?
- what content to place in which CDN node?

more .. in chapter 7

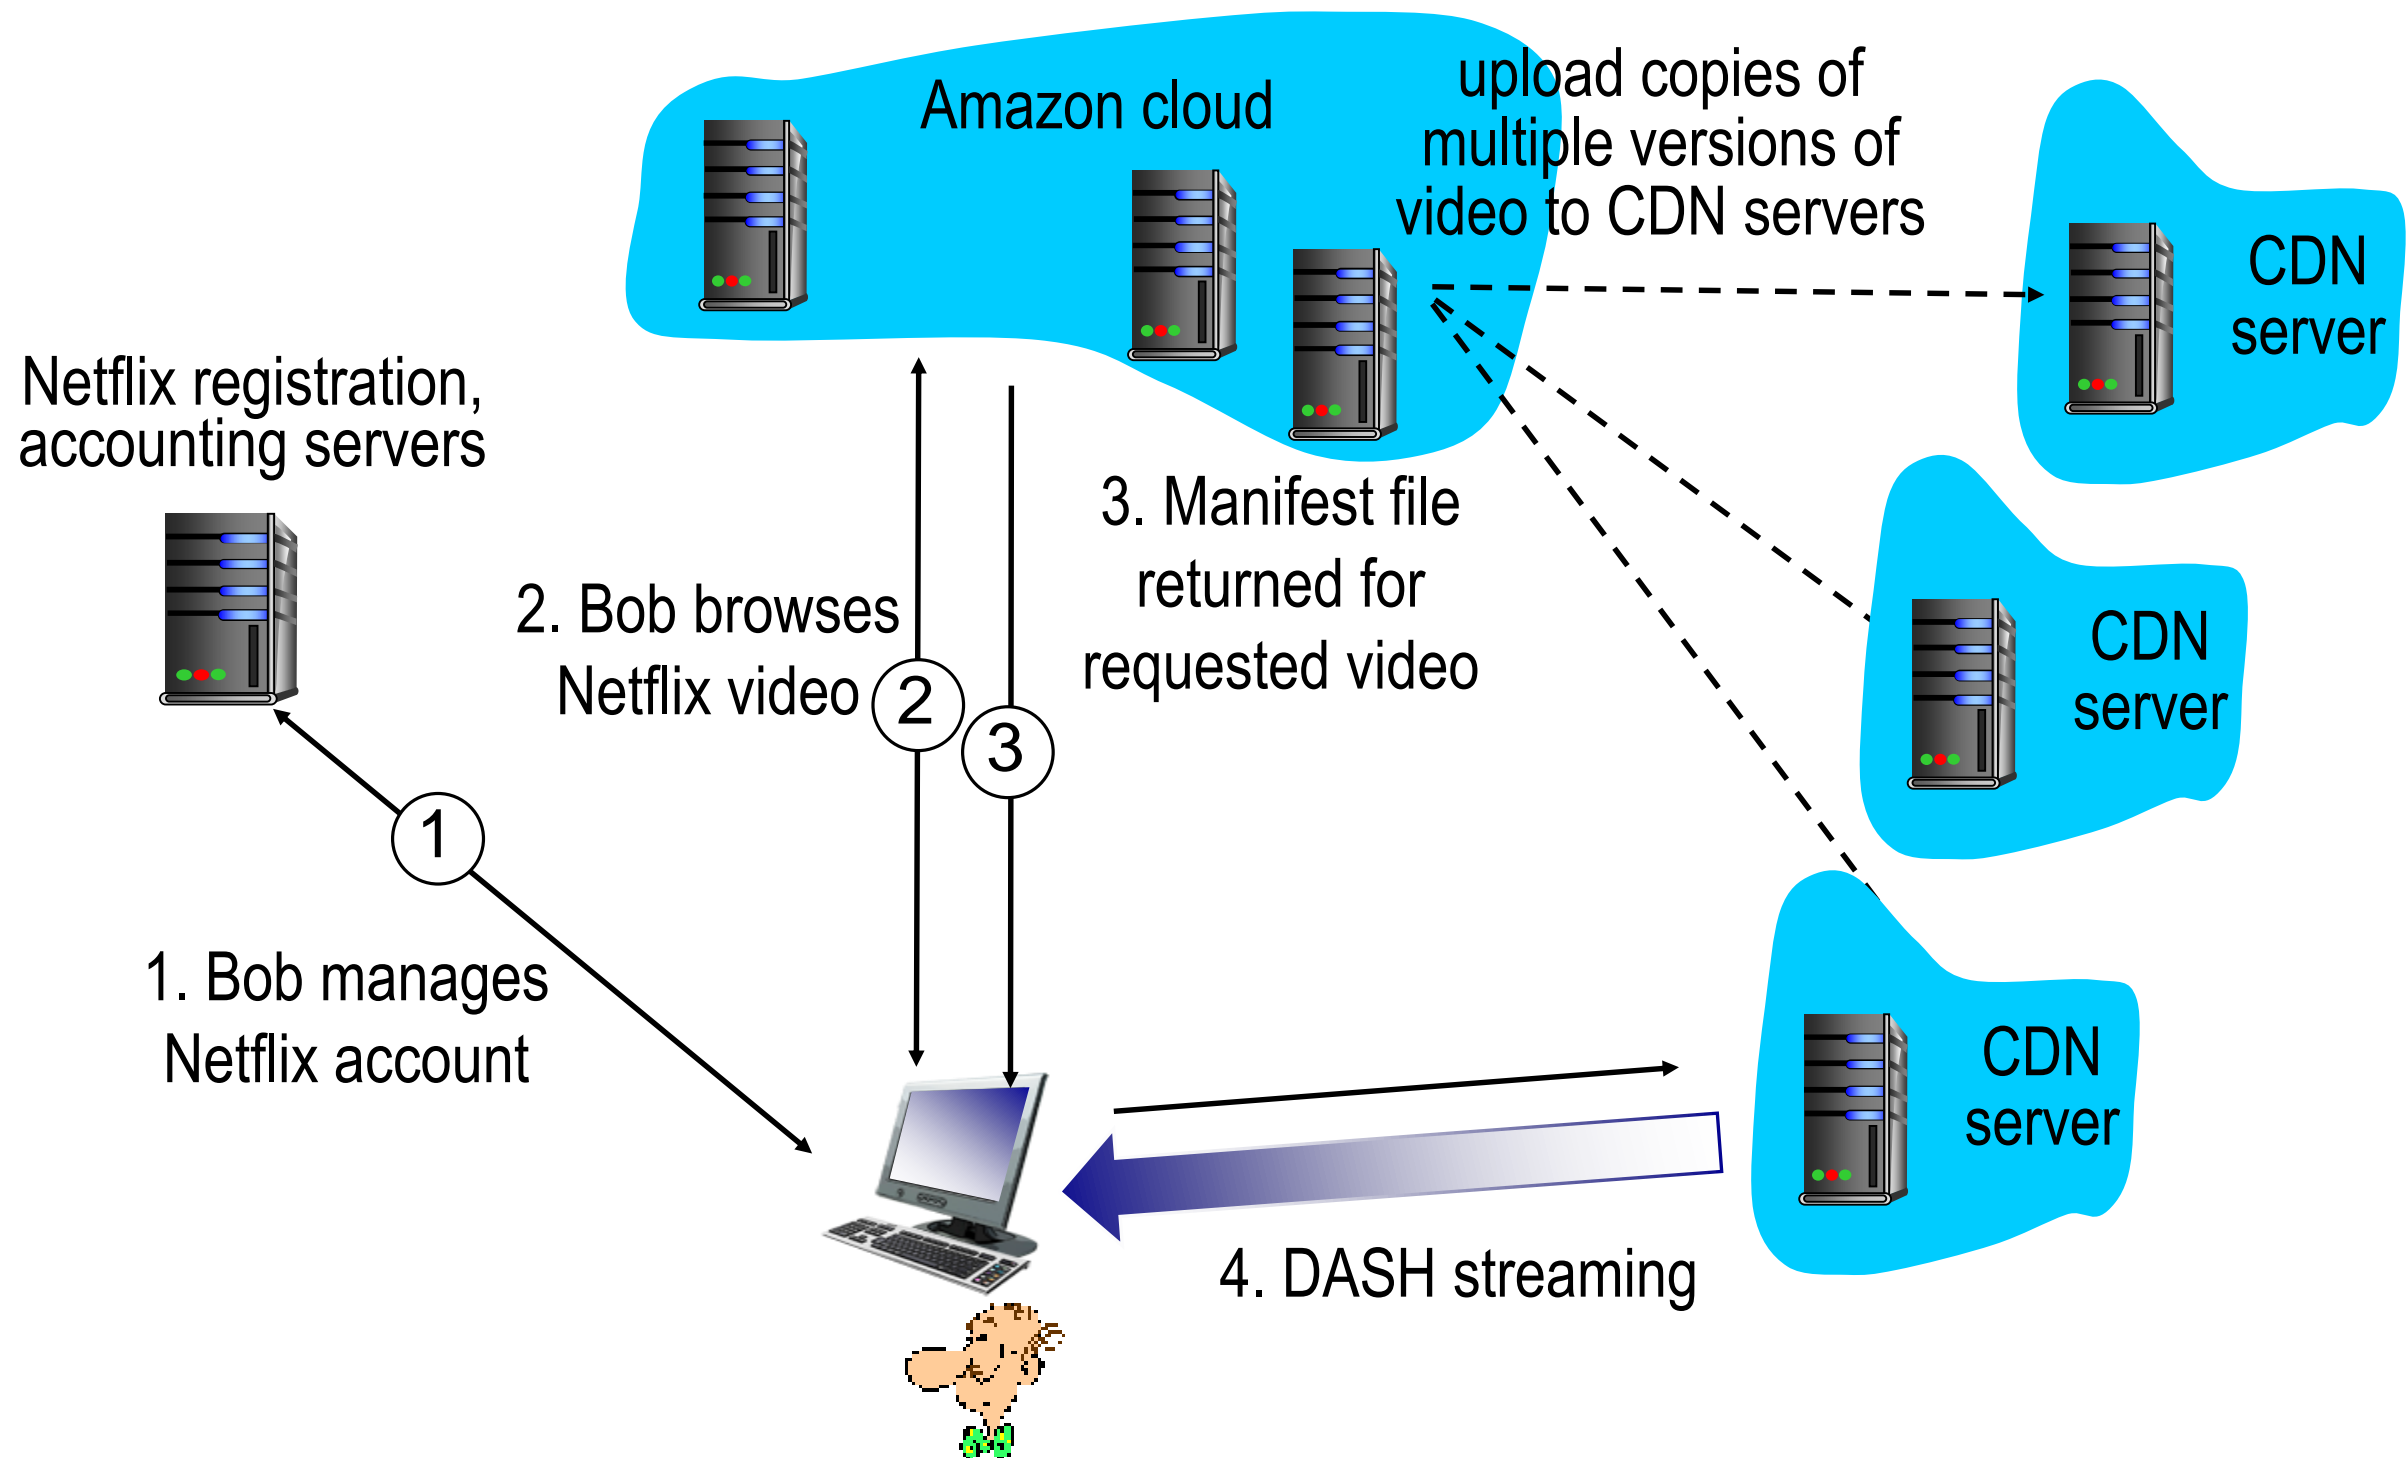
CDN content access: a closer look

Bob (client) requests video `http://netcinema.com/6Y7B23V`

- video stored in CDN at `http://KingCDN.com/NetC6y&B23V`



Case study: Netflix

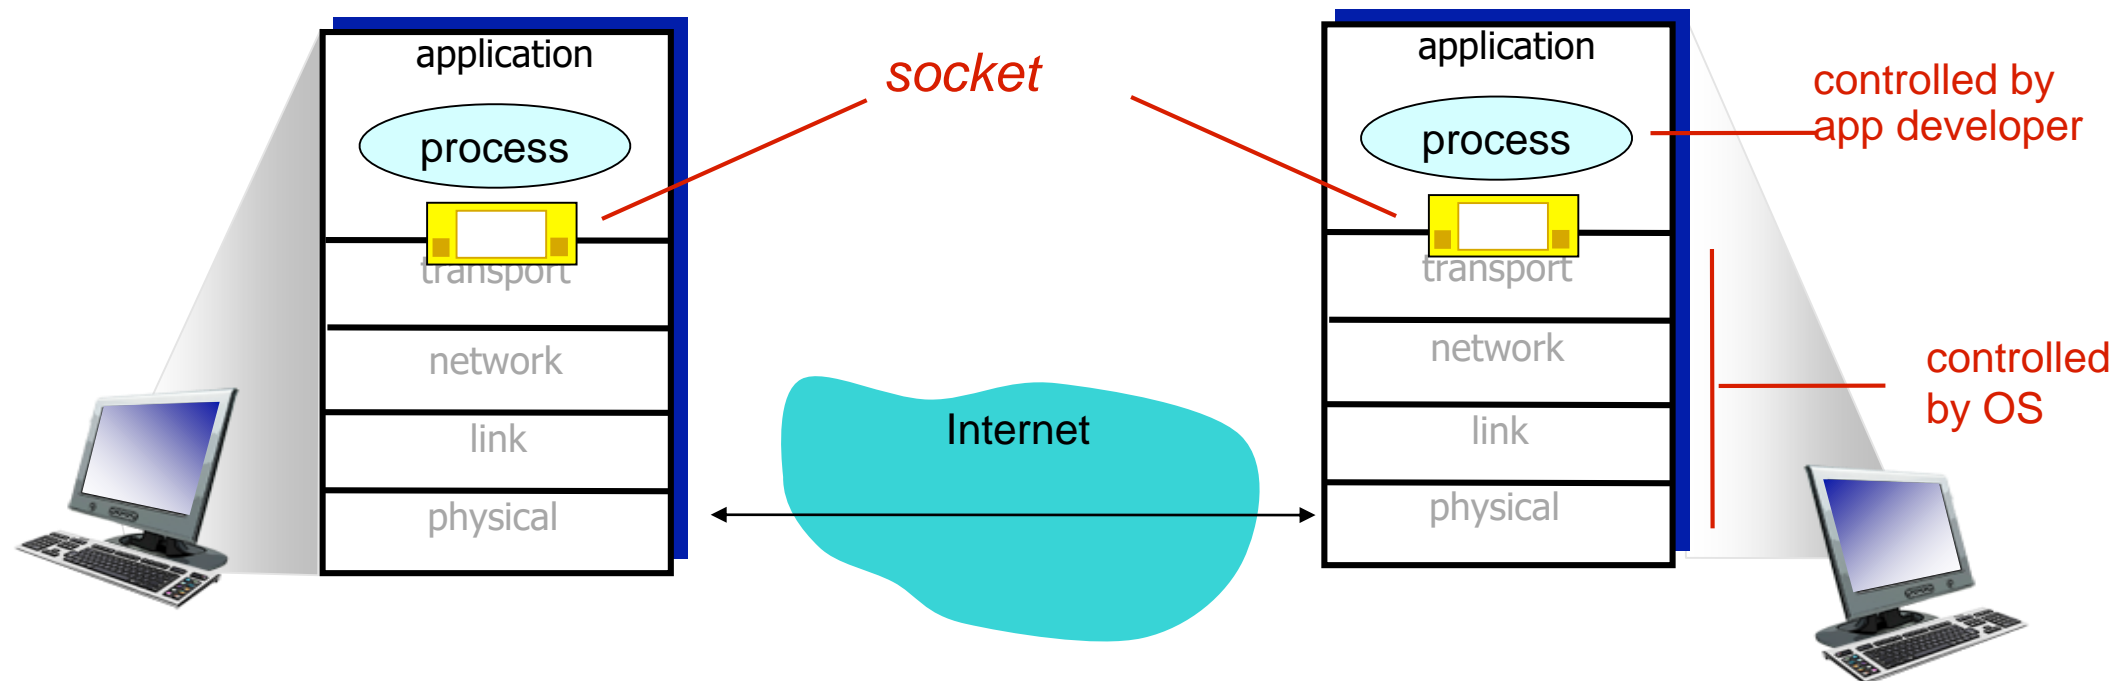


Overview of Application Layer

- **Network Application Architectures**
- **HyperText Transfer Protocol (HTTP)**
- **File Transfer and Email protocols (FTP, SMTP)**
- **Domain Name System (DNS)**
- **Peer-to-Peer Applications (P2P)**
- **Video Streaming and Content Distribution Networks**
- **Socket Programming**

Socket Programming

- **Goal:** learn how to build client/server applications that communicate using sockets
- **Socket:** door between application process and end-end-transport protocol



Socket Programming

- **Two socket types for two transport services:**
 - **UDP**: unreliable datagram
 - **TCP**: reliable, byte stream-oriented
- **Application Example:**
 - (1) Client reads a line of characters (data) from its keyboard and sends the data to a server.
 - (2) The server receives the data and converts characters to uppercase.
 - (3) The server sends the modified data to the client.
 - (4) The client receives the modified data and displays the line on its screen.

Socket Programming with UDP

- **UDP: no “connection” between client & server**
 - No handshaking before sending data
 - Sender explicitly attaches IP destination address and port number to each packet
 - Receiver extracts sender IP address and port number from received packet
- **Transmitted data may be lost or received out-of-order**
- **Application viewpoint:**
 - UDP provides unreliable transfer of datagrams between client and server

Socket Programming with TCP

- **Client must contact server**
 - Server process must first be running
 - Server must have created socket that welcomes client's connection
- **Client contacts server by:**
 - Creating TCP socket, specifying IP address, port number of server process
 - When client creates socket, client TCP establishes connection to server TCP
- **When contacted by client, server TCP creates new socket for server process to communicate with that particular client**
 - Allows server to talk with multiple clients
 - Source port numbers used to distinguish clients
- **Application viewpoint:**
 - TCP provides reliable, in-order byte-stream transfer (a "pipe") between client and server

Chapter 2: summary

most importantly: learned about protocols!

- **typical request/reply message exchange:**

- client requests info or service
- server responds with data, status code

- **message formats:**

- *headers*: fields giving info about data
- *data*: info(payload) being communicated

important themes:

- control vs. messages
 - in-band, out-of-band
- centralized vs. decentralized
- stateless vs. stateful
- reliable vs. unreliable message transfer
- “complexity at network edge”