

CS 330: Network Applications & Protocols

Transport Layer

Galin Zhelezov
Department of Physical Sciences
York College of Pennsylvania

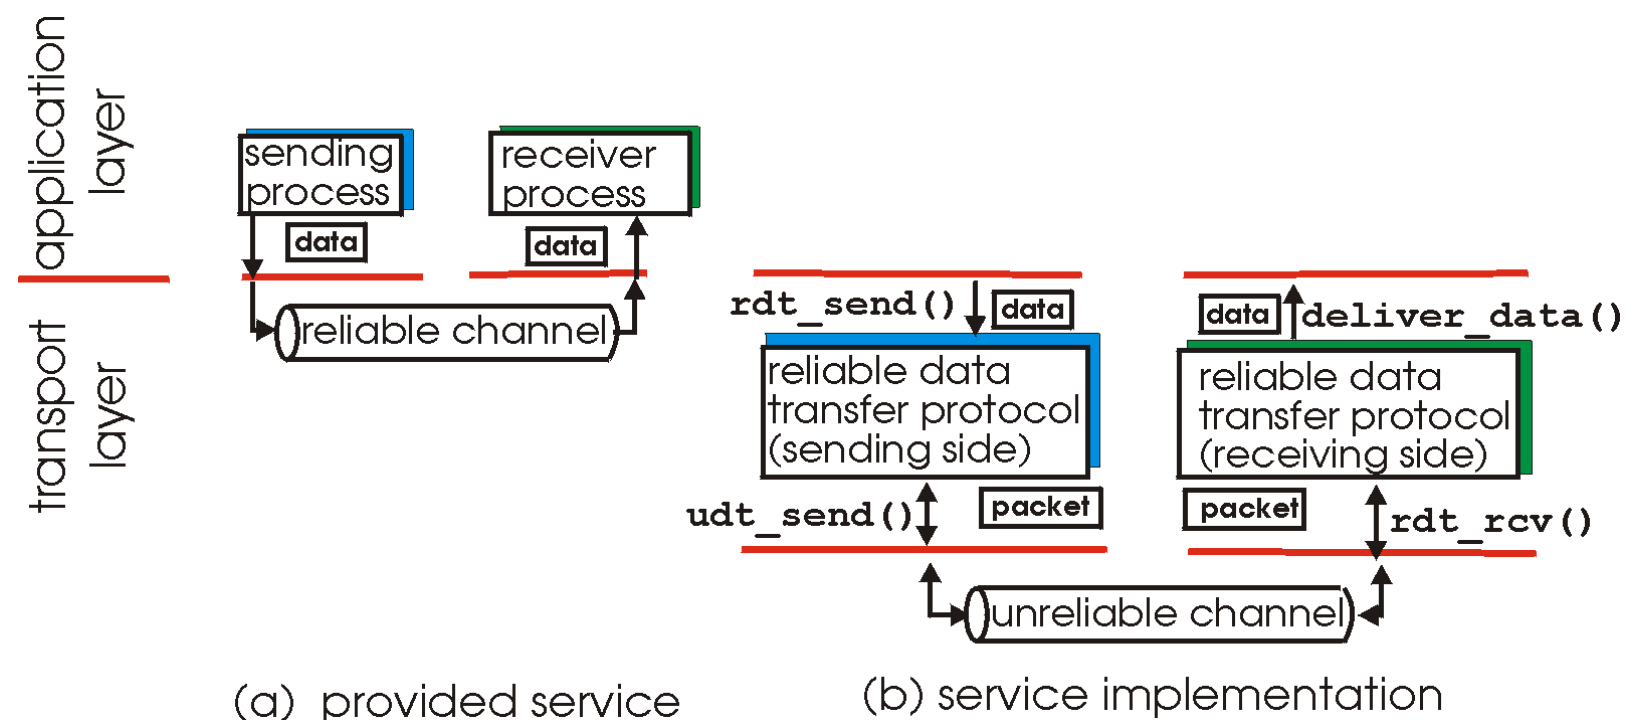


Overview of Transport Layer

- **Transport-layer Services**
- **Multiplexing and Demultiplexing**
- **Connectionless Transport: UDP**
- **Principles of Reliable Data Transfer**
 - Overview
 - Pipelined Protocols
 - Go-Back-N
 - Selective Repeat
- **Connection-oriented Transport: TCP**
- **Principles of Congestion Control**
- **TCP Congestion Control**

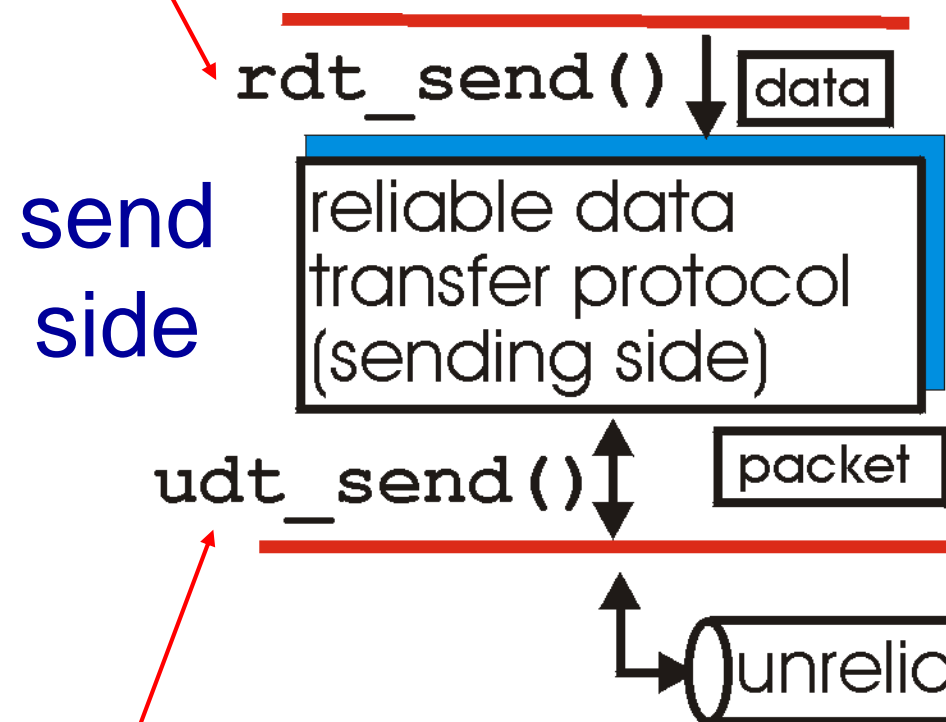
Principles of Reliable Data Transfer

- Reliable data transfer is very important application, transport, and link layers
- The characteristics of unreliable channel will determine the complexity of a **Reliable Data Transfer (RDT) protocol**
- To explore reliable data transfer, examine different types of loss and how to address them

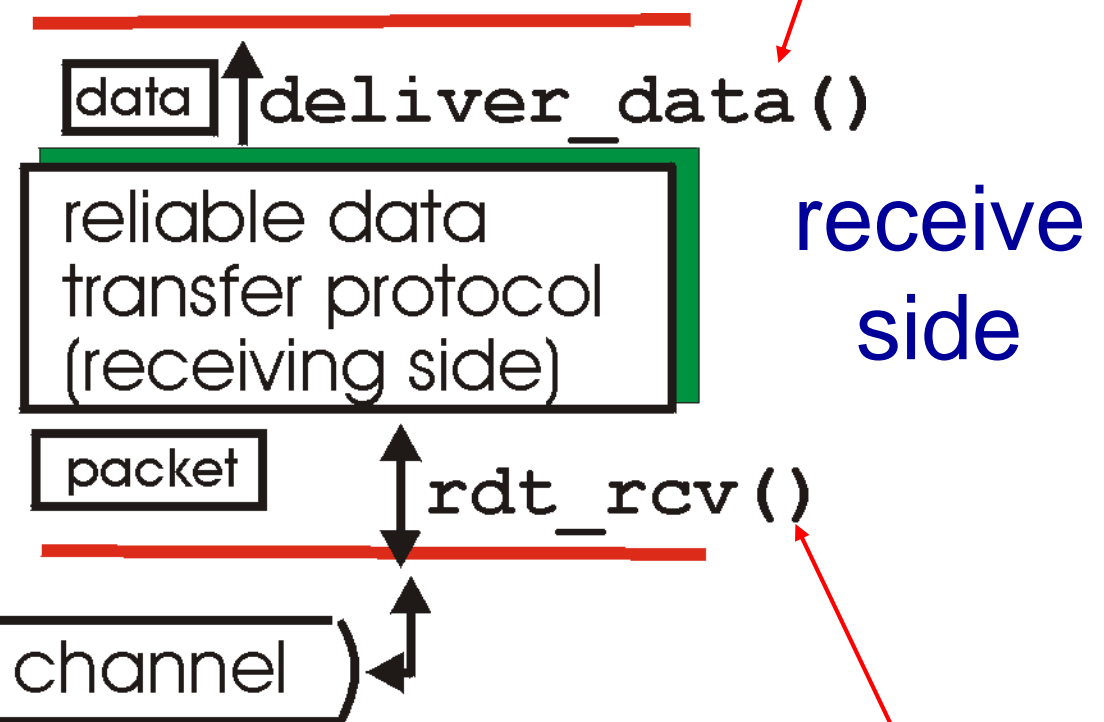


Reliable data transfer: getting started

rdt_send() : called from above, (e.g., by app.). Passed data to deliver to receiver upper layer



deliver_data() : called by **rdt** to deliver data to upper



udt_send() : called by rdt, to transfer packet over unreliable channel to receiver

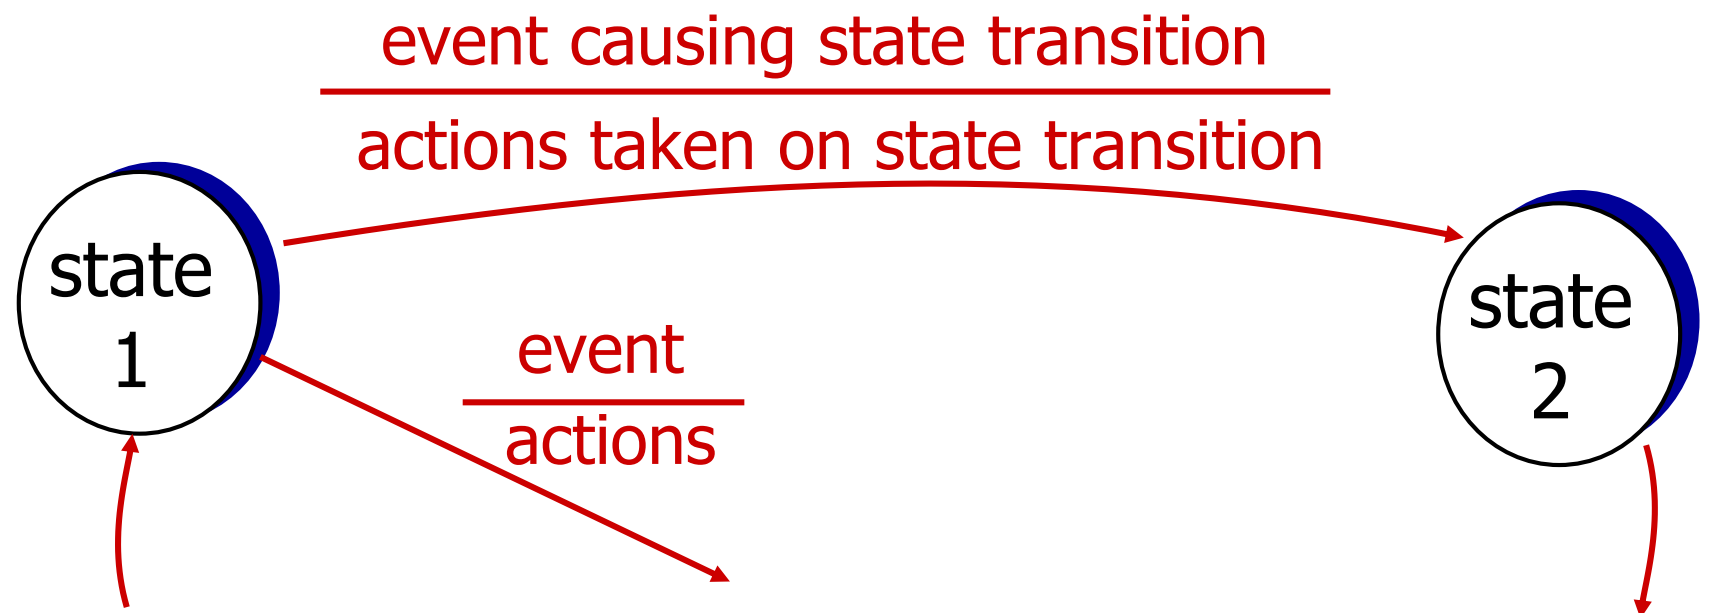
rdt_rcv() : called when packet arrives on rcv-side of channel

Reliable data transfer: getting started

we'll:

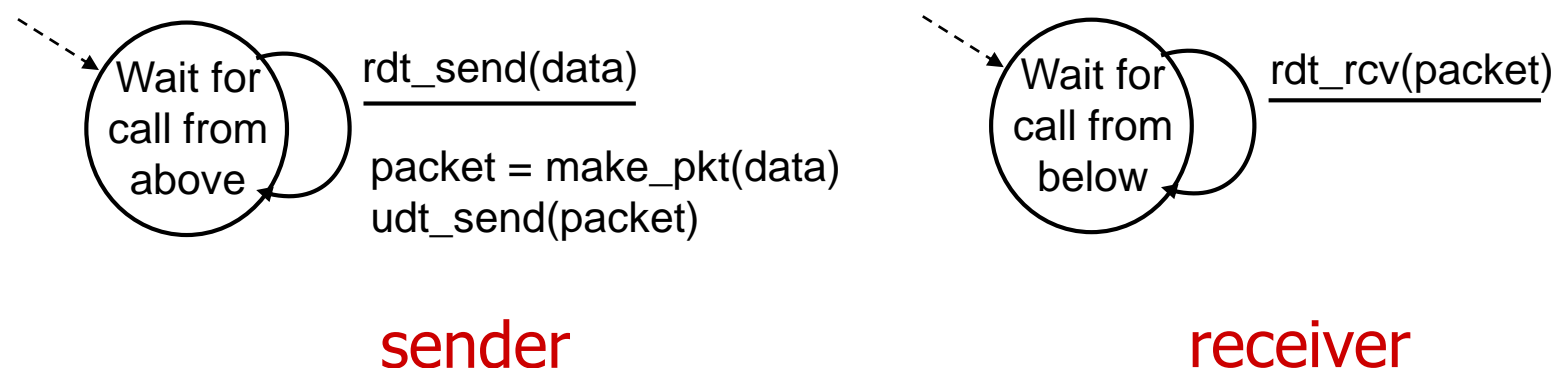
- incrementally develop sender, receiver sides of **reliable data transfer protocol (rdt)**
- **consider only unidirectional data transfer**
 - but control info will flow on both directions!
- **use finite state machines (FSM) to specify sender, receiver**

state: when in this “state” next state uniquely determined by next event



RDT 1.0: Reliable Transfer Over a Reliable Channel

- **Underlying data transmission channel is perfectly reliable**
 - No bit errors
 - No loss of packets
- **We don't have networks like this, but it's a good place to start**
- **separate FSMs for sender, receiver:**
 - sender sends data into underlying channel
 - receiver reads data from underlying channel



RDT 2.0: Data Channel With Bit Errors

- **underlying channel may flip bits in packet**
 - checksum to detect bit errors
- ***the question: how to recover from errors:***

*How do humans recover from “errors”
during conversation?*

RDT 2.0: Data Channel With Bit Errors

- **Underlying data channel may flip bits in packet**
 - Use a checksum to detect bit errors
- **Question: How should system recover from these errors?**
 - **Acknowledgements (ACKs)**: receiver explicitly tells sender that a packet is received OK
 - **Negative acknowledgements (NAKs)**: receiver explicitly tells sender that a packet had errors when it was received
 - **Sender retransmits packet on receipt of a NAK**
- **New mechanisms in RDT 2.0 (beyond RDT 1.0):**
 - Error detection
 - Feedback: send control messages (ACK, NAK) from receiver to sender

A Fatal Flaw in RDT 2.0, on to RDT 2.1

- **What happens if ACK/NAK messages are corrupted?**
 - Sender doesn't know if receiver correctly received the data
 - Data may have been corrupt on the way to receiver
 - ACK/NAK may have been corrupt on the way back to sender
 - Can't just retransmit since receiver may receive duplicate data!
- **Handling duplicates:**
 - Sender retransmits current packet if ACK/NAK is corrupted
 - Sender adds a **sequence number** to each packet
 - Receiver discards duplicate packets at Transport Layer
 - Those packets are not delivered up to the Application Layer

RDT 2.1: Discussion

- **Sender:**

- Sequence number added to packets
- Two sequence numbers (0,1) will suffice
- Must check if received ACK/NAK corrupted
- Sender must “remember” if it should be expecting a sequence number of 0 or 1

- **Receiver:**

- Must check if received packet is duplicate
 - Receiver must “remember” if it should be expecting a sequence number of 0 or 1
- Note: the receiver does not know if its last ACK/NAK message was received OK at sender

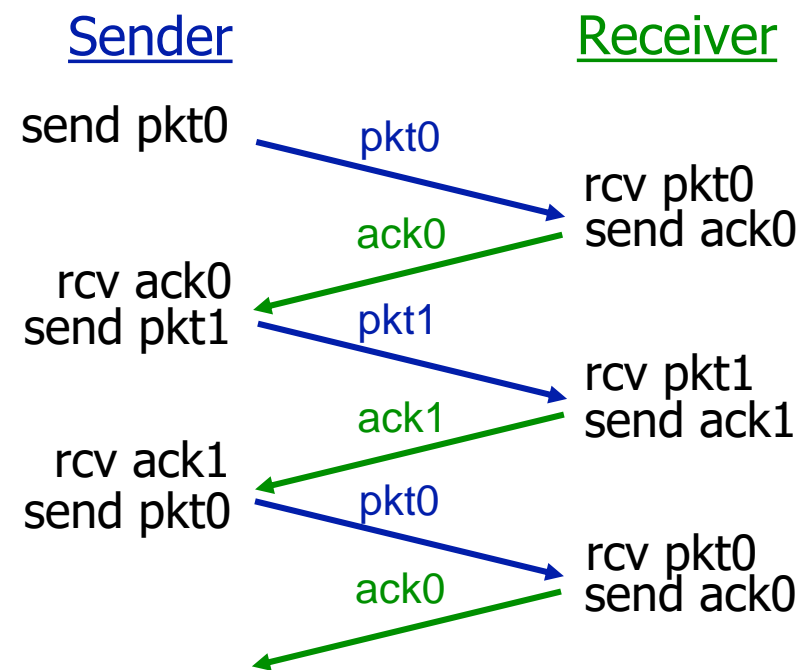
RDT 2.2: Eliminating the NAK Messages

- **Possible to achieve same functionality as RDT 2.1 using ACKs messages only**
- **Receiver sends ACK messages for last packet that was received correctly**
 - No message is sent for a packet this is received with errors
- **Duplicate ACK messages received at sender results in same action as NAK from RDT 2.1 -- retransmit packet**
 - Duplicate ACK message would be detected at sender by receiving two consecutive ACK_0 messages or two consecutive ACK_0 messages


RDT 3.0: Channels with Bit Errors & Packet Loss

- **New assumption: underlying channel can also lose packets**
 - Can lose data packets or ACK messages
 - Checksum, sequence number, ACKs, retransmissions will be of help ... but not enough
- **Approach: sender waits a “reasonable” amount of time for an ACK**
 - Retransmits packet if no ACK is received in this time
 - If packet (or ACK) is just delayed and not lost:
 - Retransmission will result in duplicate data, but sequence numbers from RDT 2.2. already handle this issue
 - Receiver must specify sequence number of packet being ACKed
 - Requires countdown timer

RDT 3.0 in Action with Packet Loss

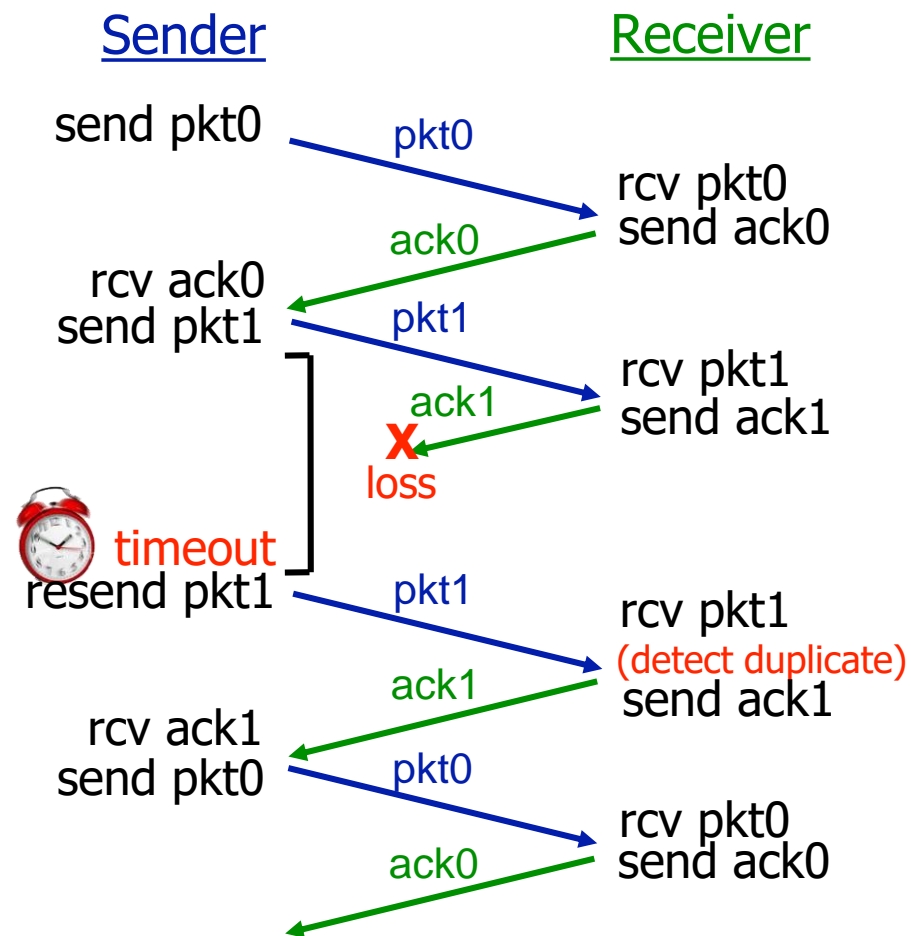


No Packet Loss




With Packet Loss

RDT 3.0 in Action with Lost/Delayed ACK



ACK Loss



Premature Timeout/Delayed ACK

Performance of RDT 3.0

- **RDT 3.0 will work reliably, but would have terrible performance**
 - RDT 3.0 utilizes a **Stop-and-Wait protocol**
 - Sender sends one packet, then waits for receiver response before sending the next packet
- **Example:** Assume a system sending 8000 bit packets on a 1 Gbps link, where there is a 15 ms propagation delay between the sender and the receiver

$$\text{Time for sender to transmit data} \quad D_{\text{trans}} = \frac{L}{R} = \frac{8000 \text{ bits}}{10^9 \text{ bits/sec}} = 8 \text{ microsecs}$$

Performance of RDT 3.0 (Cont.)

- **If Round-Trip Time (RTT) is 30 ms:**

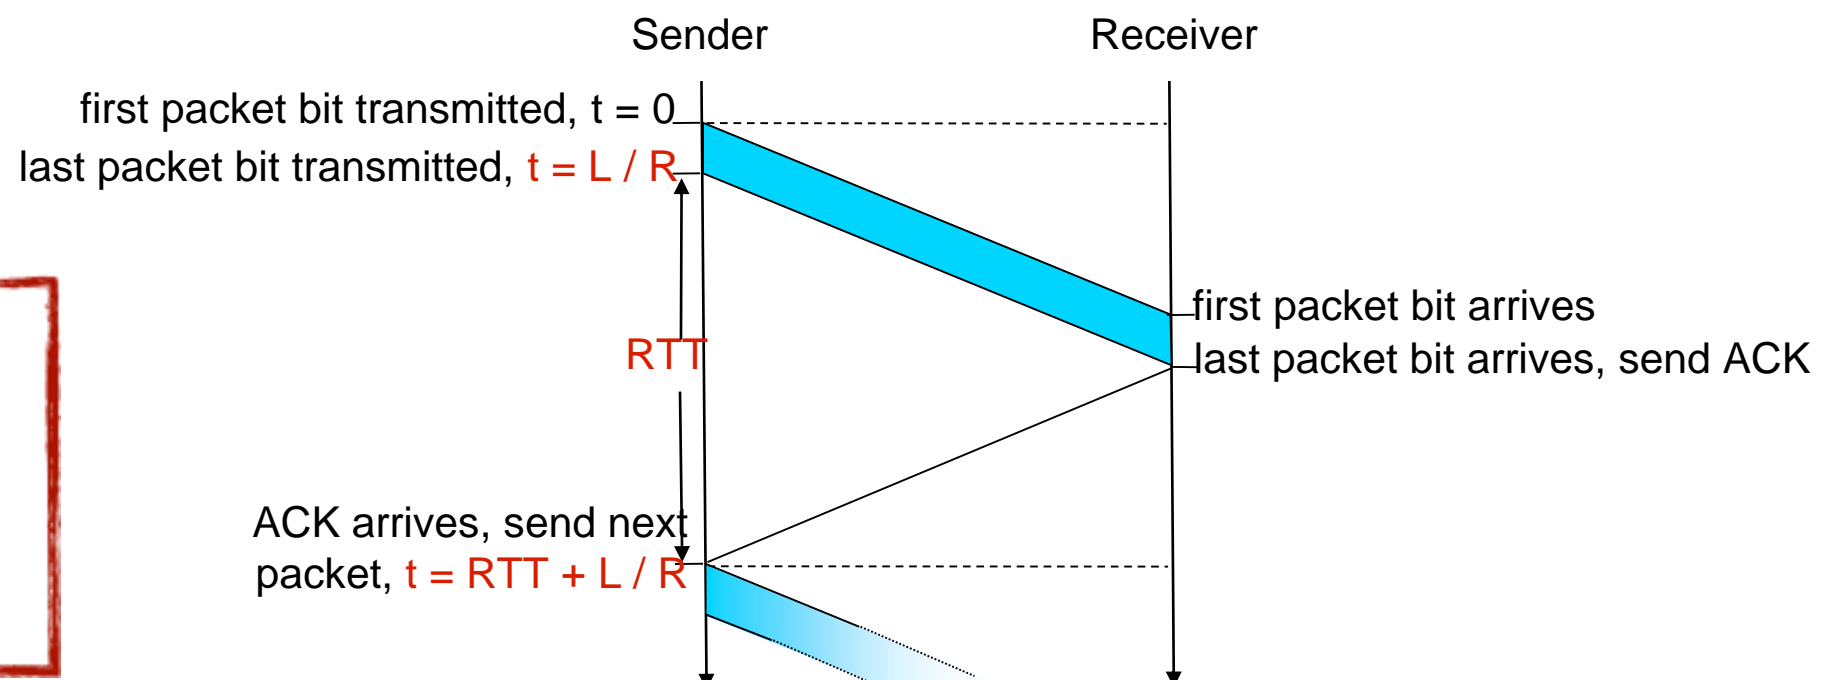
- Sender is only sending 8 microseconds
- Can only send a new packet every 30.008 microseconds
- Effectively makes 1 Gbps link run at ~270 Kbps!!!

Utilization: Fraction
of time the sender
is busy sending

$$U_{\text{sender}} = \frac{L/R}{RTT + L/R} = \frac{.008}{30.008} = 0.00027$$


.027%

The network protocol
limits use of physical
resources!




Pipelined Protocols

- **Pipelining:** sender allows multiple, “in-flight”, yet-to-be-acknowledged packets
 - Range of sequence numbers must be increased (0 and 1 will no longer suffice)
 - Buffering at sender and/or receiver is required
- **Two generic forms of pipelined protocols: Go-Back-N, and Selective Repeat**

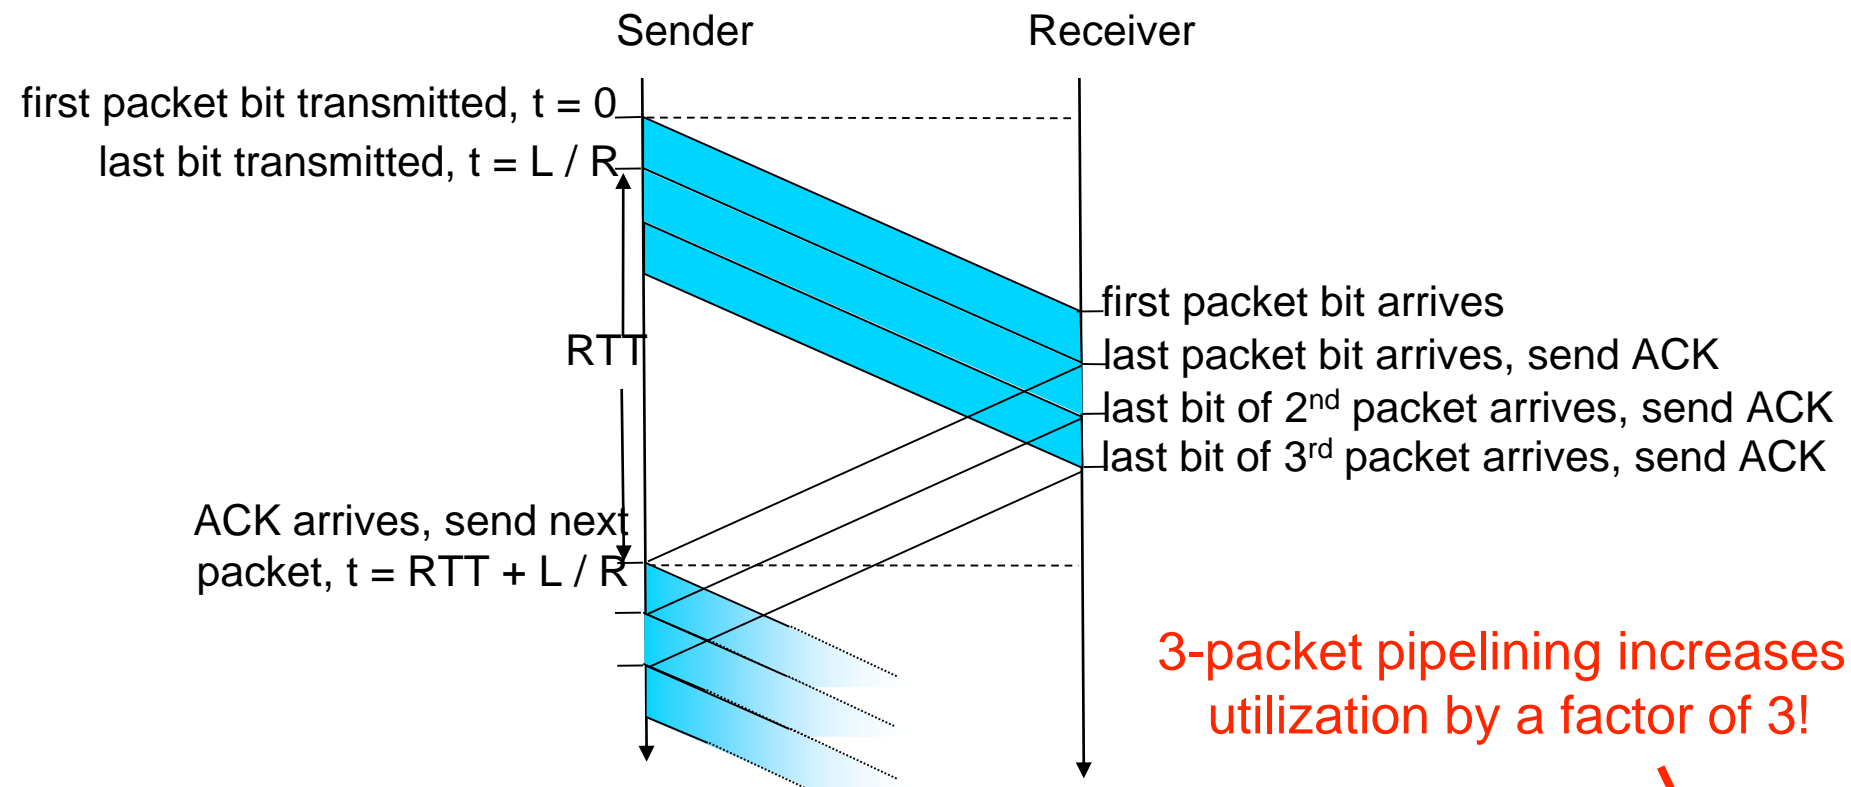


Stop-and-Wait Protocol



Go-Back-N Protocol

Pipelining: Increased Utilization



Utilization: Fraction of time the sender is busy sending

$$U_{\text{sender}} = \frac{3L/R}{RTT + L/R} = \frac{.024}{30.008} = 0.0008$$

Pipelined Protocols: Overview

- **Go-back-N:**


- Sender can have up to N unacked packets in pipeline
- Receiver only sends cumulative acks
 - Doesn't ack a packet if there is a gap
- Sender has a timer for the oldest unacked packet
 - When timer expires, retransmit **all** unacked packets

- **Selective Repeat:**

- Sender can have up to N unacked packets in pipeline
- Receiver sends individual acks for each packet
- Sender maintains timer for each unacked packet
 - When timer expires, retransmit only that unacked packet

Go-Back-N: Sender Side

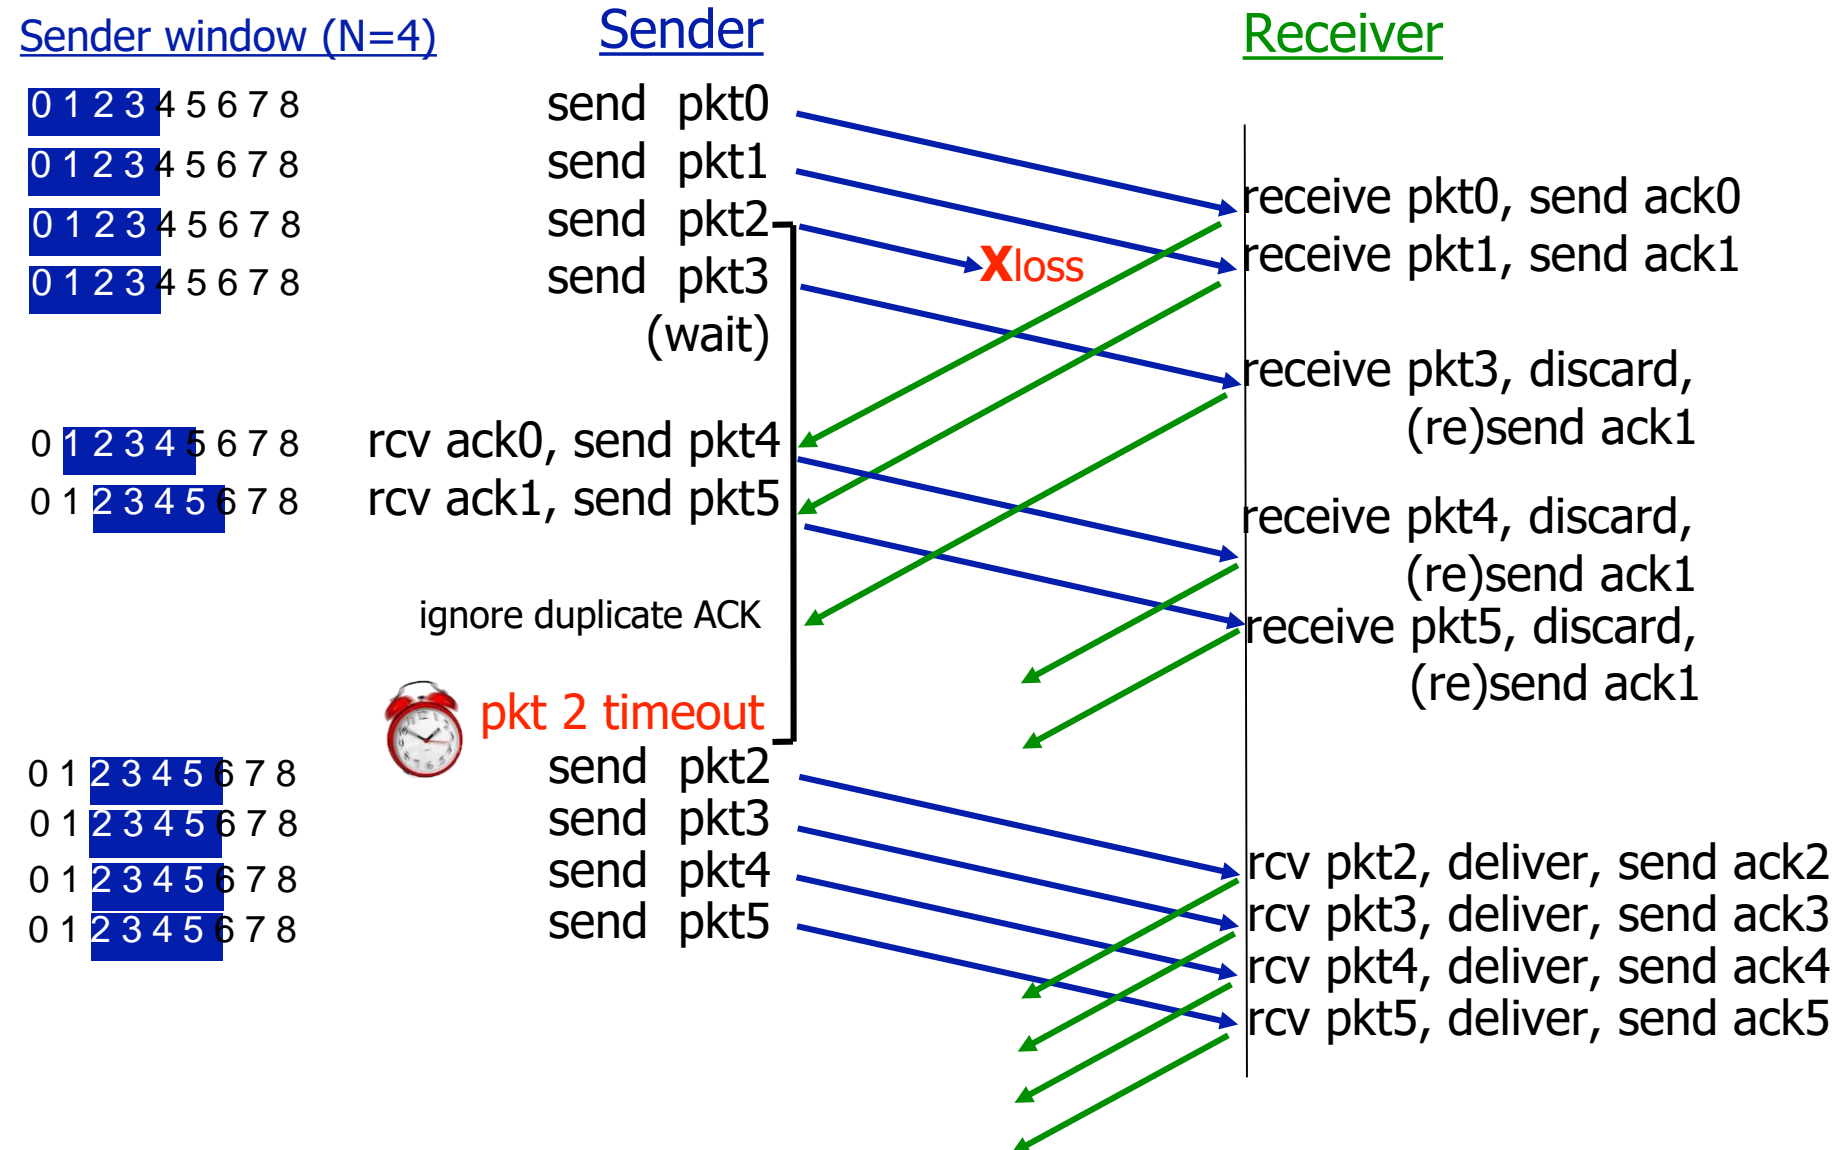
- **Sender can transmit multiple packets without waiting for an ack**
 - A “sliding window” of up to N consecutive unacked packets allowed
- **Include a k-bit sequence number in the packet header**
- **ACK(n): ACKs all packets up to and including sequence number n**
 - Cumulative ACK
 - May receive duplicate ACKs
- **Maintain timer for oldest in-flight packet**
- **Timeout(n): retransmit packet n and all higher sequence number packets in window**



Go-Back-N: Receiver

- **Send ACK message for correctly-received packet with highest in-order sequence number**
 - May generate duplicate ACKs
 - Need only remember expected sequence number
 - Must receive packets in-order to send ACK
- **Out-of-order packets:**
 - **Discard** (don't buffer) out-of-order packets (no receiver buffering)
 - Yes, even if they are correctly formatted and error-free
 - Will be retransmitted anyway based on sender's rules
 - Re-ACK packet with highest in-order sequence number

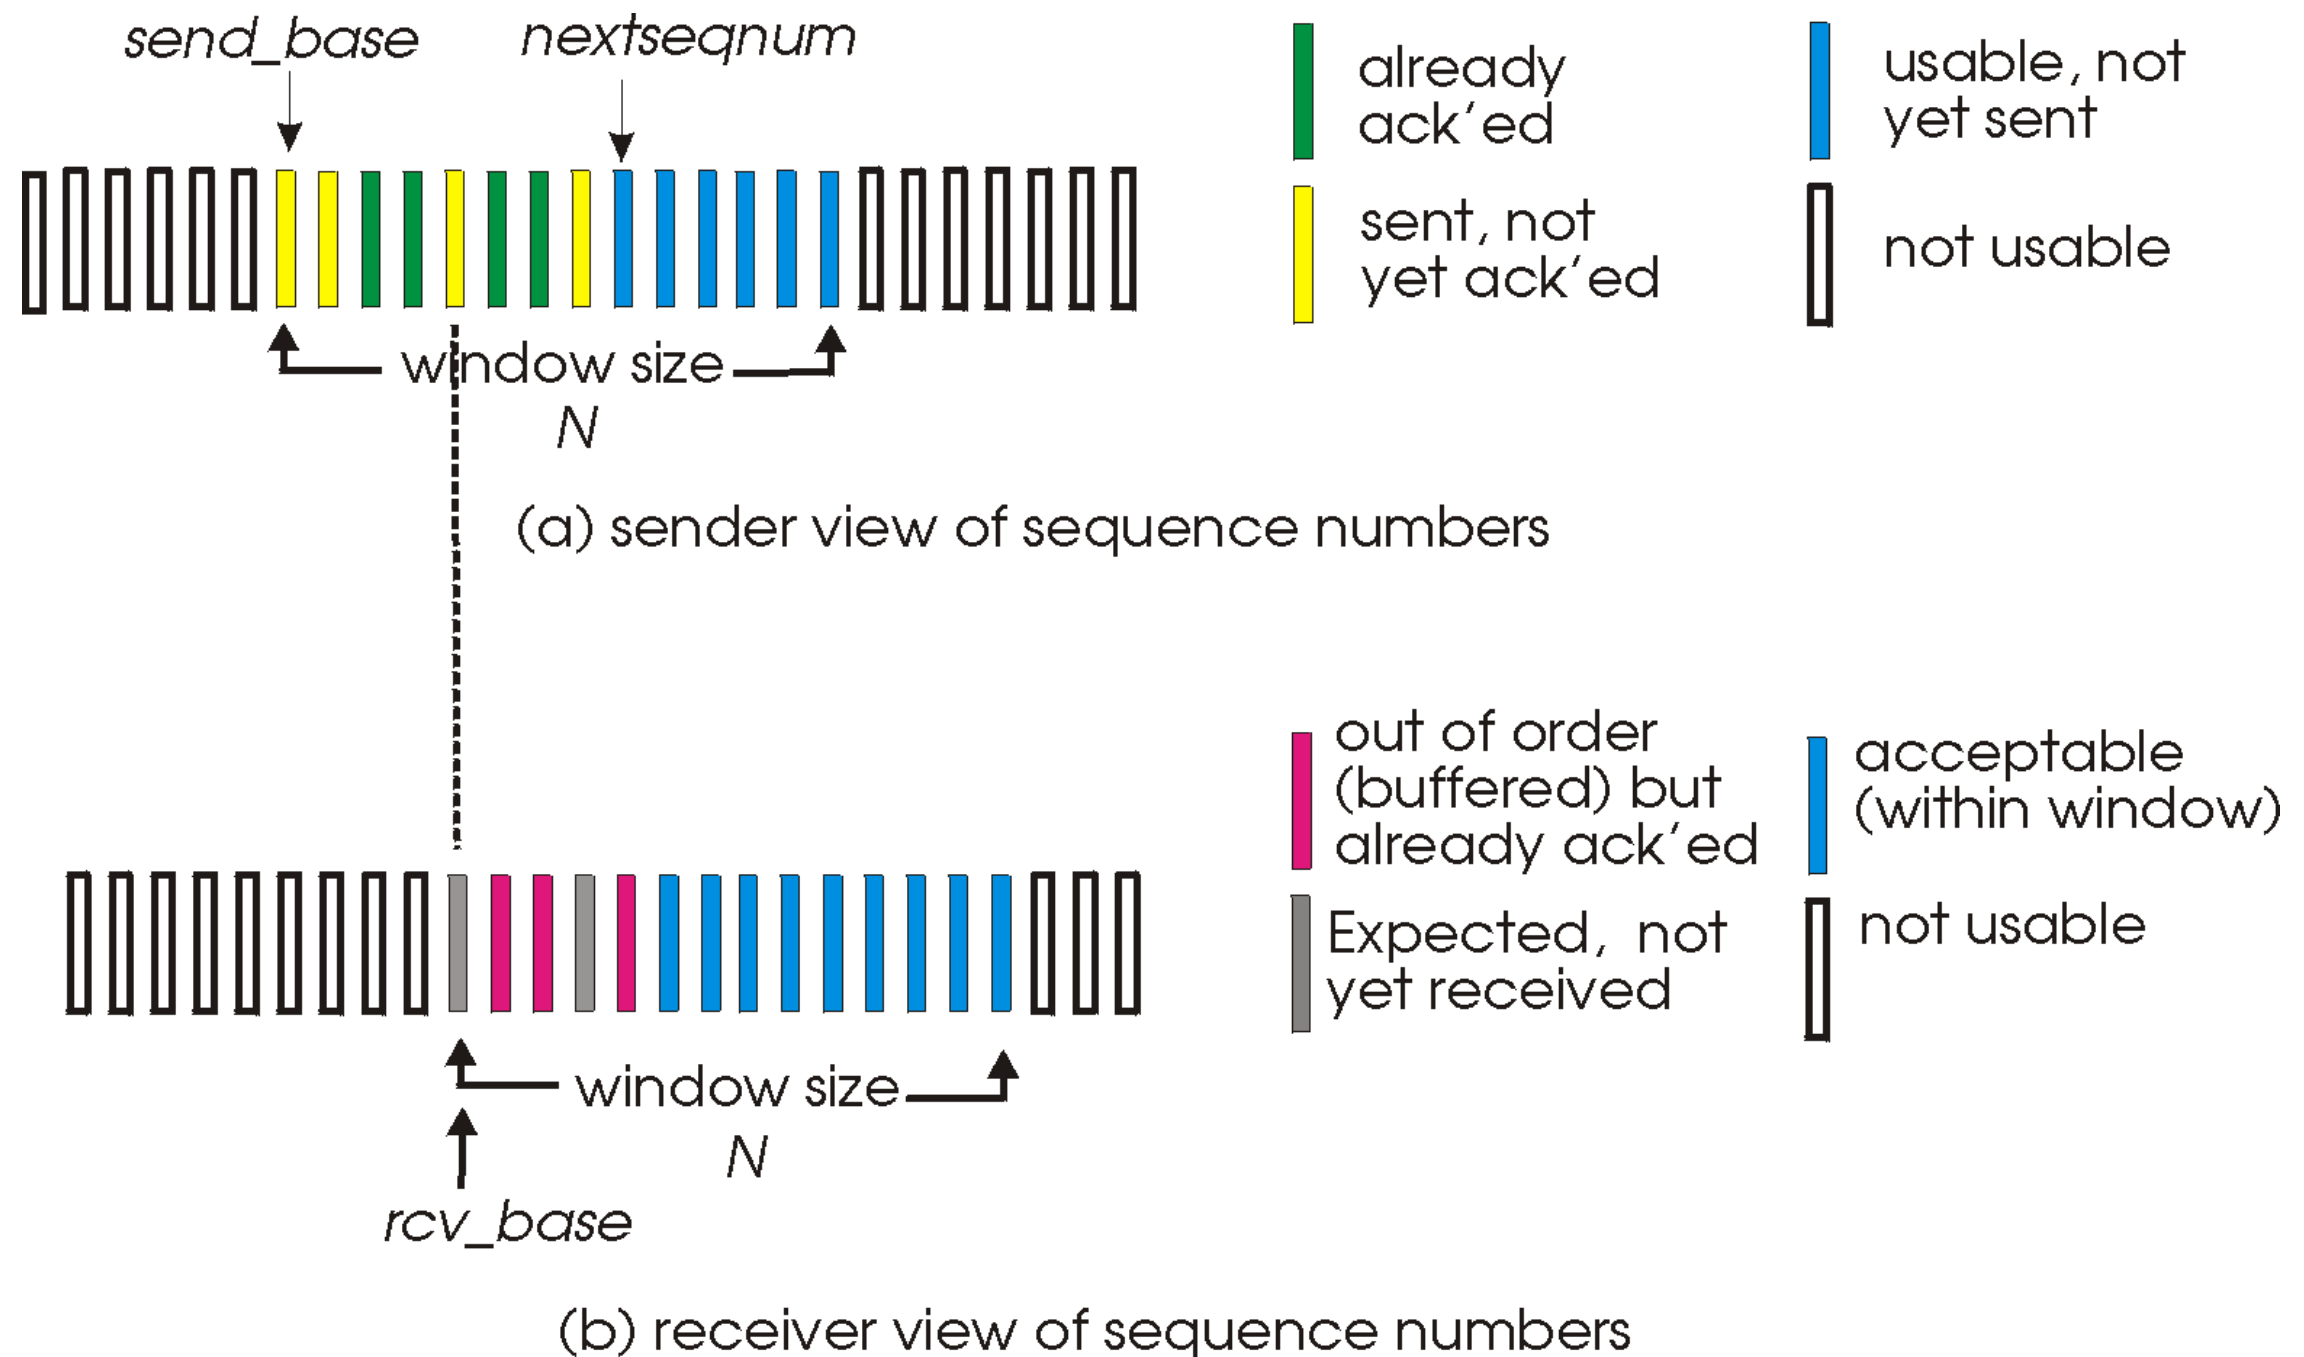
Go-Back-N in Action



Selective Repeat

- **Receiver individually acknowledges all correctly received packets**
 - Out-of-order packets are buffered at the receiver
 - Buffered packets are eventually delivered in-order to upper layer
- **Avoids unnecessary retransmission -- sender only resends packets for which an ACK was not received**
 - Sender has separate timer for each unACKed packet
- **Sender window**
 - N consecutive sequence numbers
 - Limits sequence numbers of sent, unACKed packets

Selective Repeat: Sender/Receiver Windows



Selective Repeat

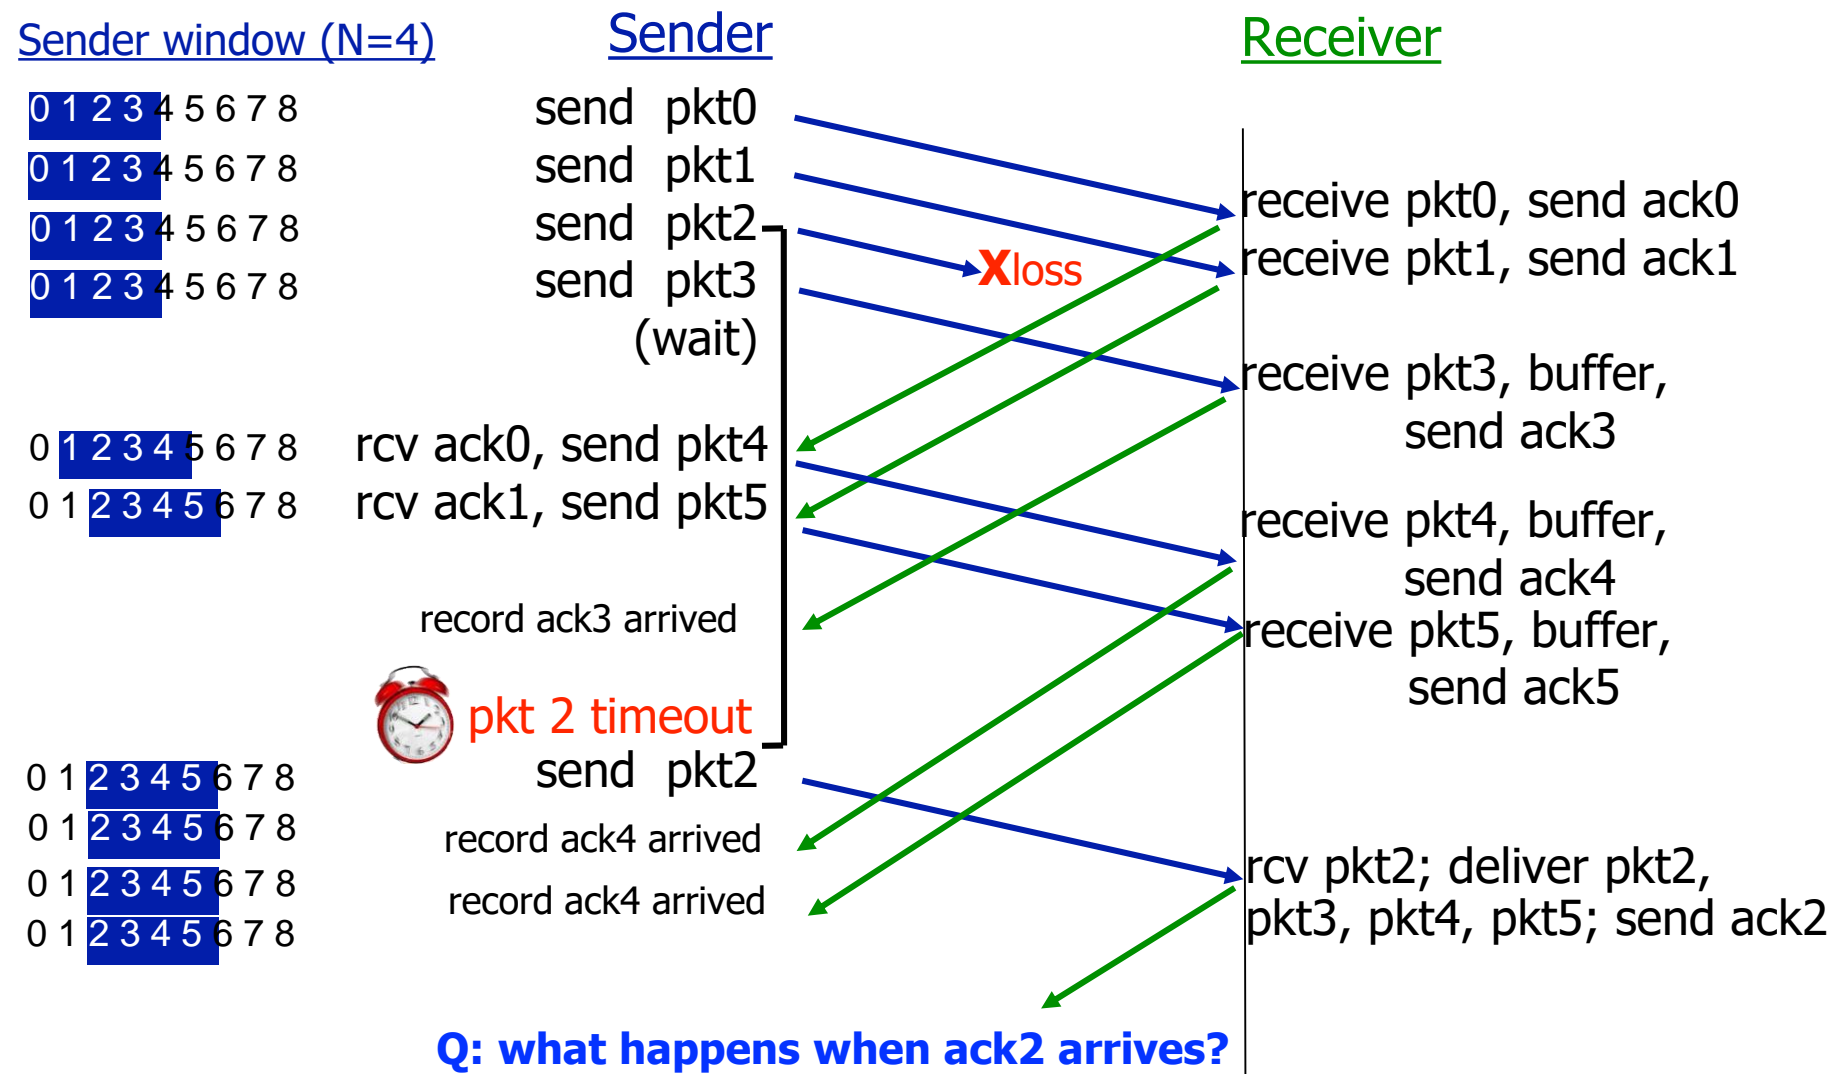
- **Sender**

- Receives data from upper layer:
 - If next available sequence number is in the sliding window, send the packet
- A timeout occurs for packet n:
 - Resend packet n, restart timer
- ACK(n) received for packet in current window:
 - Mark packet n as received
 - If n is smallest unACKed packet in window, advance the window base to next unACKed sequence number

- **Receiver**

- Receives packet n in receivers window
 - Send ACK(n)
 - If out-of-order, buffer
 - If in-order, deliver with other buffered, in-order packets to the upper layer. Also, advance window to next not-yet-received packet
- Receives packet n that has already been seen and ACKed by the receiver
 - Send ACK(n) again
- Otherwise:
 - Ignore the packet


Selective Repeat in Action



Selective repeat: dilemma

example:

- seq #'s: 0, 1, 2, 3
 - window size=3
 - receiver sees no difference in two scenarios!
 - duplicate data accepted as new in (b)
- Q:** what relationship between seq # size and window size to avoid problem in (b)?



*receiver can't see sender side.
receiver behavior identical in both cases!
something's (very) wrong!*

