

CS 330: Network Applications & Protocols

Network Layer

Galin Zhelezov
Department of Physical Sciences
York College of Pennsylvania



Overview of Network Layer

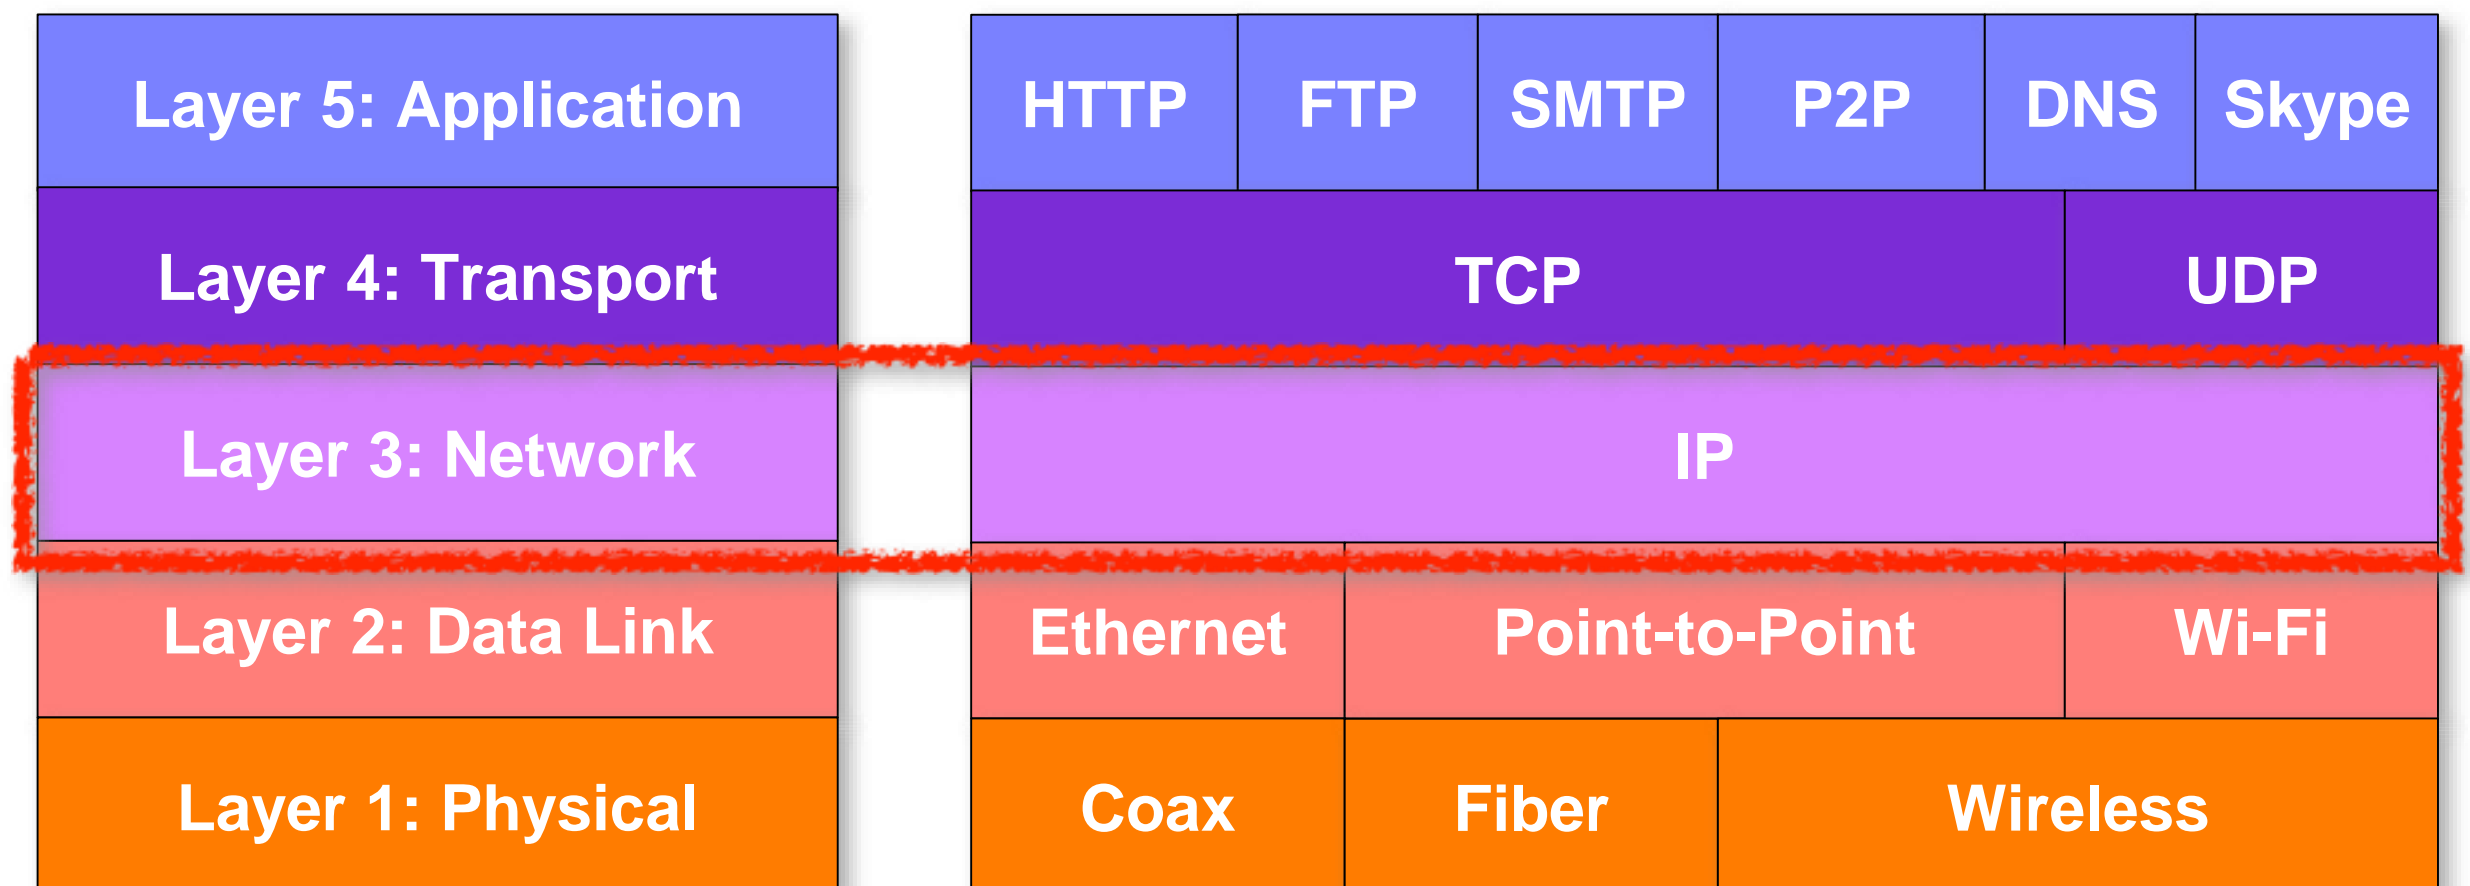
- **Introduction**
- **Virtual Circuit and Datagram Networks***
- **Router Architectures**
- **IP: Internet Protocol**
- **Routing algorithms**
- **Routing in the Internet**
- **Broadcast and multicast routing**

Overview of Network Layer

- **Introduction**
- **Virtual Circuit and Datagram Networks***
- **Router Architectures**
- **IP: Internet Protocol**
- **Routing algorithms**
- **Routing in the Internet**
- **Broadcast and multicast routing**

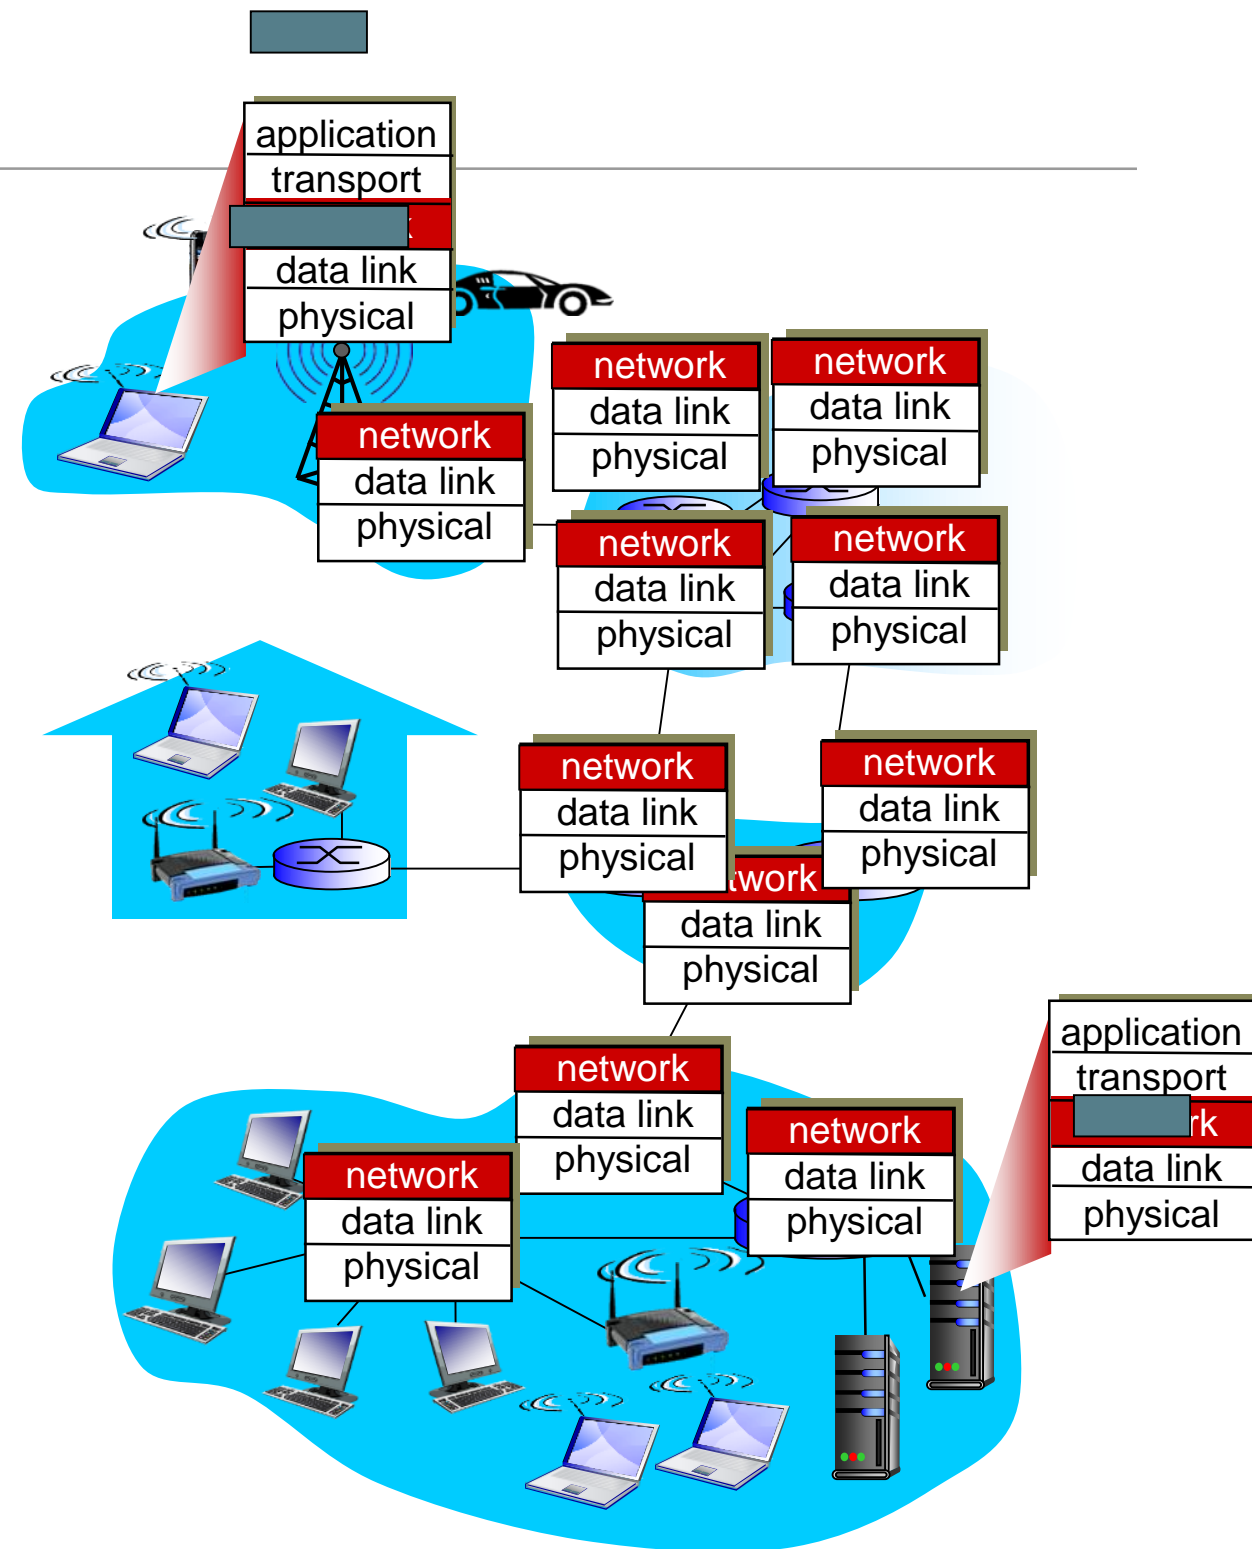
Protocol Layers

- **Top-Down Approach**

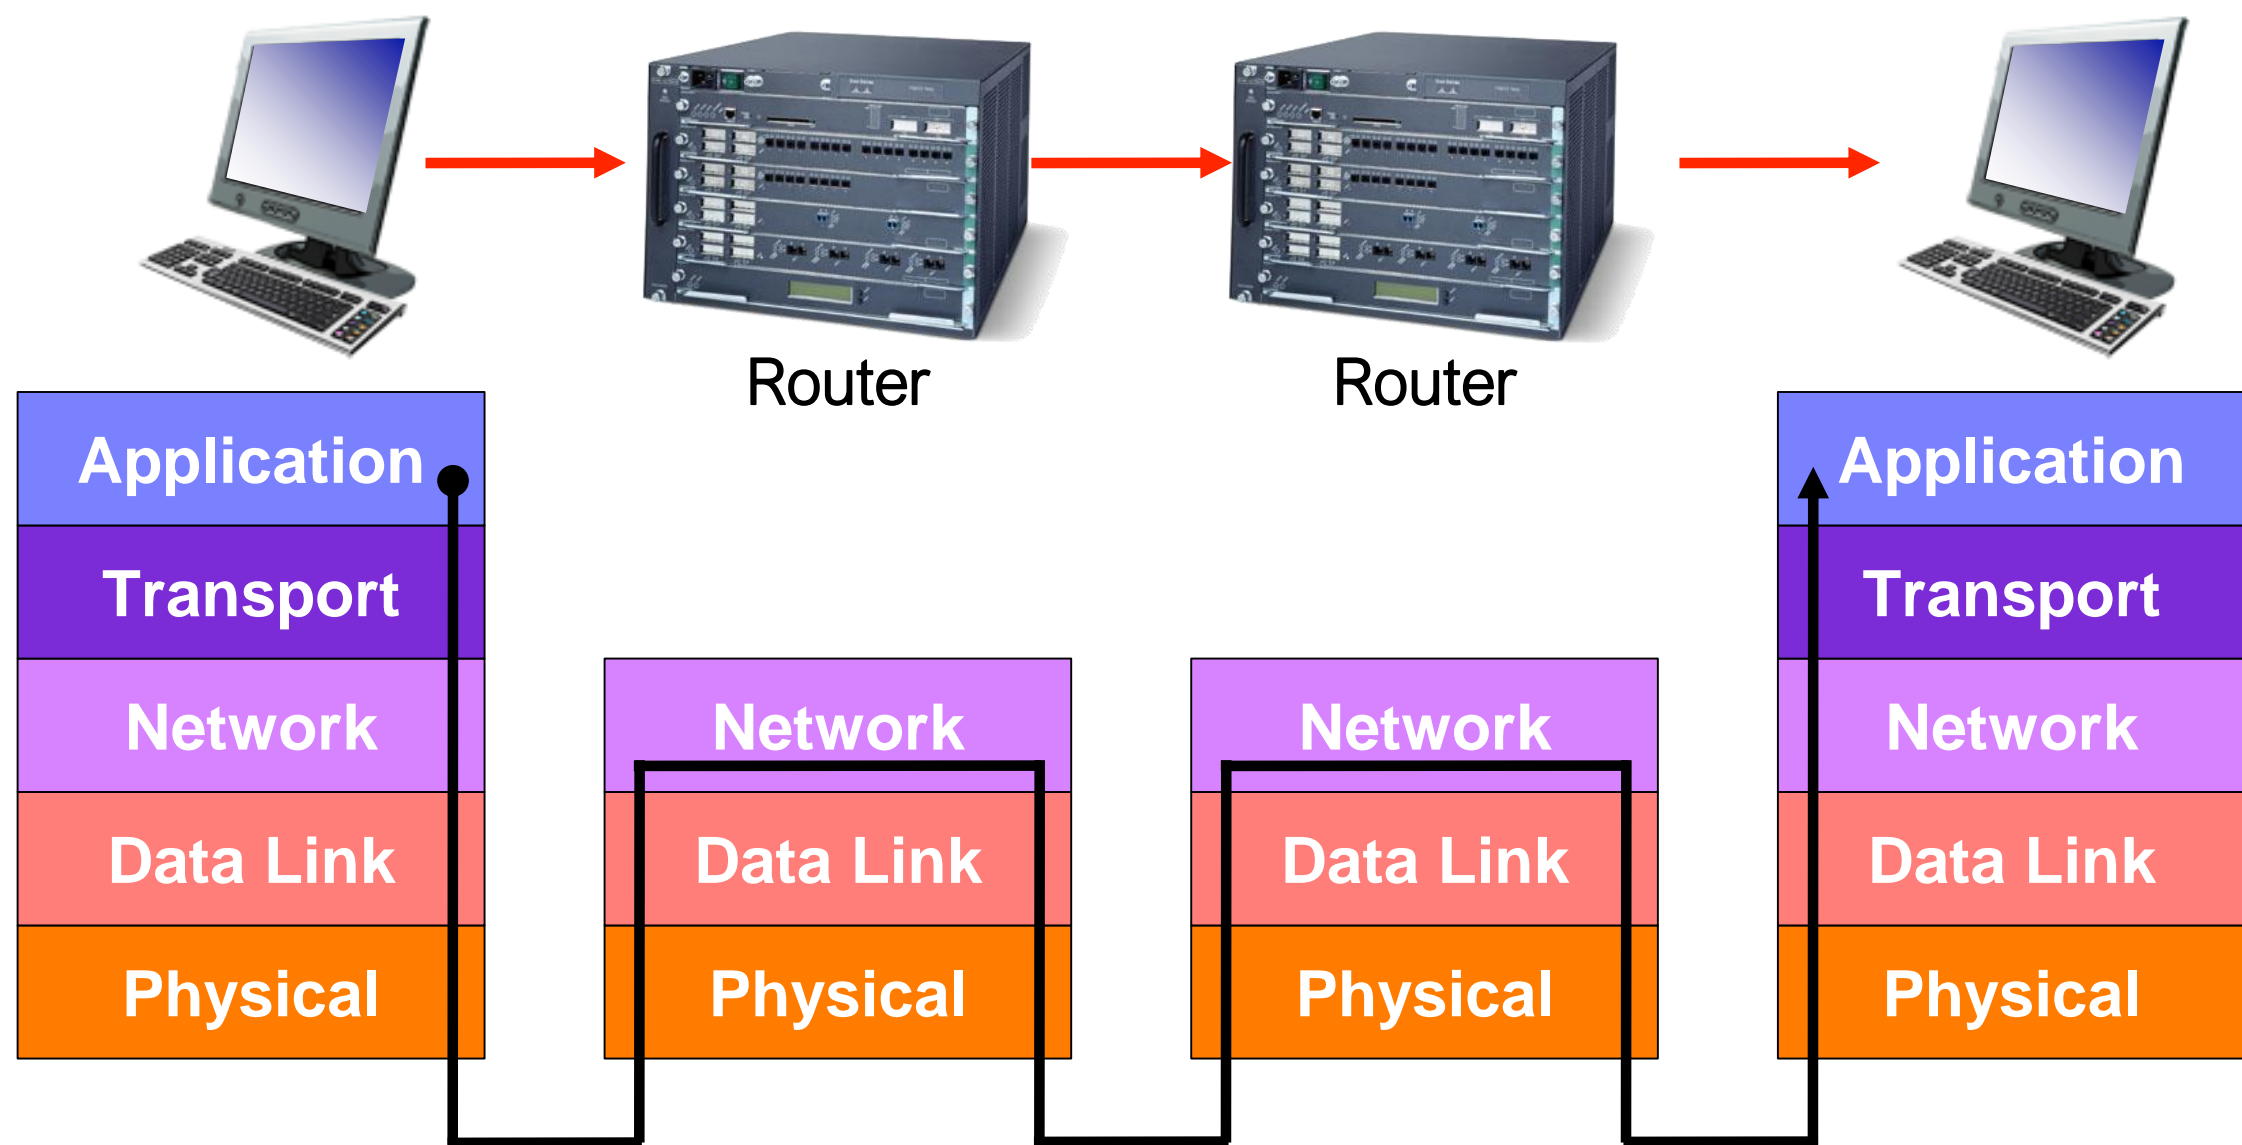


Network layer

- transport segment from sending to receiving host
- on sending side encapsulates segments into datagrams
- on receiving side, delivers segments to transport layer
- network layer protocols in **every** host, router
- router examines header fields in all IP datagrams passing through it



Network Layer



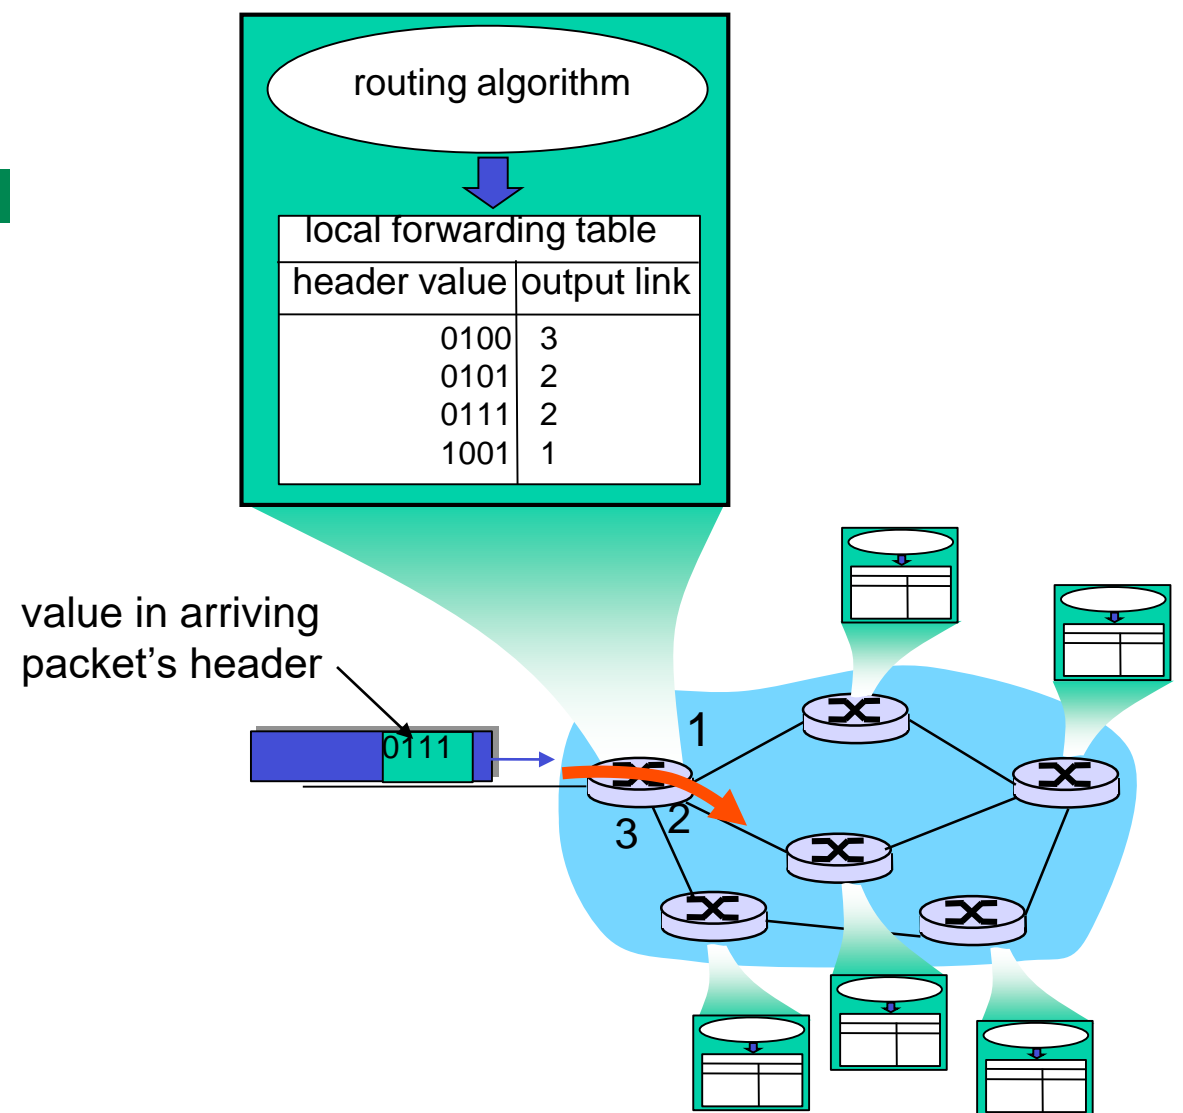
- **Network layer provides host to host communication service**
 - Network layer protocols in *every* host, router
 - A *router* is used to examine the header fields of all IP datagrams passing through it

Two Key Network-Layer Functions

- **Forwarding** - move packets from a single router's input to the appropriate router output
 - Happens within a single router
 - Routers contain a forwarding table to determine output
- **Routing** - determine the *complete route* taken by packets from the source host to the destination host
 - Complete route may include multiple routers
 - Requires routing algorithms and knowledge of other routers on the network
 - Routing algorithms are used to update a router's forwarding tables

Interplay Between Routing and Forwarding

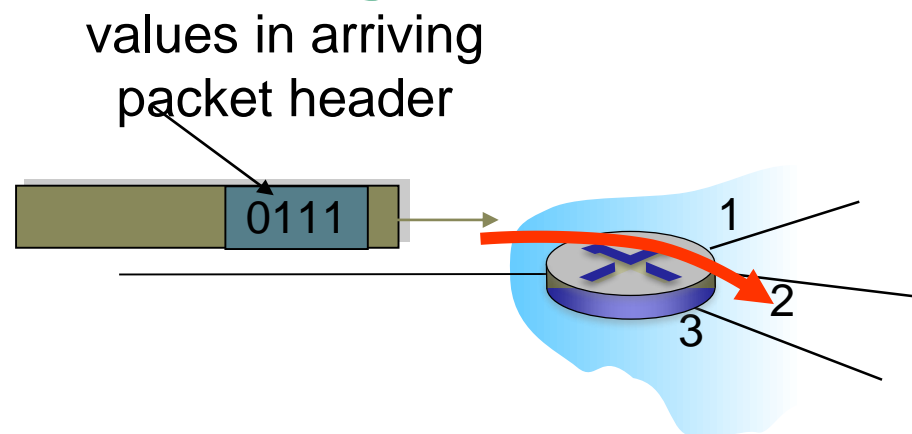
- **Routing algorithm determines end-end-path through network**
- **Forwarding table determines local forwarding at this router**
- **Example:**
 - Packet arrives with a header value of 0111, it is output on port 2 of the router



Network layer: data plane, control plane

Data plane

- **local, per-router function**
- **determines how datagram arriving on router input port is forwarded to router output port**
- **forwarding function**

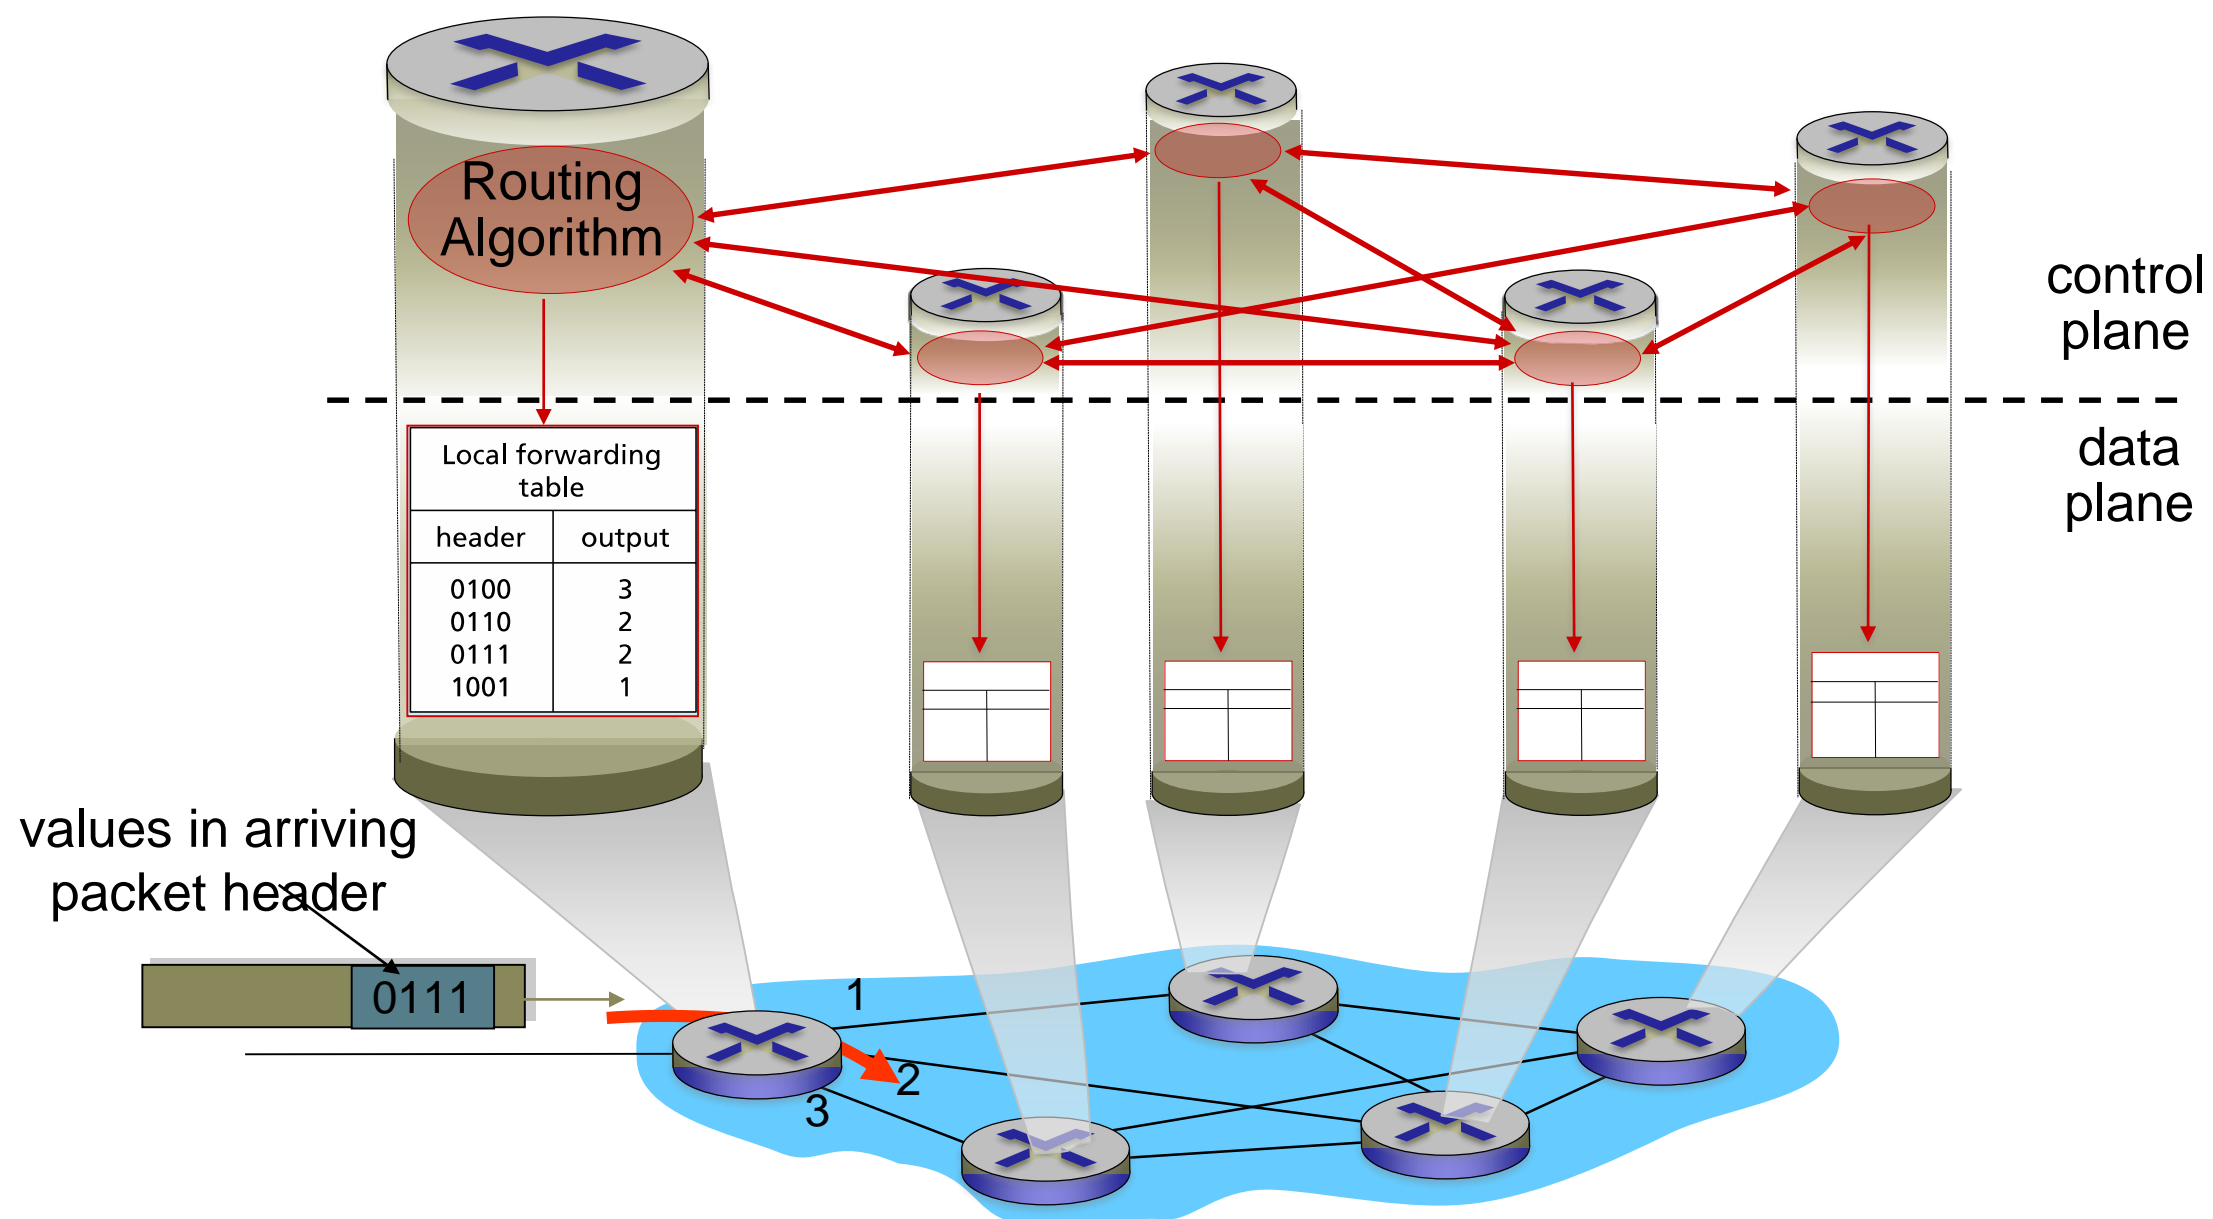


Control plane

- network-wide logic
- determines how datagram is routed among routers along end-end path from source host to destination host
- two control-plane approaches:
 - *traditional routing algorithms*: implemented in routers
 - *software-defined networking (SDN)*: implemented in (remote) servers

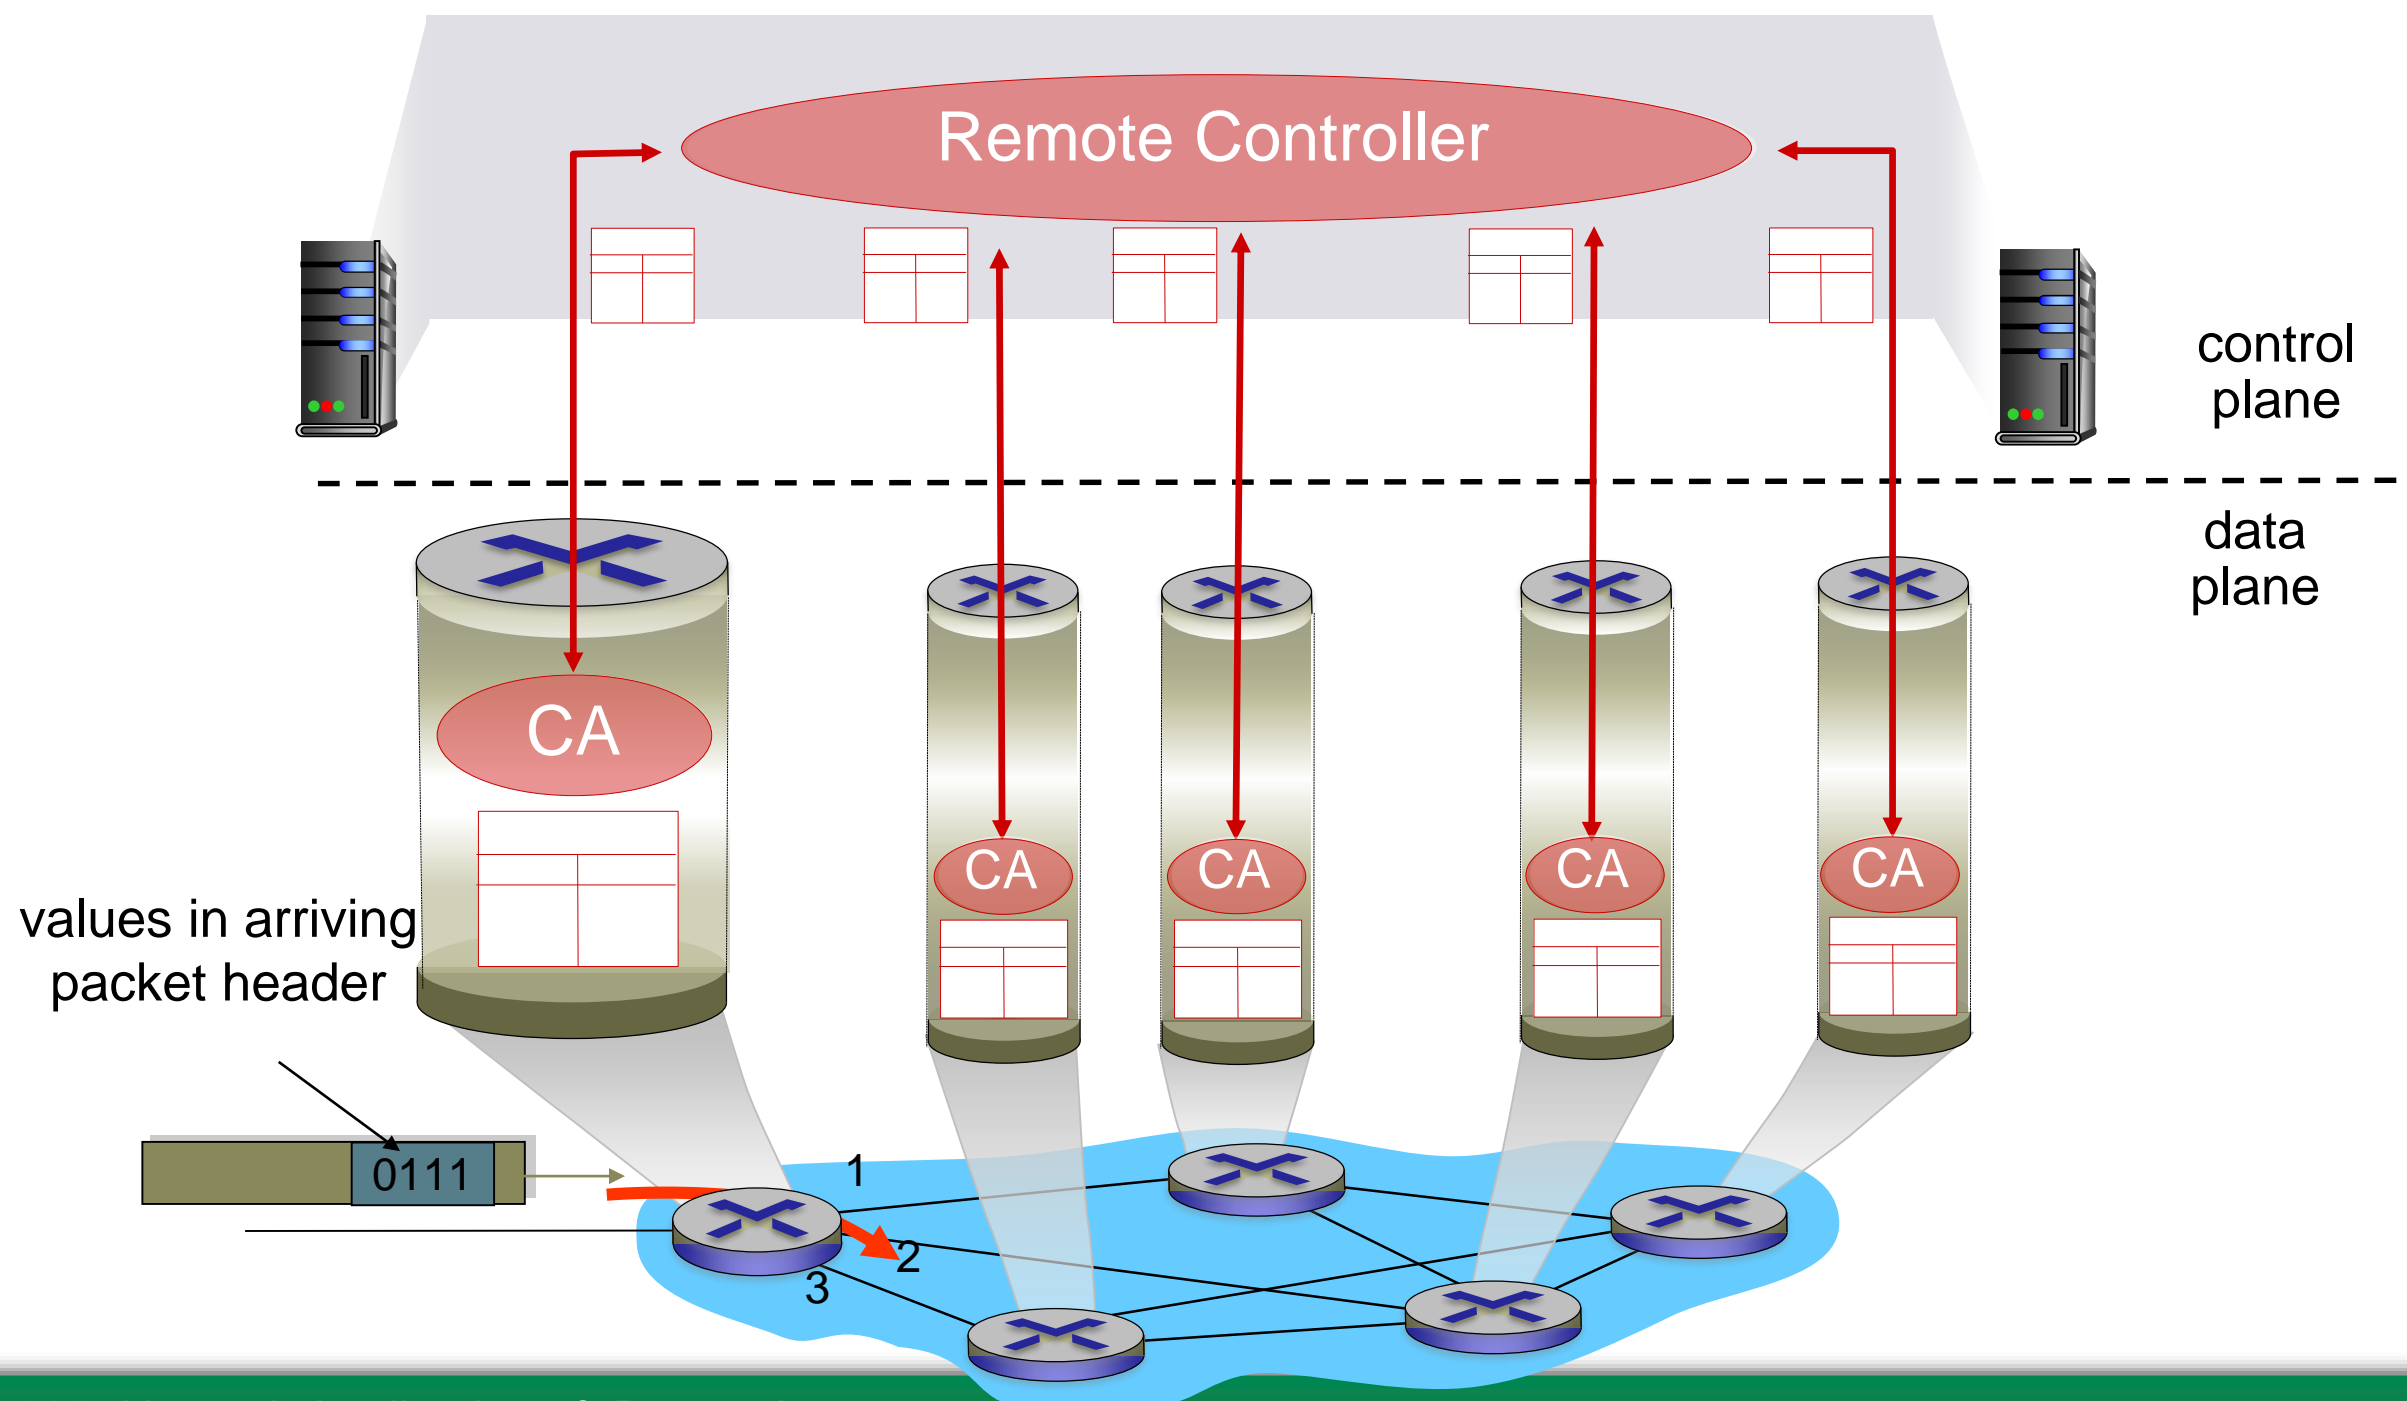
Per-router control plane

Individual routing algorithm components *in each and every router* interact in the control plane



Logically centralized control plane

A distinct (typically remote) controller interacts with local control agents (CAs)



A Third Network-Layer Function

- **Connection Setup**

- 3rd important function in some network architectures:
 - ATM, frame relay, X.25
- Before datagrams flow, two end hosts and intervening routers establish virtual connection
- Network vs. Transport layer connection service:
 - Network - connection is between two hosts (may also involve intervening routers in case of Virtual Circuits)
 - Transport - connection is between two processes

Overview of Network Layer

- **Introduction**
- **Virtual Circuit and Datagram Networks***
- **Router Architectures**
- **IP: Internet Protocol**
- **Routing algorithms**
- **Routing in the Internet**
- **Broadcast and multicast routing**

Connection-oriented or Connectionless Service

- **Network Layer can provide either connection-oriented or connectionless service (just like in the Transport Layer, TCP/UDP)**
 - Virtual Circuit Networks provide Network Layer connection-oriented service
 - Datagram Networks provide Network Layer connectionless service
- **Unlike the Transport Layer**
 - Services are offered as a host-to-host services
 - The network can provide *either* connection-oriented *or* connectionless service, but not both
 - Implemented in the core of the network (i.e. the routers)

Virtual Circuit Networks

- **Provide Network Layer connection-oriented service**
- **Requires that each connection be *setup* before data can flow**
- **Each packet carries a Virtual Circuit identifier (not destination host address)**
- **Every router on source-to-destination path maintains “state” for each connection that passes through it**
- **Router resources (e.g. bandwidth, buffers) may be allocated to a Virtual Circuit during setup**
 - Dedicated resources = Predictable service

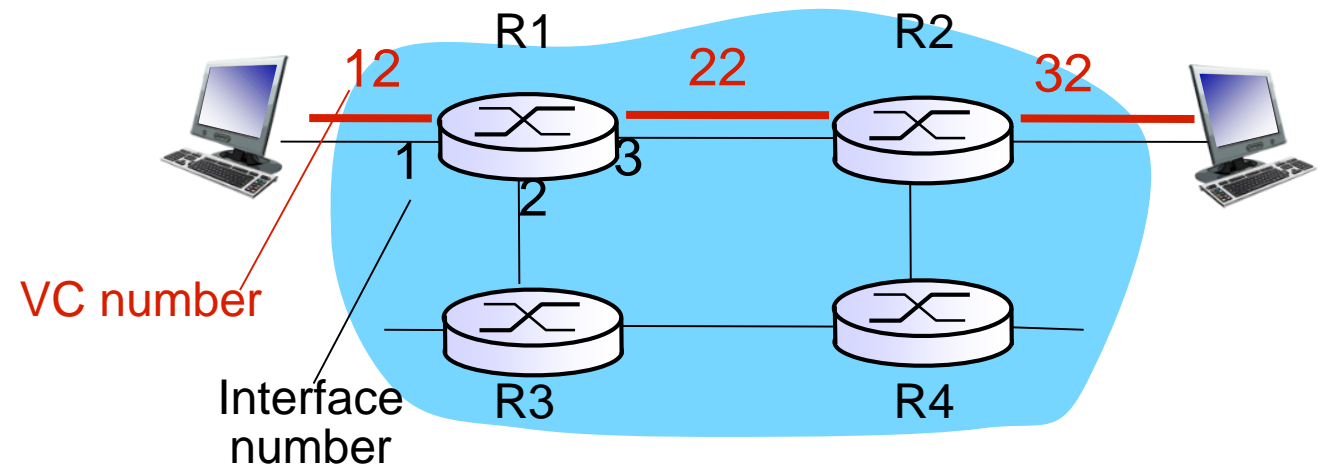
Virtual Circuit Implementation

- **A Virtual Circuit consists of:**
 - (1) A path between the source and destination hosts
 - This path is made of a series of links and routers
 - (2) Virtual Circuit numbers, one number for each link along the path
 - (3) Entries in the forwarding tables of the routers along the path
- **A packet belonging to virtual circuit carries a virtual circuit number (rather than dest address)**
 - VC number can be changed on each link
 - At each router along the path, a new VC number comes from the forwarding table

Virtual Circuit Forwarding Table

- **Each VC router maintains connection state information**

- When a connection is setup, entries are added to the VC forwarding table
- When a connection is torn down, entries are removed from the table

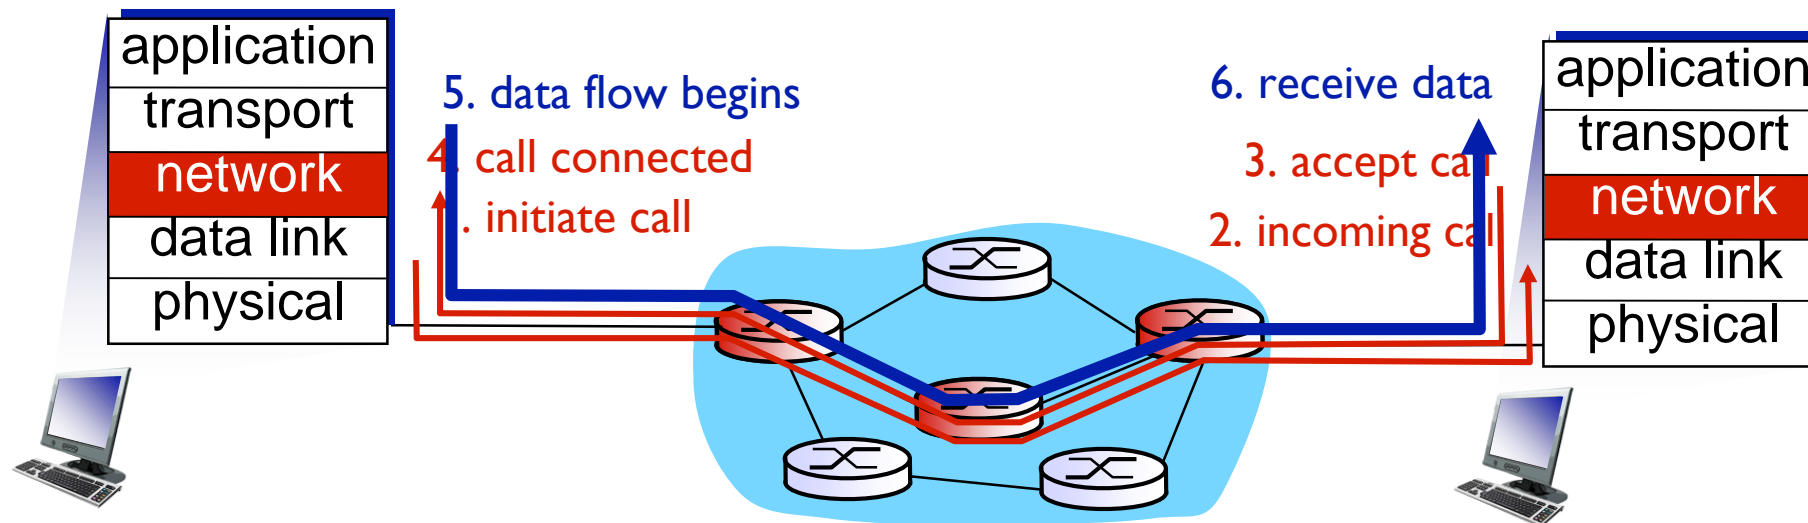


Incoming Interface	Incoming VC #	Outgoing Interface	Outgoing VC #
1	12	3	22
2	63	1	18
3	7	2	17
1	97	3	87
...

Virtual Circuits Signaling Protocols

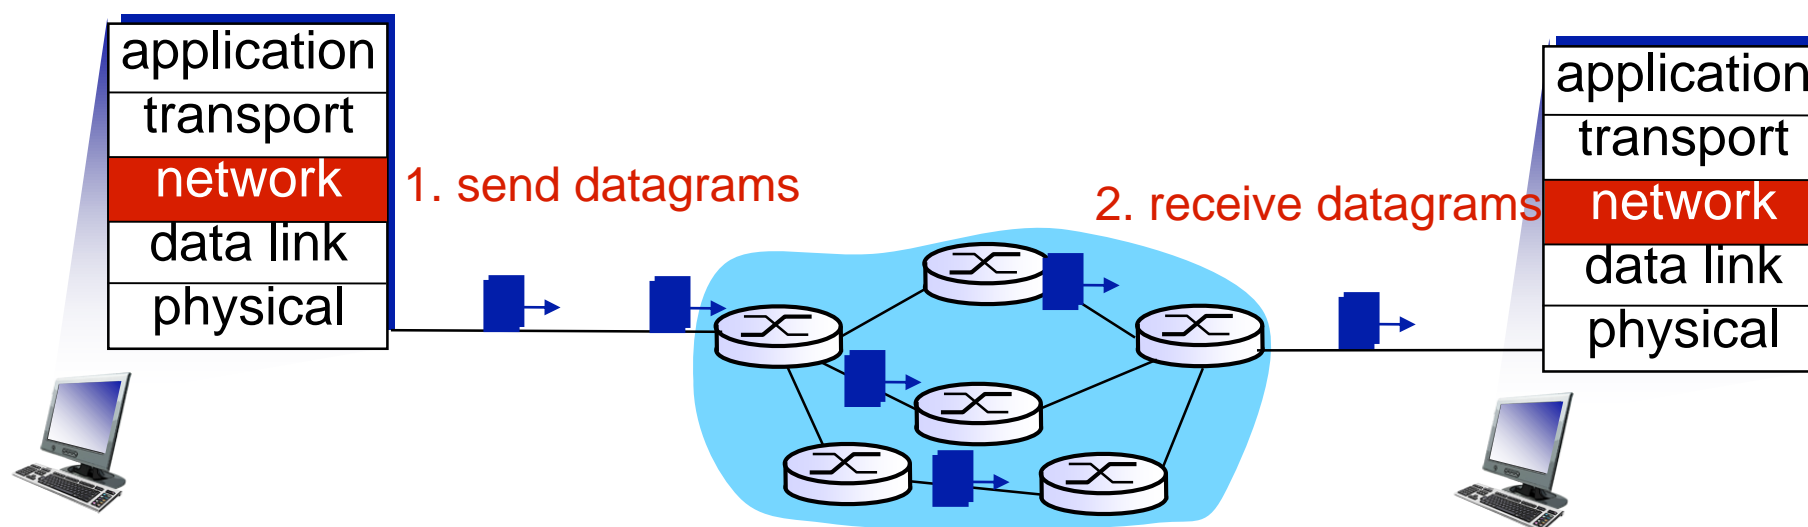
- **Protocols used to setup, maintain, and teardown VCs**

- All protocol message pass through all routers on the way to the destination
- As messages pass through routers, routers setup the VC



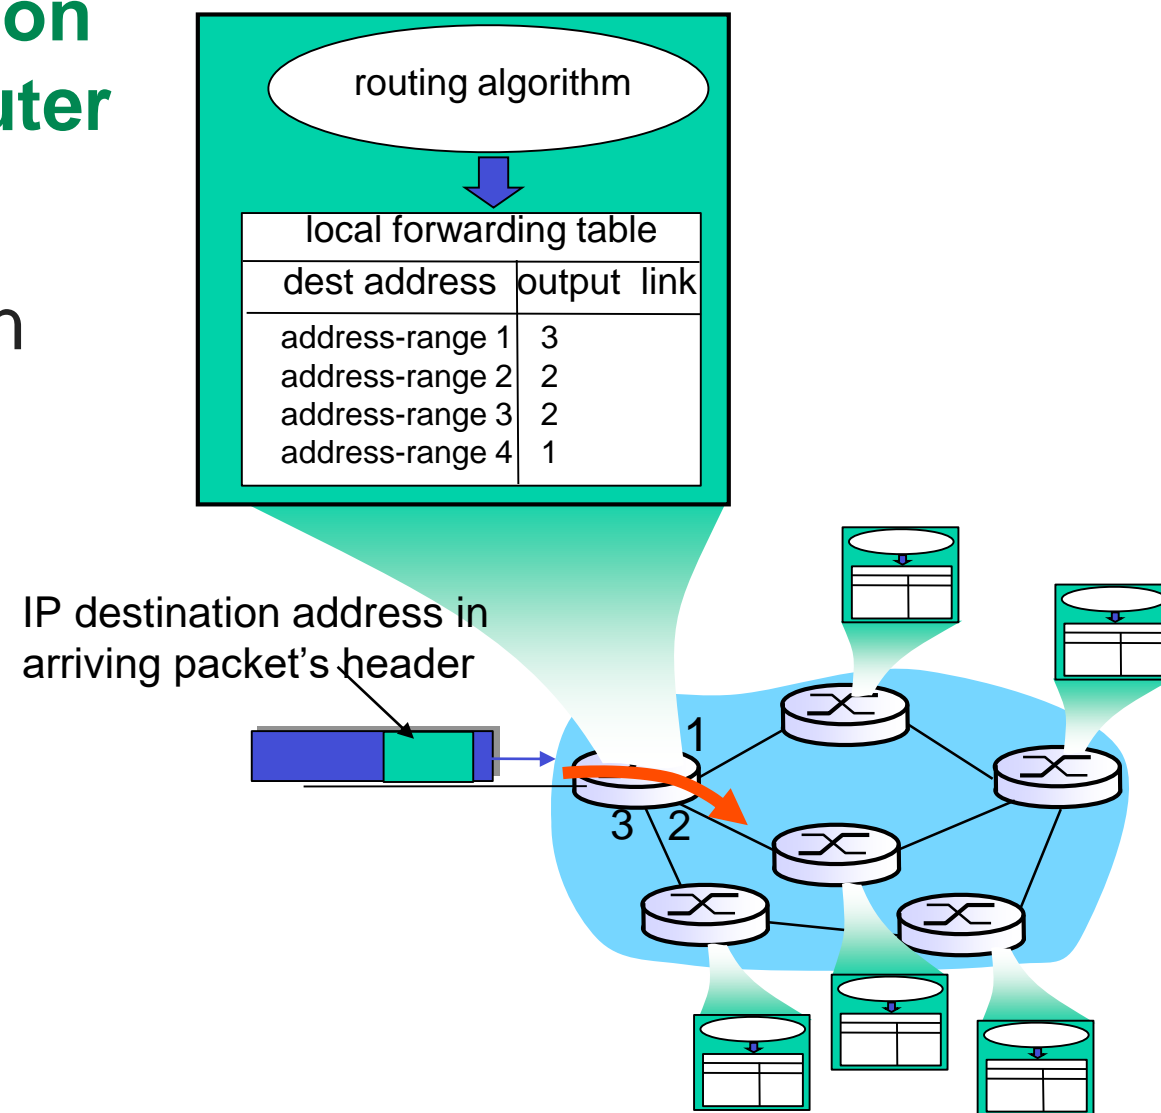
Datagram Networks

- Requires no call setup at network layer
- Routers contain no state about end-to-end connections
- No network-level concept of a “connection”
- Packets are forwarded using destination host address
 - To send a packet, a host puts a destination address on the packet and simply inserts it into the network



Datagram Forwarding Table

- **Each router contains a forwarding table that maps a packet's destination address to an output link on the router**
 - 4 billion possible IP addresses so rather than list individual destination addresses, list range of addresses (aggregate table entries)
 - Use **Longest-Prefix Matching** to determine which address-range a packet falls into



Building a Router Forwarding Table

- **Given a router with multiple output links, specify ranges of addresses to determine which addresses are output on which output link**
 - Example below shows a router with four output links
 - Packet with destination **200.23.24.17** would go out on **interface 1**
 - Packet with destination **122.17.10.1** would go out on **interface 3**

Destination Address Ranges	Destination Address Ranges (again)	Outgoing Link Interface
11001000 00010111 00010000 00000000 - 11001000 00010111 00010111 11111111	200.23.16.0 - 200.23.23.255	0
11001000 00010111 00011000 00000000 - 11001000 00010111 00011000 11111111	200.23.24.0 - 200.23.24.255	1
11001000 00010111 00011001 00000000 - 11001000 00010111 00011111 11111111	200.23.25.0 - 200.23.31.255	2
otherwise		3

Longest-Prefix Matching

- An implementation of ranged matching
- When searching the forwarding table for given destination address, use the longest address prefix that matches the destination address
 - Example 1: 11001000 00010111 00010110 10100001 → 0
 - Example 2: 11001000 00010111 00011000 10101010 → 1
 - Matches multiple table entries, send out on interface that has the *longest* match

Destination Address Ranges	Destination Address Ranges (again)	Outgoing Link Interface
11001000 00010111 00010*** *****	200.23.16.0 - 200.23.23.255	0
11001000 00010111 00011000 *****	200.23.24.0 - 200.23.24.255	1
11001000 00010111 00011*** *****	200.23.24.0 - 200.23.31.255	2
otherwise		3

Datagram Network or Virtual Circuit Network

- **Internet (Datagram Network)**

- Data exchange among computers
 - “Elastic” service, no strict timing requirements
- Many link types
 - Different characteristics
 - Uniform service difficult
- “Smart” end systems (computers)
 - Can adapt, perform control, error recovery
 - Simple inside network, complexity at “edge”

- **ATM (Virtual Circuit Network)**

- Evolved from telephony
- Like human conversation:
 - Strict timing, reliability requirements
 - Need for guaranteed service
- “Dumb” end systems
 - Telephones
 - Complexity inside network

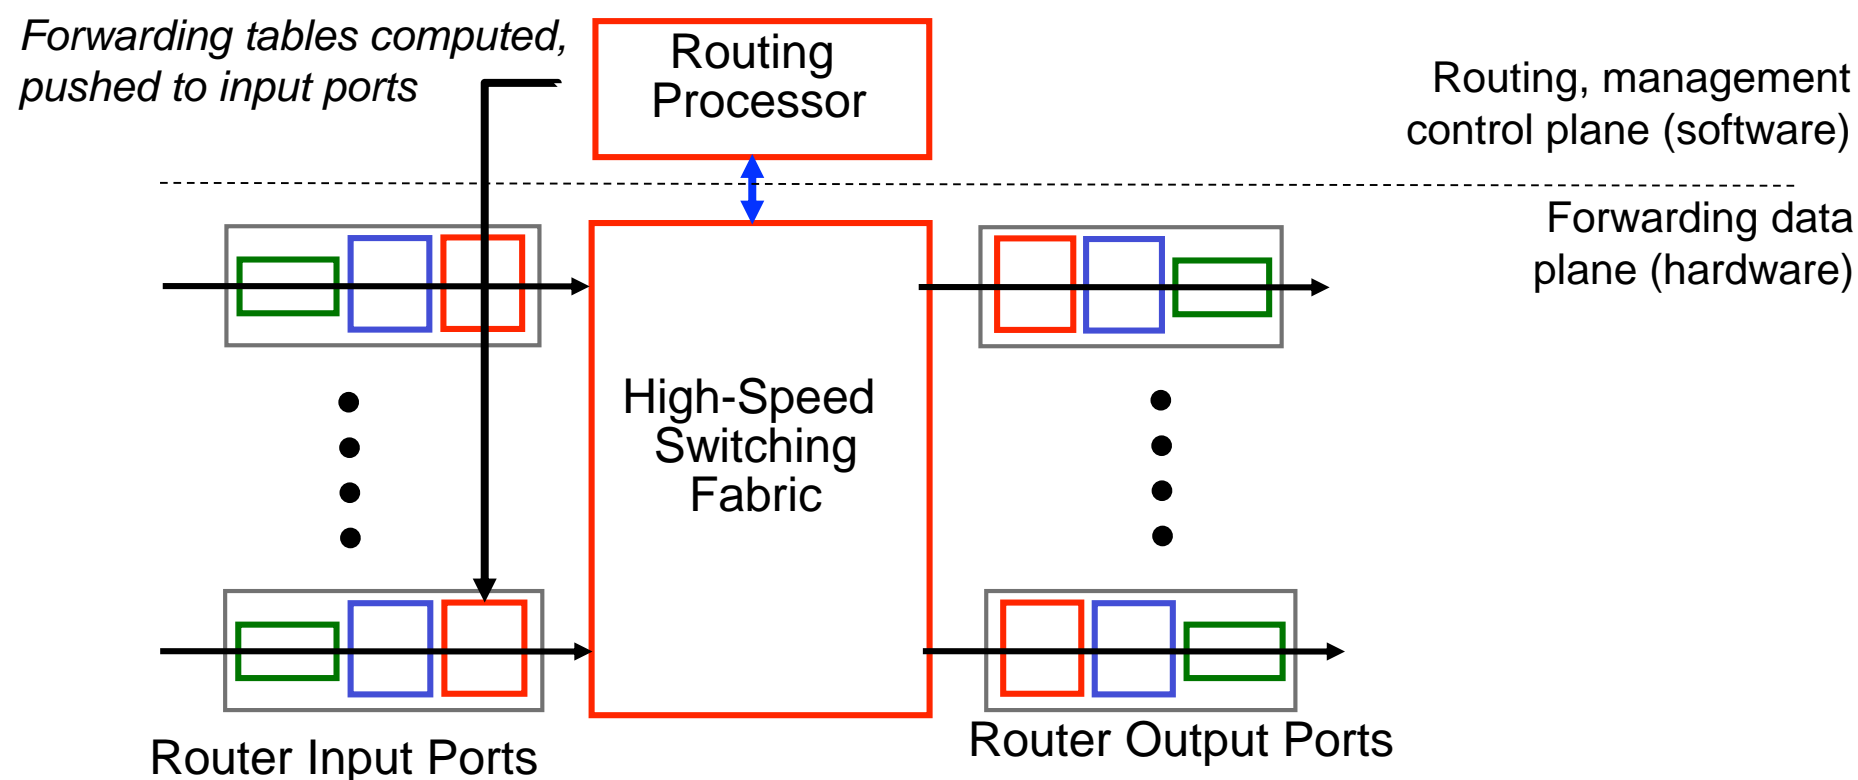
Overview of Network Layer

- **Introduction**
- **Virtual Circuit and Datagram Networks***
- **Router Architectures**
- **IP: Internet Protocol**
- **Routing algorithms**
- **Routing in the Internet**
- **Broadcast and multicast routing**

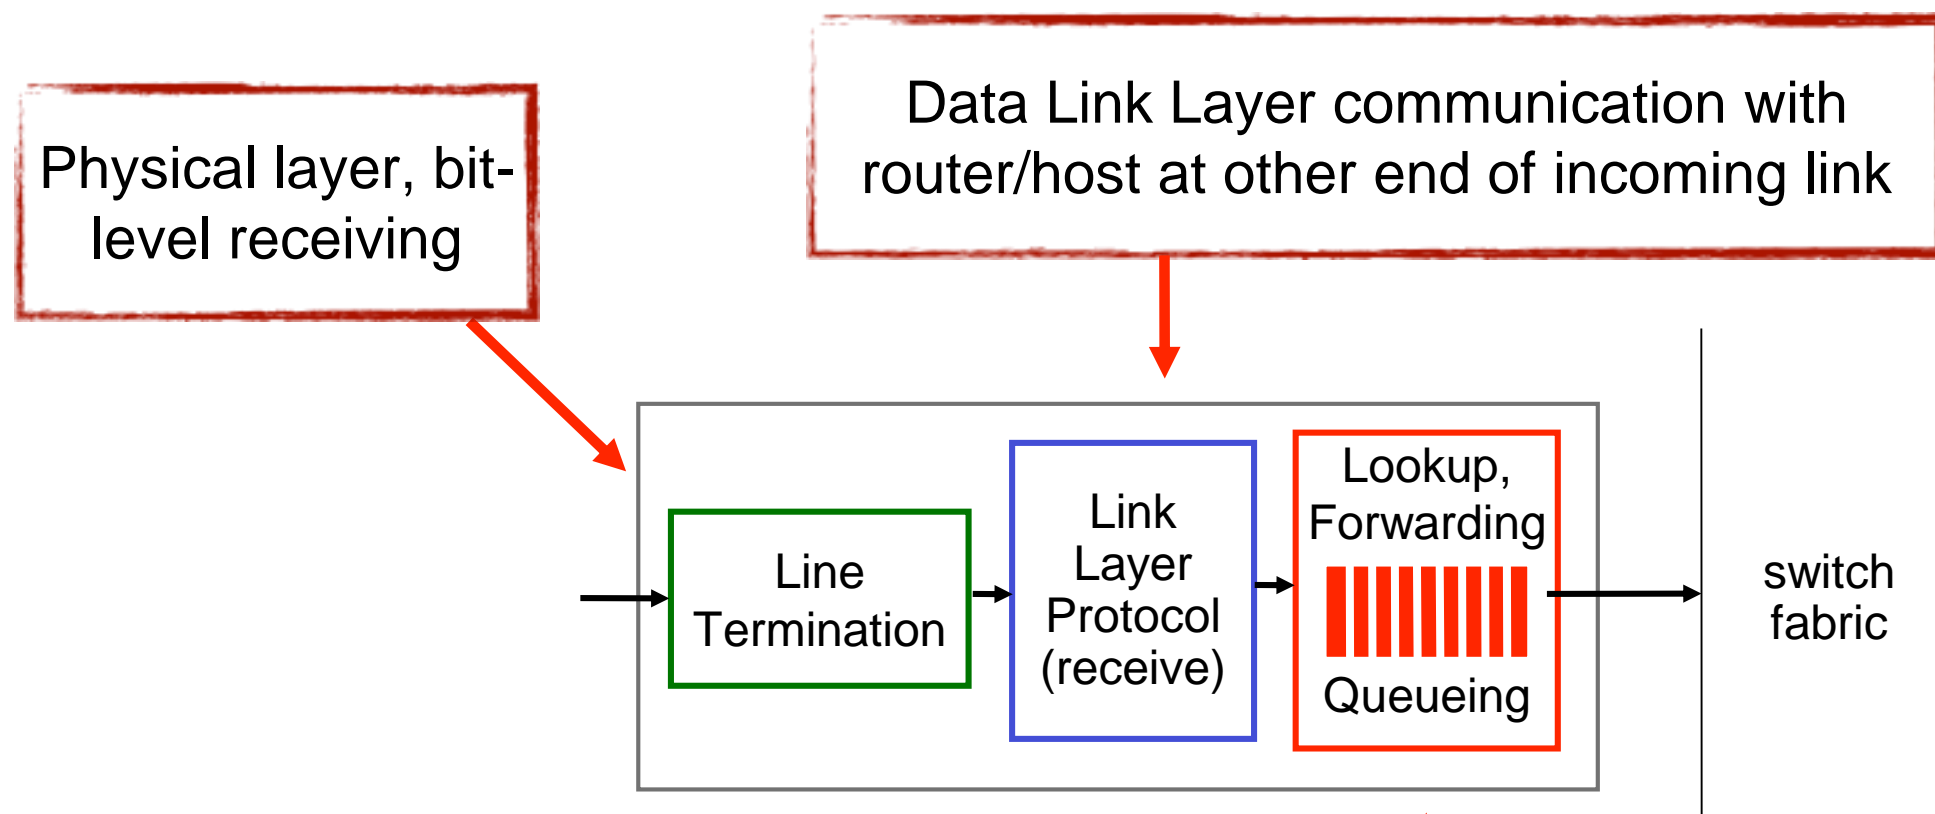
Router Architecture Overview

- **Routers have two key functions**

- Run routing algorithms
- Forward datagrams from incoming to outgoing link



Input Port Functions

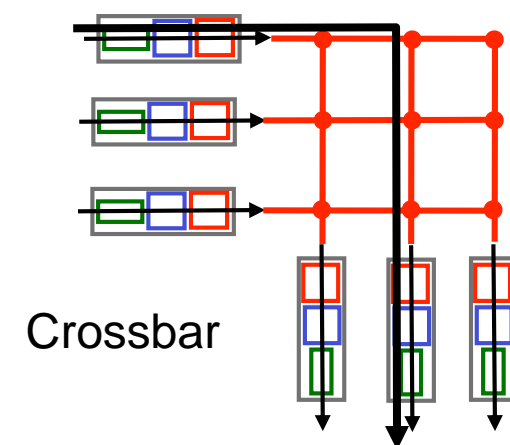
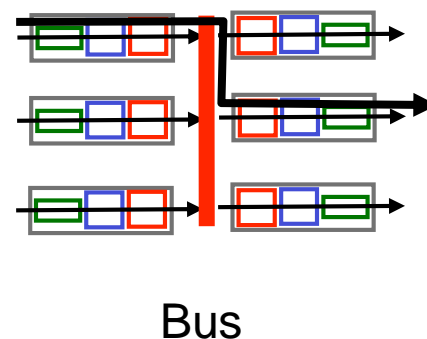
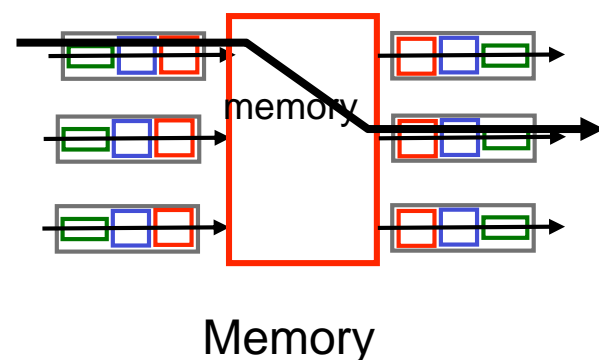


Decentralized Switching

- Given datagram destination, **lookup output port using the forwarding table** in input port memory
- Complete input port processing at 'line speed', or close to it
- If datagrams arrive faster than they can be processed, **queue them up**

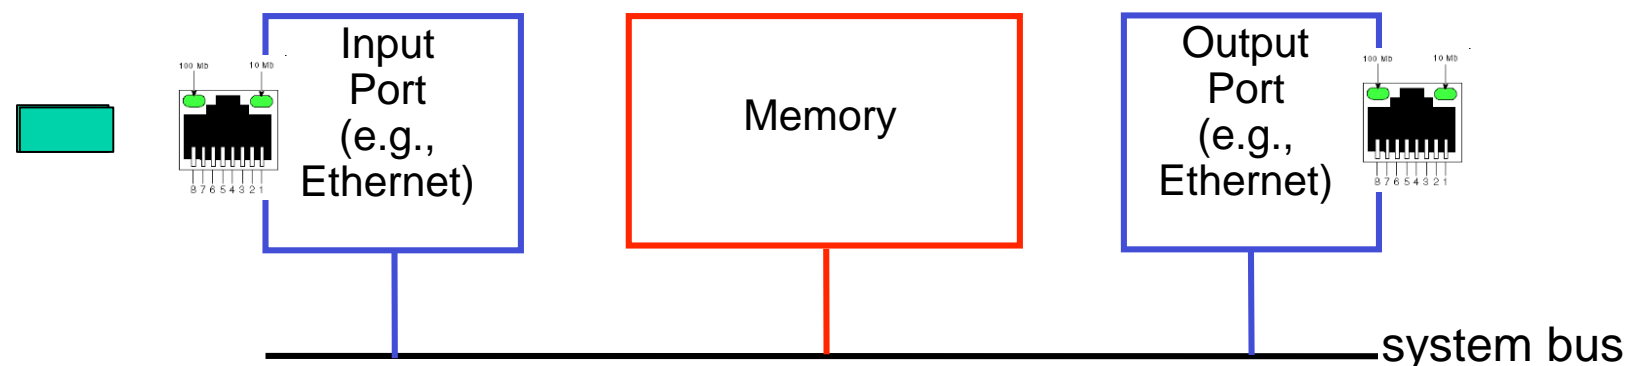
Switching Fabrics

- Used to transfer packets from input buffers to appropriate output buffers
- **Switching rate** is the rate at which packets can be transfer from inputs to outputs
 - Often measured as multiple of input/output line rate
 - If N inputs, switching rate N times line rate is desirable
- Three types of switching fabrics



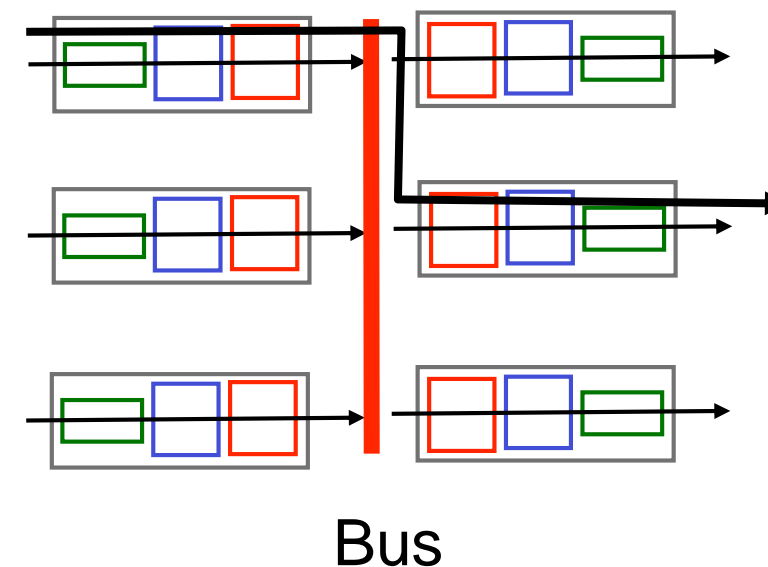
Switching Via Memory

- **Used in first generation routers**
 - Traditional computers with switching under direct control of CPU
 - Packet copied to system's memory
 - Speed limited by memory bandwidth (2 bus crossings per datagram)
- **Still used in some systems today**



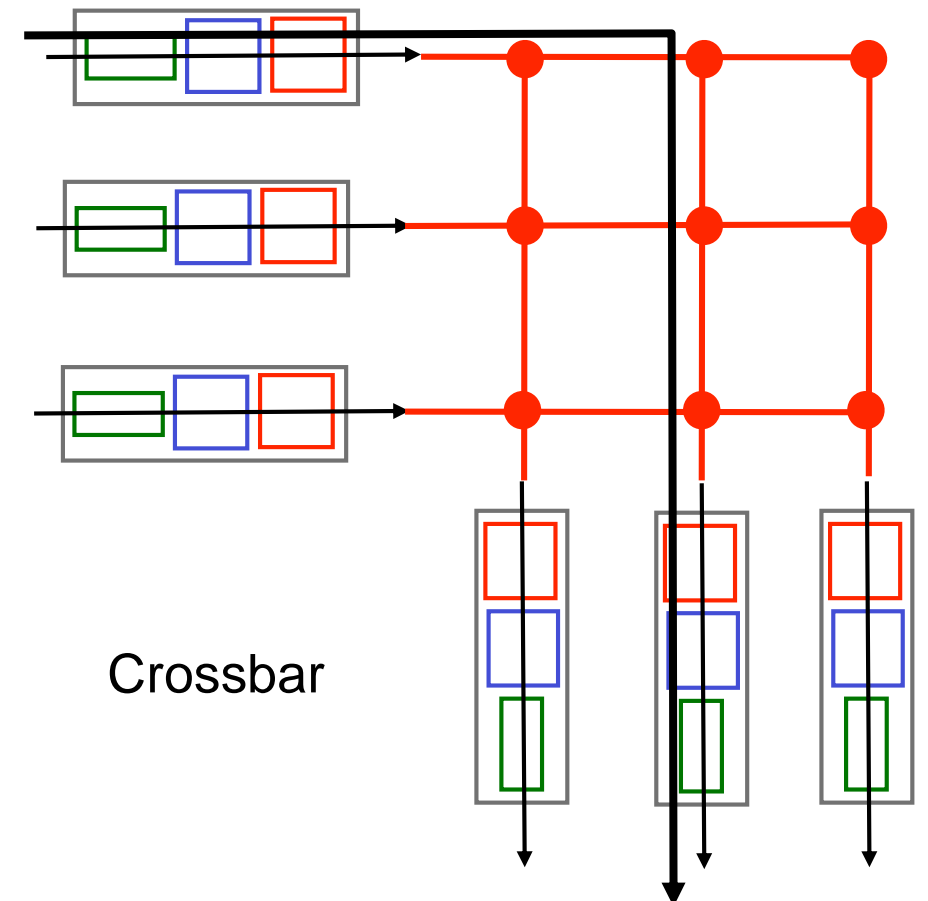
Switching Via a Bus

- **Datagram is transferred from input port memory to output port memory via a shared bus**
 - Prepend header internal to router to get packet to appropriate output port
- **Must deal with bus contention**
 - Switching speed limited by bus bandwidth
- **Example: 32 Gbps bus, Cisco 5600: sufficient speed for access and enterprise routers**



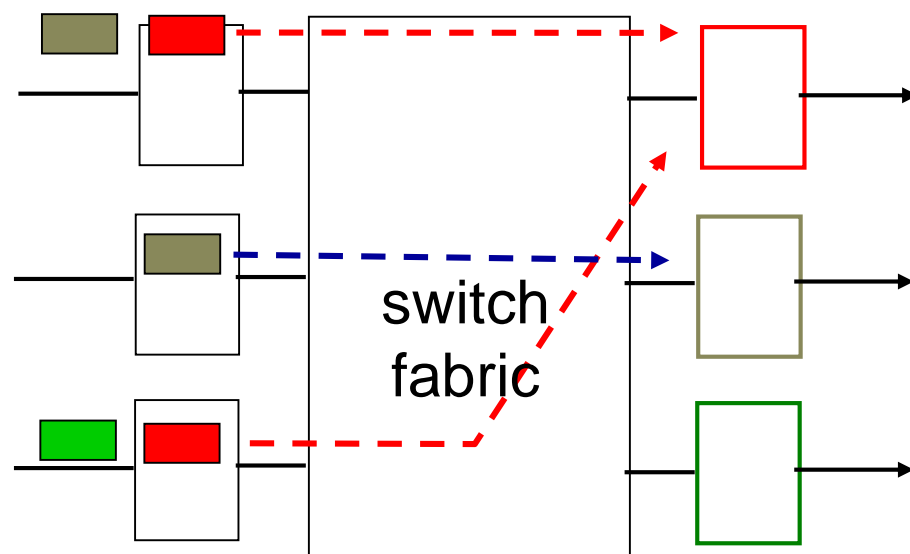
Switching Via Interconnection Network

- **Overcomes bus bandwidth limitations**
- **Banyan networks, crossbar, other interconnection nets**
 - Interconnect contains $2 \cdot N$ buses that connects N input ports to N output ports
- **Can forward multiple packets at the same time (unless they are destined for same output)**
- **Example: The Cisco 12000 switches 60 Gbps through the interconnection network**

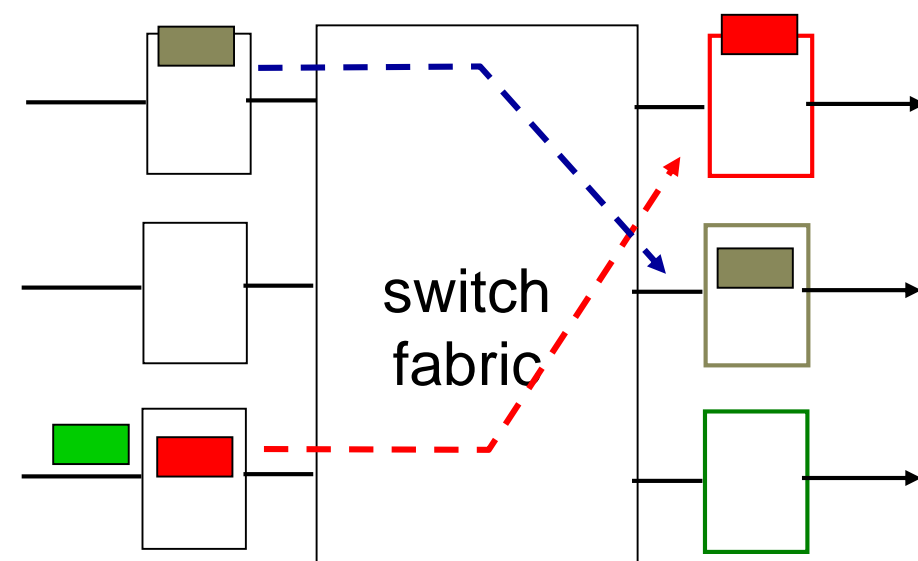


Input port queuing

- fabric slower than input ports combined -> queueing may occur at input queues
 - *queueing delay and loss due to input buffer overflow!*
- **Head-of-the-Line (HOL) blocking:** queued datagram at front of queue prevents others in queue from moving forward

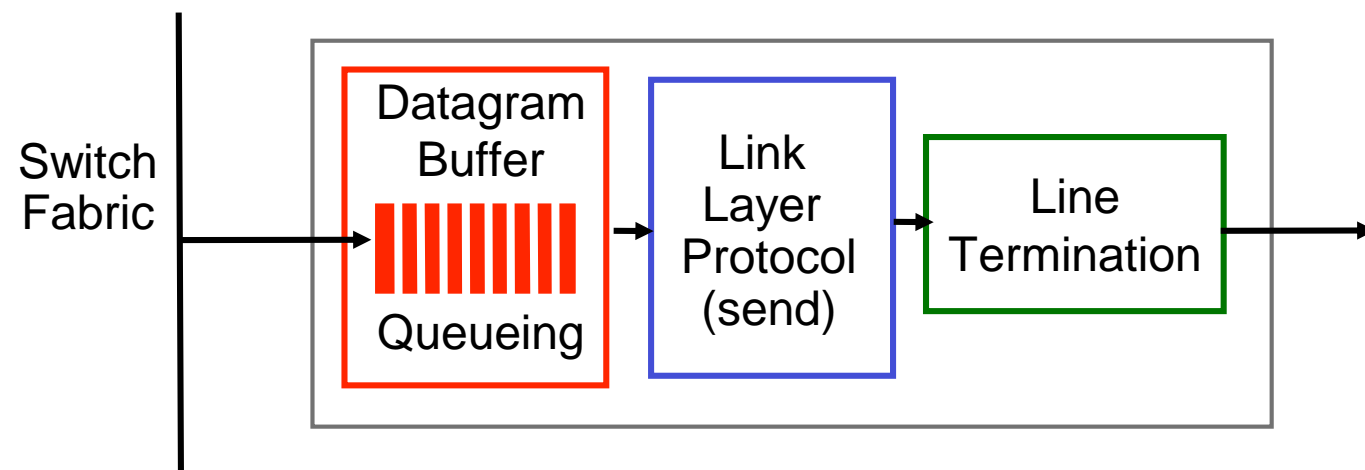


output port contention:
only one red datagram can be
transferred.
lower red packet is blocked



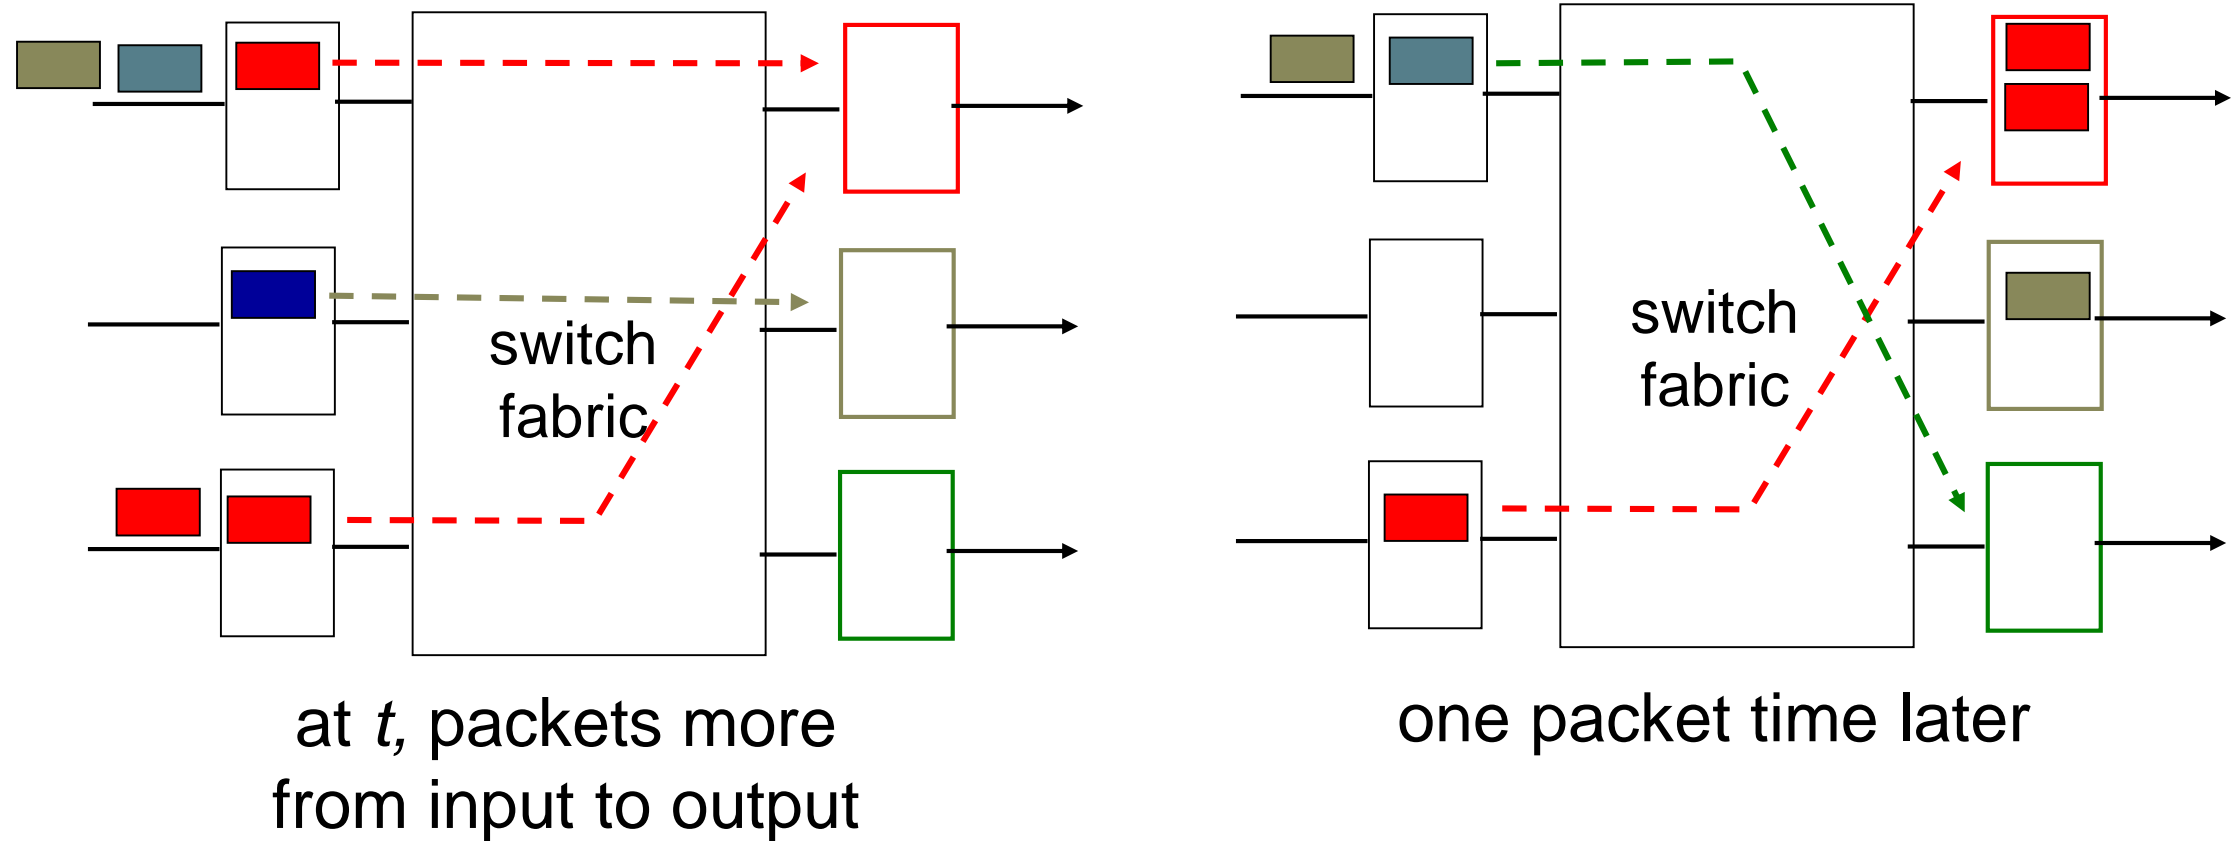
one packet time later:
green packet
experiences HOL
blocking

Output Ports



- **Buffers outgoing datagrams as required**
 - When datagrams arrive from fabric faster than they can be transmitted
 - Datagram (packets) can be lost due to congestion, lack of buffers
- **Scheduling to choose among queued datagrams for transmission**
 - Priority scheduling – who gets best performance, network neutrality
- **Data Link Layer communication with router/host at other end of outgoing link**
- **Physical Layer bit transmission**

Output port queueing



- **buffering when arrival rate via switch exceeds output line speed**
- ***queueing (delay) and loss due to output port buffer overflow!***

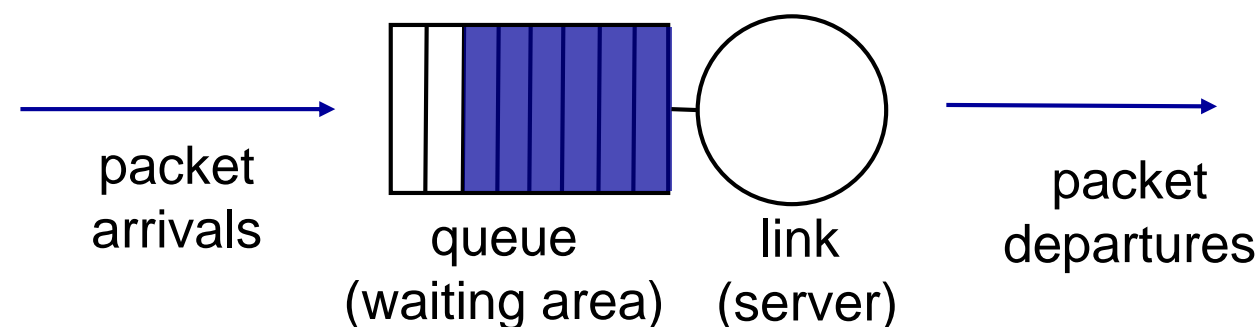
How much buffering?

- **RFC 3439 rule of thumb: average buffering equal to “typical” RTT (say 250 msec) times link capacity C**
 - e.g., C = 10 Gpbs link: 2.5 Gbit buffer
- **recent recommendation: with N flows, buffering equal to**

$$\frac{RTT \cdot C}{\sqrt{N}}$$

Scheduling mechanisms

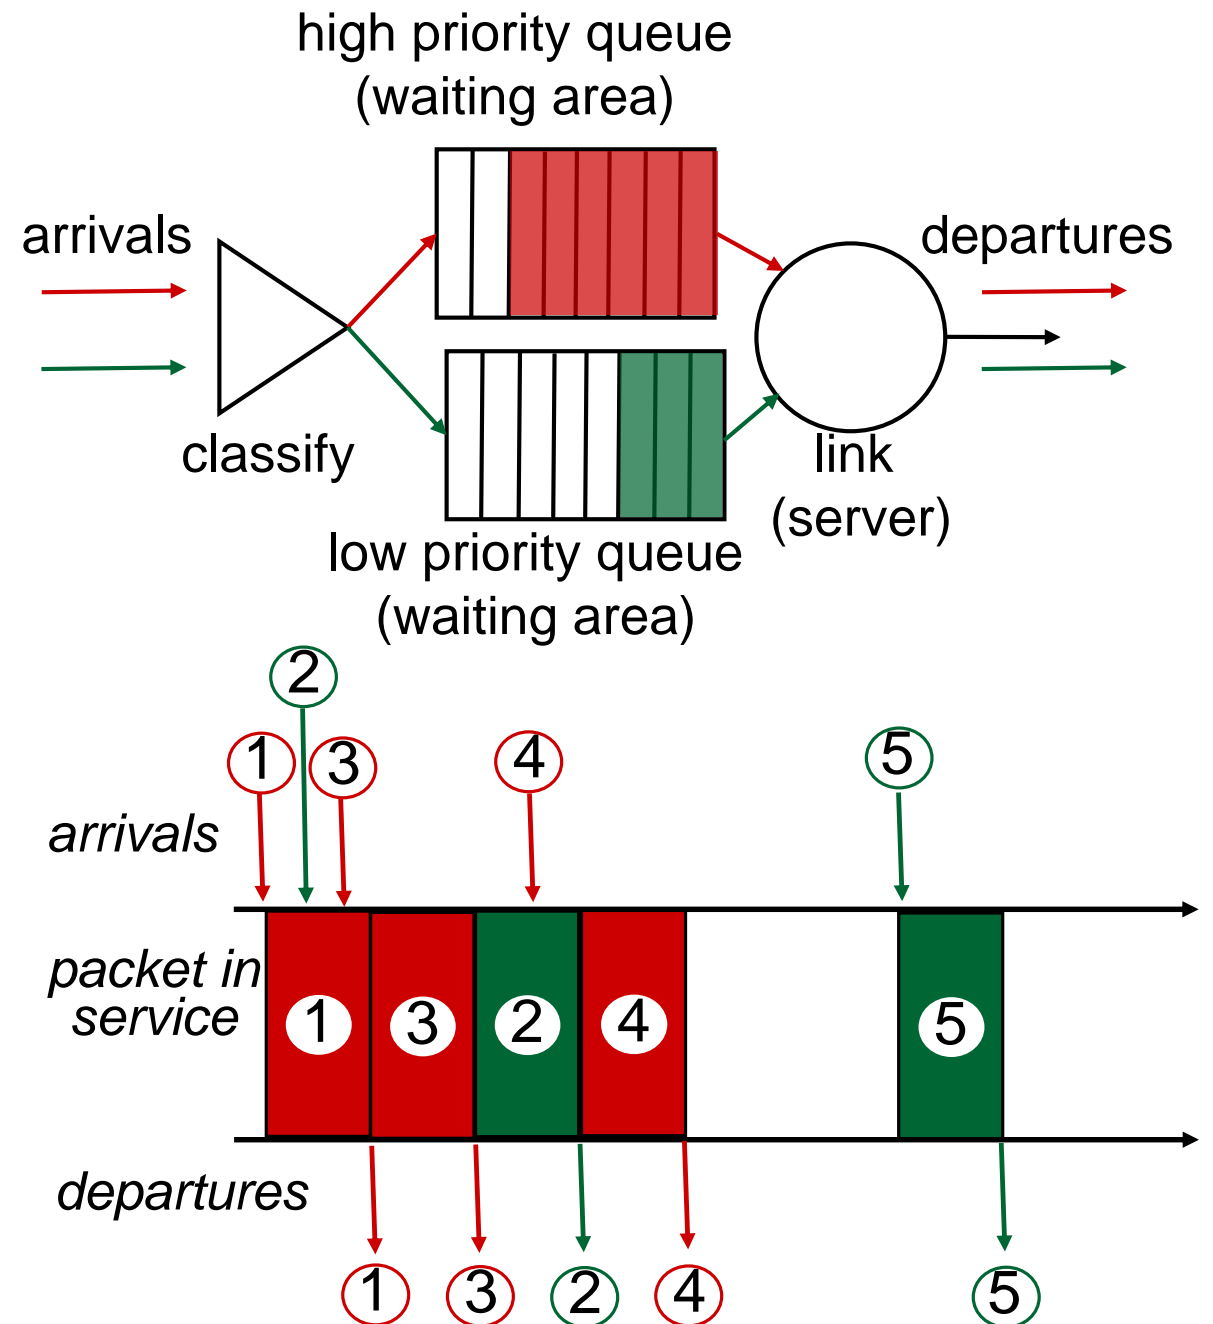
- **scheduling**: choose next packet to send on link
- **FIFO (first in first out) scheduling**: send in order of arrival to queue
 - real-world example?
 - **discard policy**: if packet arrives to full queue: who to discard?
 - **tail drop**: drop arriving packet
 - **priority**: drop/remove on priority basis
 - **random**: drop/remove randomly



Scheduling policies: priority

priority scheduling: send highest priority queued packet

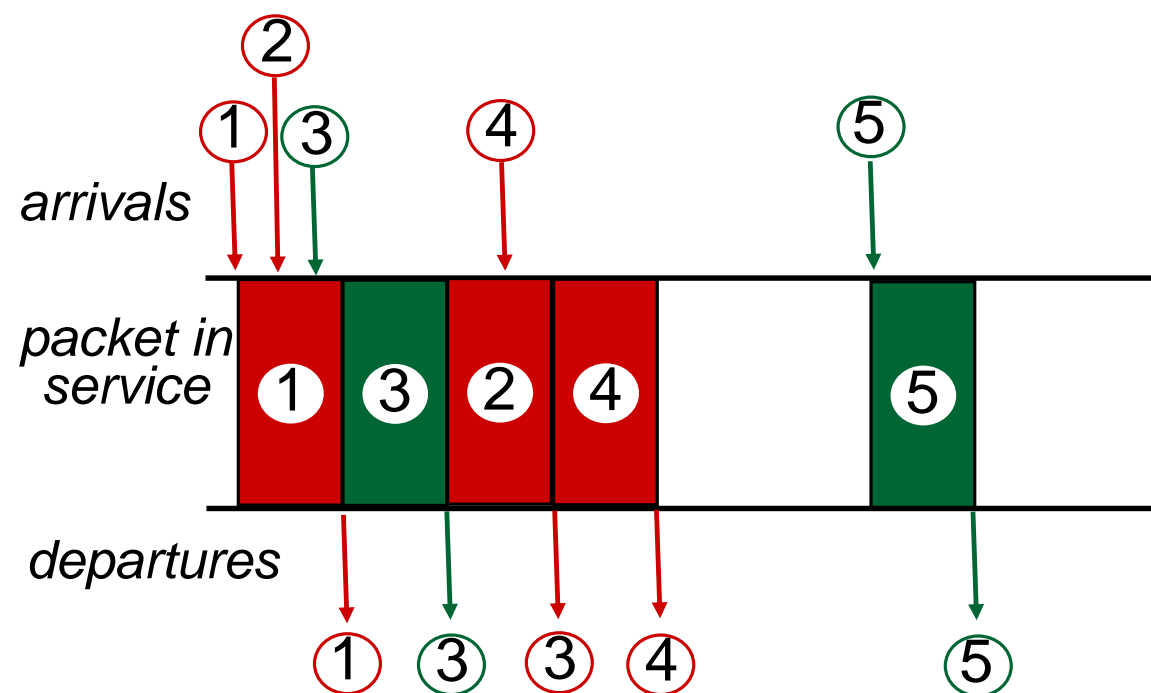
- multiple *classes*, with different priorities
 - class may depend on marking or other header info, e.g. IP source/dest, port numbers, etc.
 - real world example?



Scheduling policies: still more

Round Robin (RR) scheduling:

- multiple classes
- cyclically scan class queues, sending one complete packet from each class (if available)
- real world example?



Scheduling policies: still more

Weighted Fair Queuing (WFQ):

- generalized Round Robin
- each class gets weighted amount of service in each cycle
- real-world example?

