

CS 330: Network Applications & Protocols

Network Layer

Galin Zhelezov
Department of Physical Sciences
York College of Pennsylvania



Overview of Network Layer

- **Virtual Circuit and Datagram Networks**
- **Router Architectures**
- **IP: Internet Protocol**
- **Routing algorithms**
 - Overview
 - Link state
 - Distance vector
- **Routing in the Internet**
- **Broadcast and multicast routing**

Network-layer functions

Recall: two network-layer functions:

- **forwarding:** move packets from router's input to appropriate router output
- **routing:** determine route taken by packets from source to destination

data plane

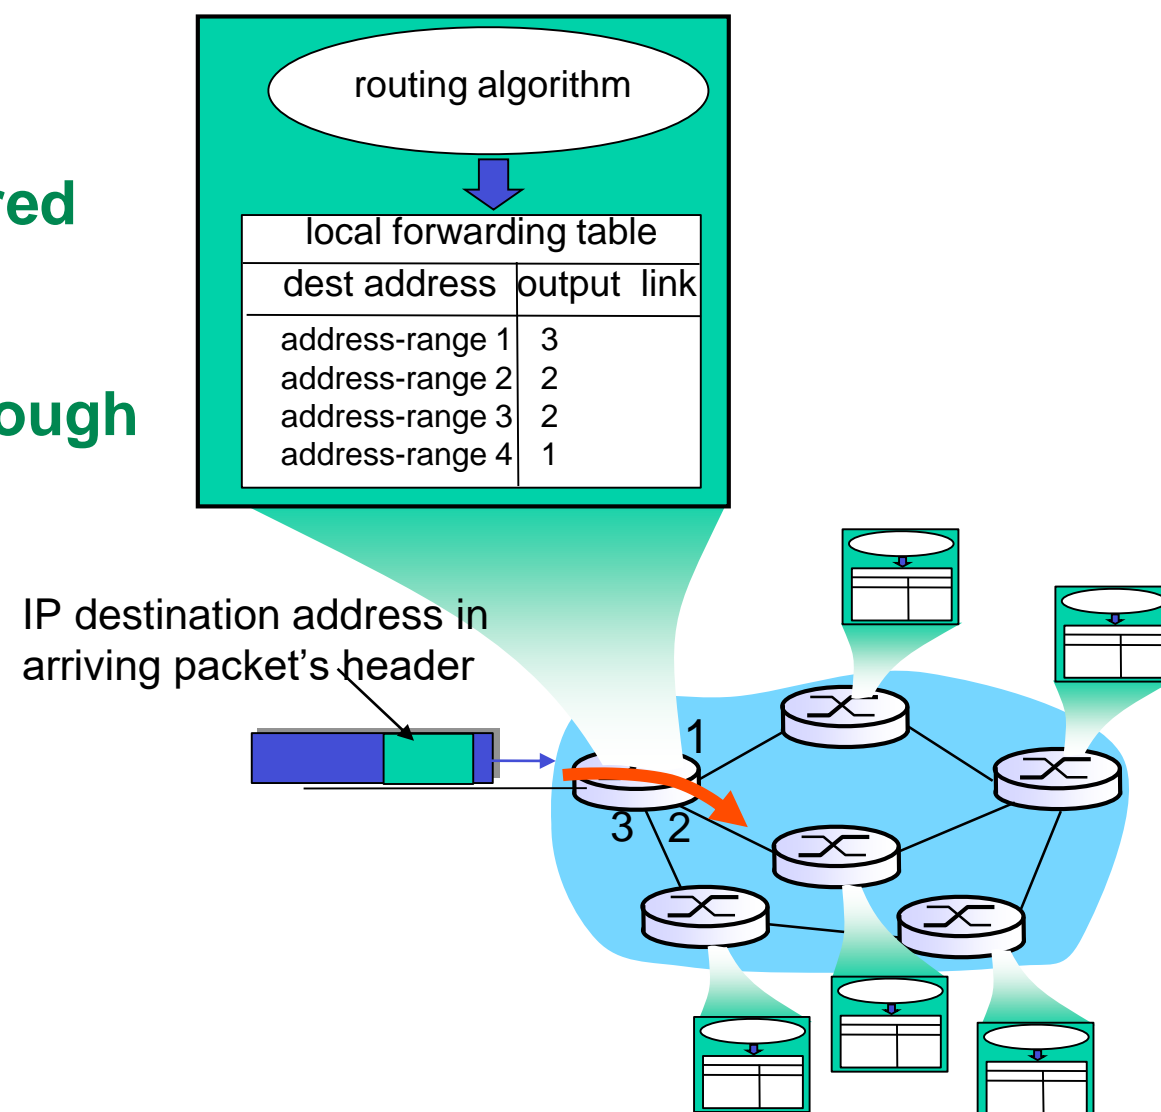
control plane

Two approaches to structuring network control plane:


- per-router control (traditional)
- logically centralized control (software defined networking)

Overview

- Routers use **routing algorithms** to maintain **forwarding tables**
- Routers exchange information with other routers to compute information that is entered into forwarding table
- Routing algorithms determine **best path** through a network from some source to some destination
 - **Least-cost path** is best path
 - Cost of path can depend on: distance, congestion, \$\$ cost, or other factors
 - Networks are represented as graphs
 - Routers are nodes in the graph
 - Links between routers are edges in graph



Graph abstraction of the network



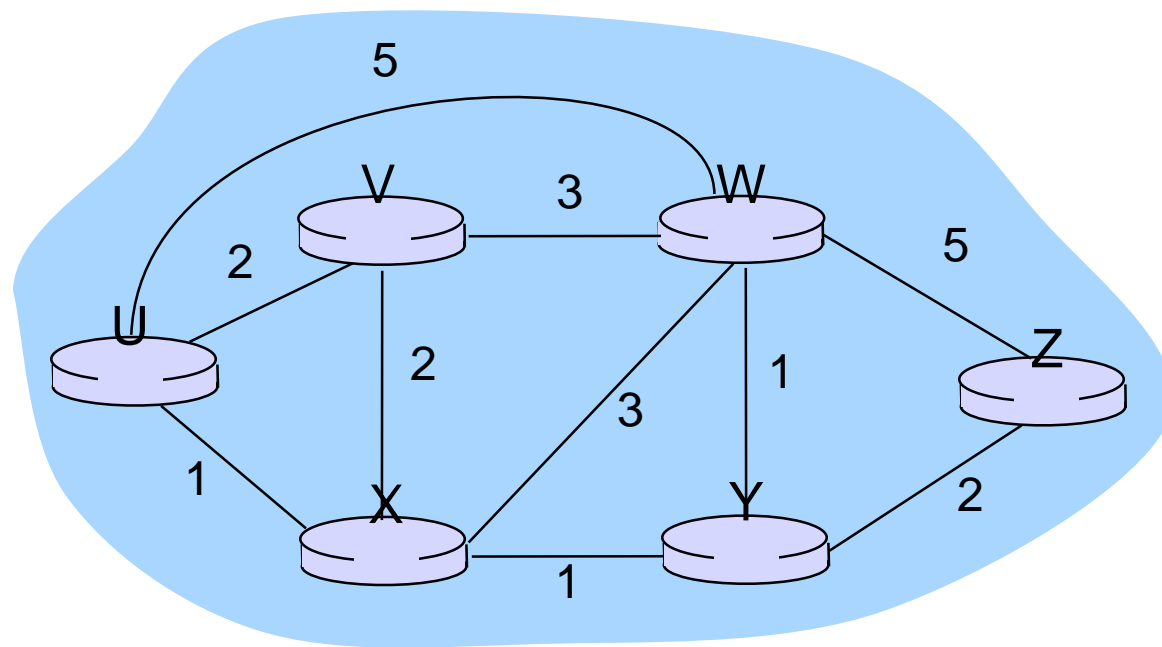
graph: $G = (N, E)$

N = set of routers = $\{ u, v, w, x, y, z \}$

E = set of links = $\{ (u,v), (u,x), (v,x), (v,w), (x,w), (x,y), (w,y), (w,z), (y,z) \}$

aside: graph abstraction is useful in other network contexts, e.g., P2P, where N is set of peers and E is set of TCP connections

Graph abstraction: costs



$c(x, x') = \text{cost of link } (x, x')$
e.g., $c(w, z) = 5$

cost could always be 1, or
inversely related to bandwidth,
or inversely related to
congestion

cost of path $(x_1, x_2, x_3, \dots, x_p) = c(x_1, x_2) + c(x_2, x_3) + \dots + c(x_{p-1}, x_p)$

key question: what is the least-cost path between u and z ?
routing algorithm: algorithm that finds that least cost path

Types of Routing Algorithms

- **Global Routing Algorithm**

- All routers know complete network topology, and link cost info
- Often called **Link-state algorithms** since the algorithm must know the state (i.e. cost) of all links in the network

- **Decentralized Routing Algorithm**

- No router has complete information about the network or links costs
- Routers know information about physically-connected neighbors, such as link costs to neighbors
 - **Neighbors also exchange information about each others' neighbors**
- Calculation of least-cost path is determined iteratively in a distributed fashion
- An example is the **Distance-vector algorithm**

Types of Routing Algorithms

- **Static Routing Algorithm**

- Routes change slowly over time
- Static routes are often entered by a human into a forwarding table

- **Dynamic Routing Algorithm**

- Routes are not dependent on human input
- Routes are periodically recomputed based on changes in the network
 - A new router is added
 - A router crashes
 - Too much congestion in part of the network
 - etc.
- Must be intelligent to avoid creating routing loops in network

Overview of Network Layer

- **Virtual Circuit and Datagram Networks**
- **Router Architectures**
- **IP: Internet Protocol**
- **Routing algorithms**
 - Overview
 - **Link state**
 - Distance vector
- **Routing in the Internet**
- **Broadcast and multicast routing**

Link-State Routing Algorithm

- **All routers know complete network topology, and link cost info**
 - Periodically, **routers broadcast link-state information** about each of their links to all other routers on the network
 - Network topology and link costs known to all routers
 - **All nodes have same information about the network**
 - Each router computes least-cost path from itself (the source) to *ALL* other routers on the network (the destinations)
 - **Typically done using Dijkstra's algorithm**
 - **Result provides forwarding table for that node**

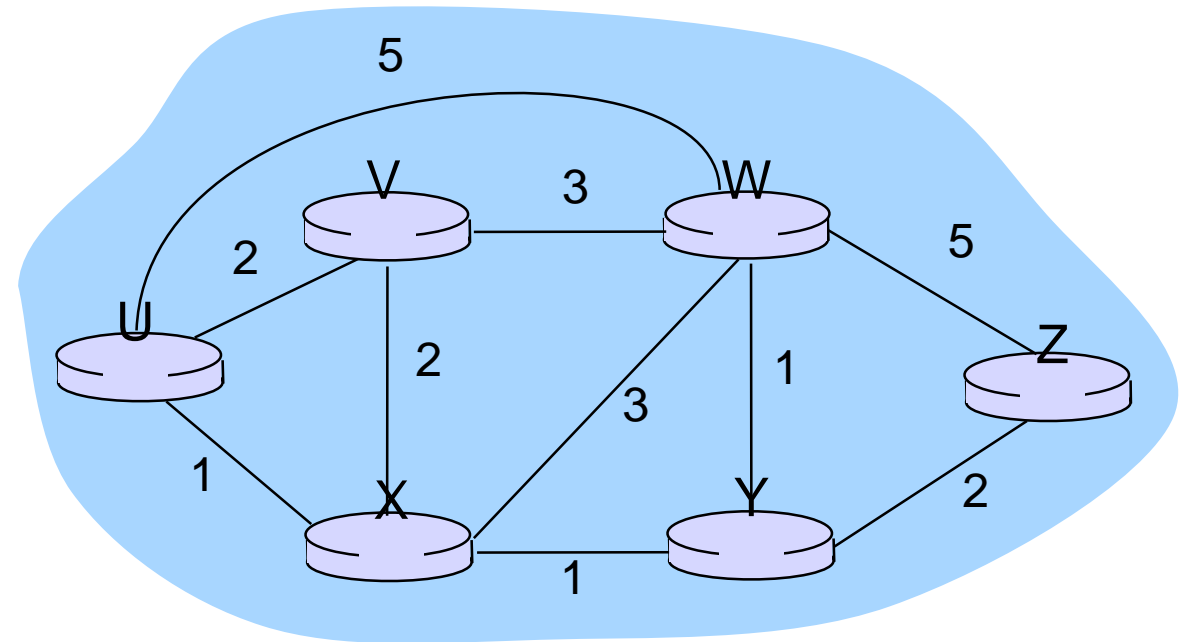
Dijkstra's Algorithm

- **Algorithm works as follows:**

- Starts by assigning some initial distance value for each node in the graph
 - Distance from node s to itself is 0
 - Distances from node s to all other nodes in graph are initialized to INFINITY
- Operates in steps, where at each step the algorithm improves the distance values for nodes in the graph
- At each step the shortest distance from node s to another node in the graph is determined

Dijkstra's Algorithm Example

- **Network is represented as a graph**
 - Routers are nodes in graph
 - Links are edges in graph
 - Cost of edge is labeled
- **Each router computes distance to all other routers in network**
 - Example shows computation done by router U



A link-state routing algorithm

Dijkstra's algorithm

- **net topology, link costs known to all nodes**
 - accomplished via “link state broadcast”
 - all nodes have same info
- **computes least cost paths from one node (‘source’) to all other nodes**
 - gives *forwarding table* for that node
- **iterative: after k iterations, know least cost path to k dest.’s**

notation:

- **$c(x,y)$: link cost from node x to y ; $= \infty$ if not direct neighbors**
- **$D(v)$: current value of cost of path from source to dest. v**
- **$p(v)$: predecessor node along path from source to v**
- **N' : set of nodes whose least cost path definitively known**

Dijkstra's algorithm

1 **Initialization:**

2 $N' = \{u\}$

3 for all nodes v

4 if v adjacent to u

5 then $D(v) = c(u, v)$

6 else $D(v) = \infty$

7

8 **Loop**

9 find w not in N' such that $D(w)$ is a minimum

10 add w to N'

11 update $D(v)$ for all v adjacent to w and not in N' :

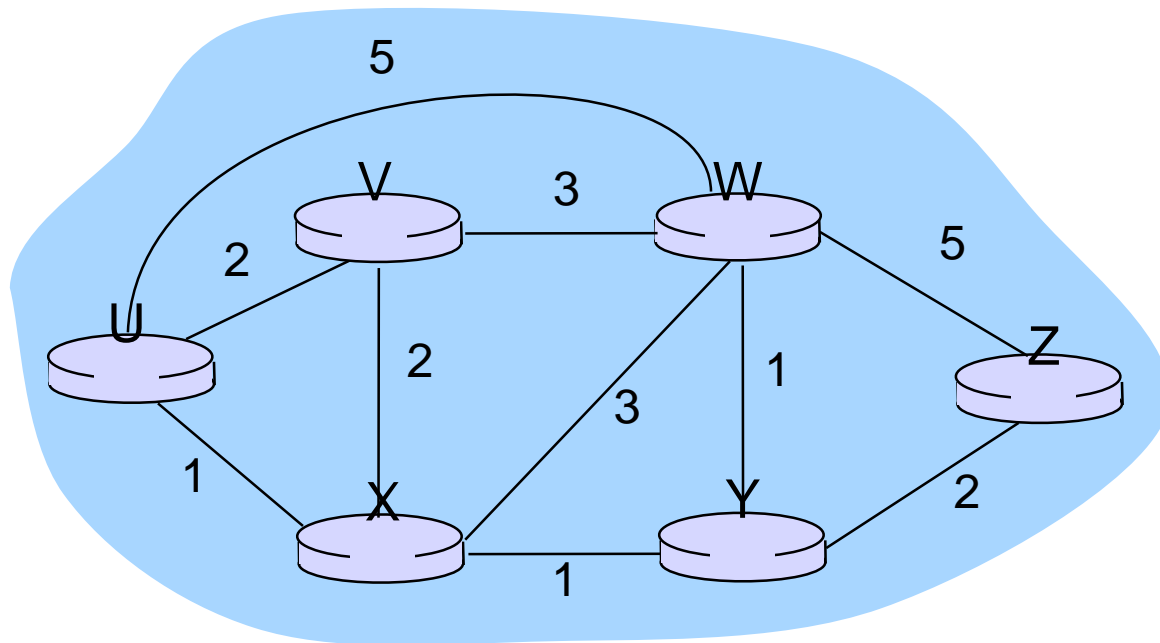
12 **$D(v) = \min(D(v), D(w) + c(w, v))$**

13 /* new cost to v is either old cost to v or known

14 shortest path cost to w plus cost from w to v */

15 **until all nodes in N'**

Dijkstra's Algorithm Example (Node U)




Initialize distances to U:

- Distance to itself is 0
- Distance to all other nodes is ∞

Node	Distance Computation from U						Path
U	0						
V	∞						
W	∞						
X	∞						
Y	∞						
Z	∞						

Dijkstra's Algorithm Example (Node U)




Select node with shortest distance to U (currently U) and determine shortest distance of its neighbors from U

- If node is unreachable it is still ∞
- Record the path

Node	Distance Computation from U						Path
U	0	-					-
V	∞	2					U → V
W	∞	5					U → W
X	∞	1					U → X
Y	∞	∞					-
Z	∞	∞					-

Dijkstra's Algorithm Example (Node U)




Select node with shortest distance to U (currently X) and determine shortest distance of its neighbors from U

- Node selected is min distance

Node	Distance Computation from U						Path
U	0	-	-				-
V	∞	2	2				U \rightarrow V
W	∞	5	4				U \rightarrow X \rightarrow W
X	∞	1	-				U \rightarrow X
Y	∞	∞	2				U \rightarrow X \rightarrow Y
Z	∞	∞	∞				-

Dijkstra's Algorithm Example (Node U)

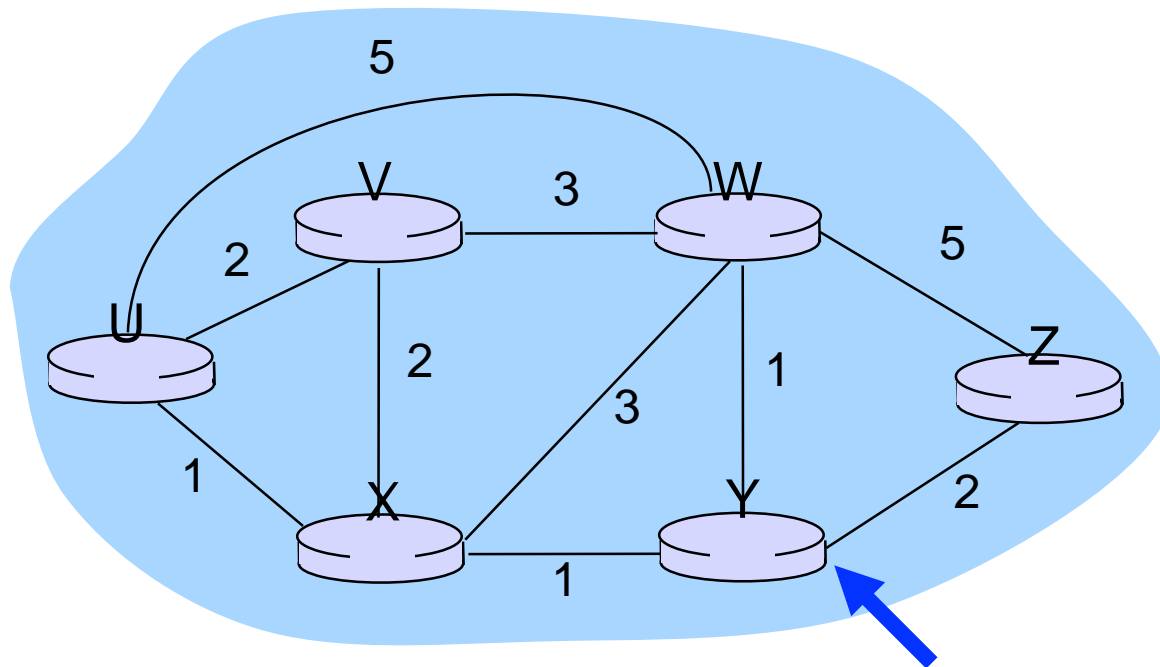


Select node with shortest distance to U (currently X) and determine shortest distance of its neighbors from U

- Node selected is min distance

Node	Distance Computation from U						Path
U	0	-	-	-	-	-	-
V	∞	2	2	-	-	-	U → V
W	∞	5	4	4	-	-	U → X → W
X	∞	1	-	-	-	-	U → X
Y	∞	∞	2	2	-	-	U → X → Y
Z	∞	∞	∞	∞	-	-	-

Dijkstra's Algorithm Example (Node U)




Select node with shortest distance to U (currently Y) and determine shortest distance of its neighbors from U

- Node selected is min distance

Node	Distance Computation from U						Path
U	0	-	-	-	-	-	-
V	∞	2	2	-	-	-	U → V
W	∞	5	4	4	3	-	U → X → Y → W
X	∞	1	-	-	-	-	U → X
Y	∞	∞	2	2	-	-	U → X → Y
Z	∞	∞	∞	∞	4	-	U → X → Y → Z

Dijkstra's Algorithm Example (Node U)

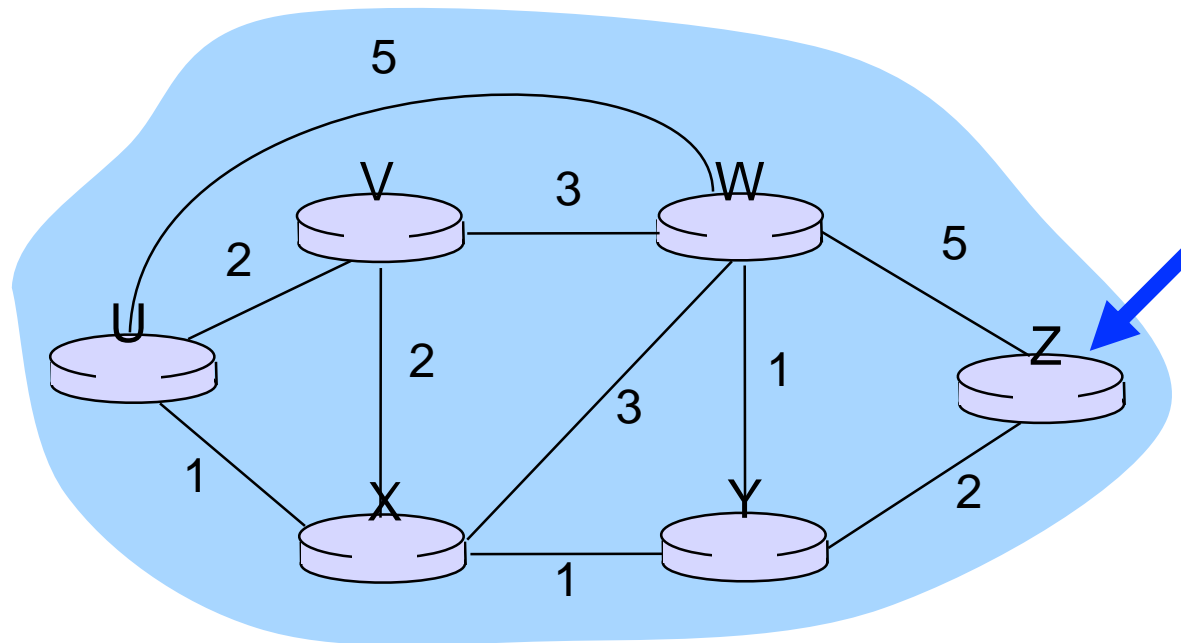


Select node with shortest distance to U (currently W) and determine shortest distance of its neighbors from U

- Node selected is min distance

Node	Distance Computation from U						Path
U	0	-	-	-	-	-	-
V	∞	2	2	-	-	-	U → V
W	∞	5	4	4	3	-	U → X → Y → W
X	∞	1	-	-	-	-	U → X
Y	∞	∞	2	2	-	-	U → X → Y
Z	∞	∞	∞	∞	4	4	U → X → Y → Z

Dijkstra's Algorithm Example (Node U)



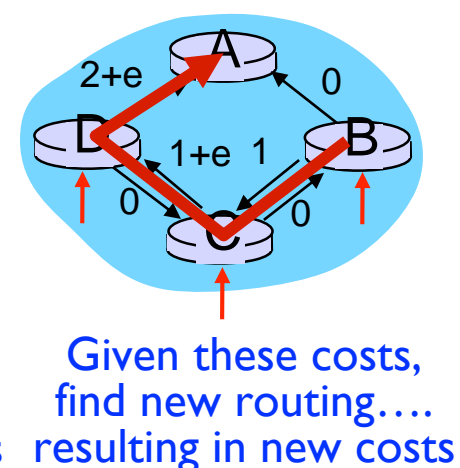
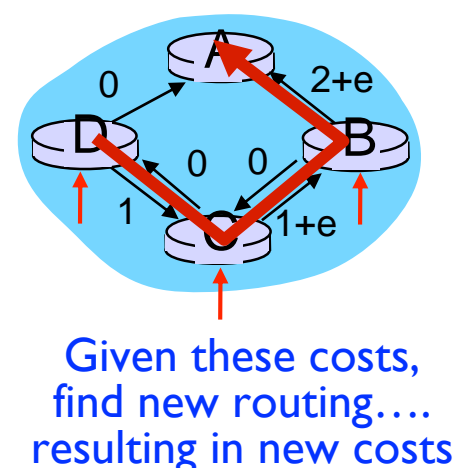
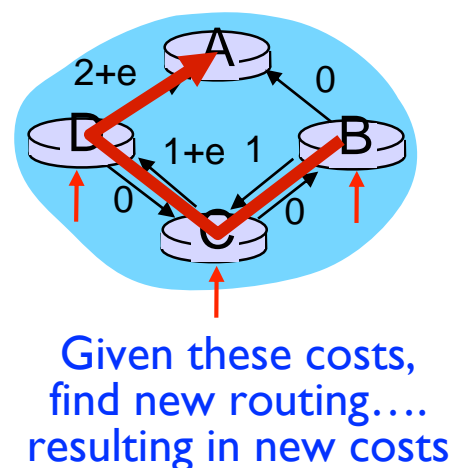
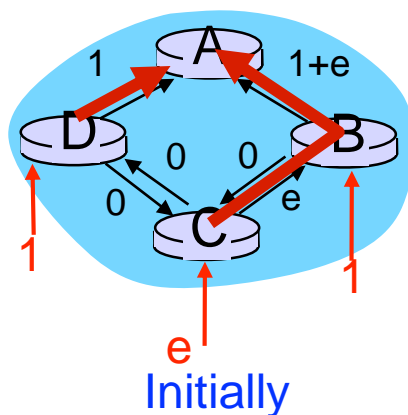
Select node with shortest distance to U (currently Z) and determine shortest distance of its neighbors from U

- All nodes have been accounted for, so terminate

Node	Distance Computation from U						Path
U	0	-	-	-	-	-	-
V	∞	2	2	-	-	-	U → V
W	∞	5	4	4	3	-	U → X → Y → W
X	∞	1	-	-	-	-	U → X
Y	∞	∞	2	2	-	-	U → X → Y
Z	∞	∞	∞	∞	4	4	U → X → Y → Z

Dijkstra's Algorithm Considerations

- When using Dijkstra's algorithm dynamically, **oscillations are possible**
 - For example, if link cost equals amount of carried traffic
 - Links with no traffic have low cost ..
 - So all traffic is rerouted down those links
 - Then a different set of links will have lower cost



Overview of Network Layer

- **Virtual Circuit and Datagram Networks**
- **Router Architectures**
- **IP: Internet Protocol**
- **Routing algorithms**
 - Overview
 - Link state
 - Distance vector
- **Routing in the Internet**
- **Broadcast and multicast routing**

Distance Vector Algorithm

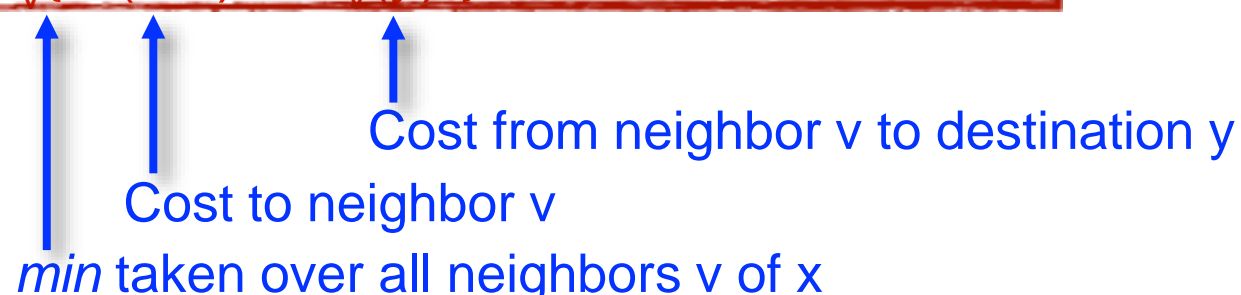
- No router has complete information about the network or links costs
- Routers know information about physically-connected neighbors, such as link costs to neighbors
- Key idea
 - From time-to-time, each node sends its own distance vector estimate to its neighbors
 - When x receives new DV estimate from neighbor, it updates its own DV using the **Bellman-Ford algorithm**

Let

$d_x(y) :=$ cost of least-cost path from x to y

then

$$d_x(y) = \min_v \{ c(x,v) + d_v(y) \}$$


 \min taken over all neighbors v of x
Cost to neighbor v
Cost from neighbor v to destination y

Distance vector algorithm

key idea:

- from time-to-time, each node sends its own distance vector estimate to neighbors
- when x receives new DV estimate from neighbor, it updates its own DV using B-F equation:

$$D_x(y) \leftarrow \min_v \{c(x, v) + D_v(y)\} \text{ for each node } y \in N$$

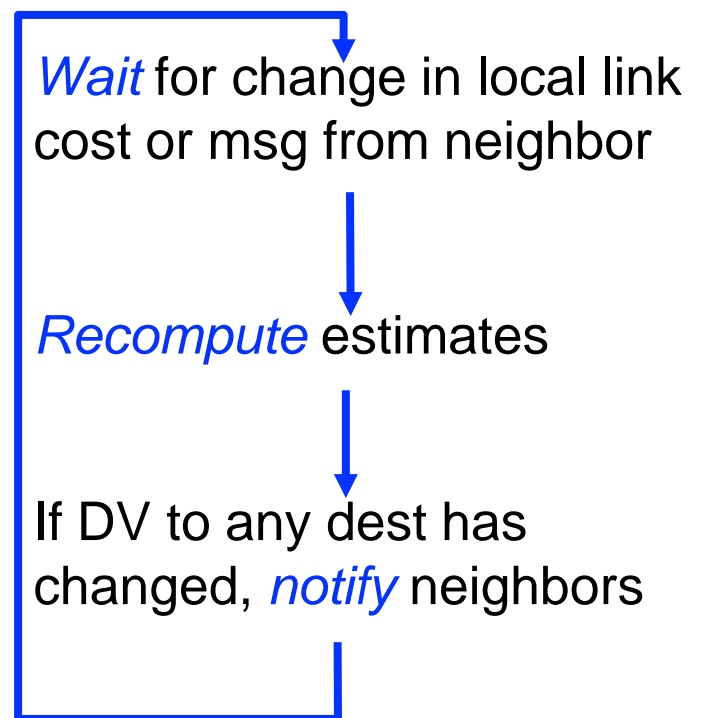
- ❖ under minor, natural conditions, the estimate $D_x(y)$ converge to the actual least cost $d_x(y)$

Distance Vector Algorithm

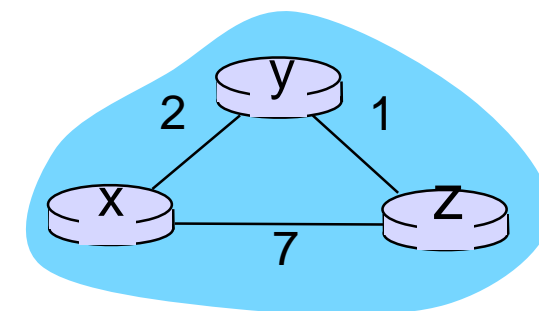
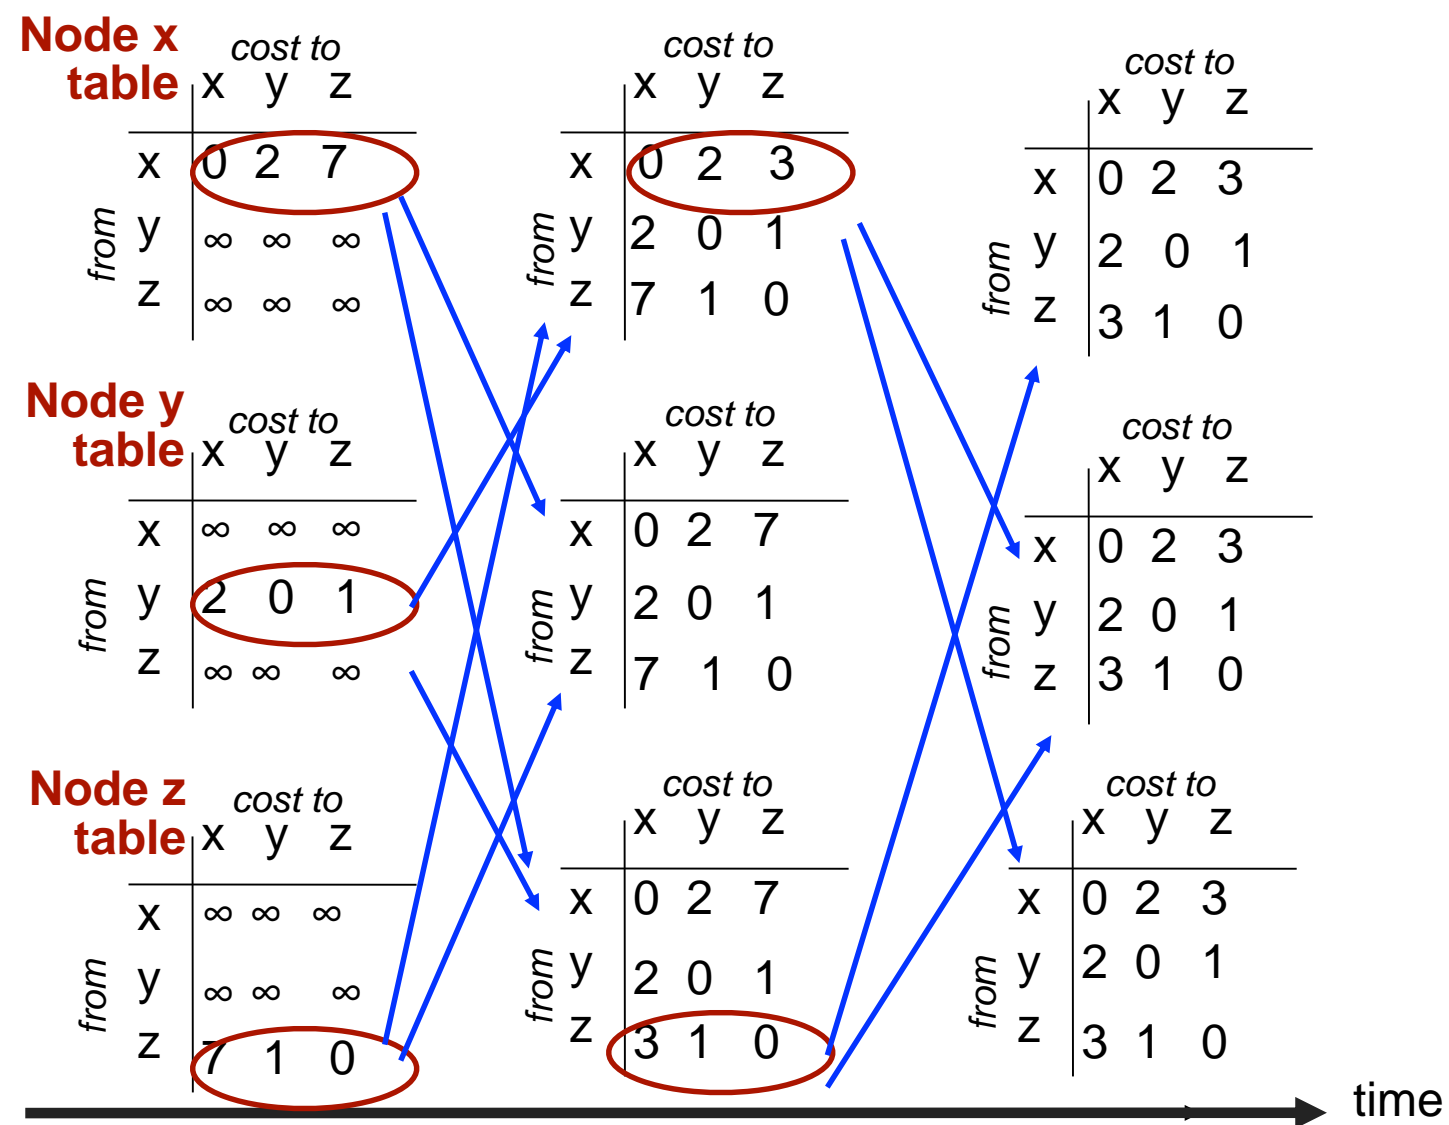
- **Algorithm is:**

- Iterative and asynchronous
 - Each local iteration caused by:
 - Local link cost change
 - DV update message from neighbor
- Distributed
 - Each node notifies neighbors only when its DV changes
 - Neighbors then notify their neighbors if necessary

At Each Node:



Distance Vector Example



Distance Vector Considerations

- **When link costs change**

- Good news travels quickly

- When link cost decreases, notifying neighbors is quick

- Bad news travels slowly

- When link cost increases, can cause **routing loops** that last for long periods of time (see text section 5.5.2)

Comparison of LS and DV algorithms

message complexity

- **LS:** with n nodes, E links, $O(nE)$ msgs sent
- **DV:** exchange between neighbors only
 - convergence time varies

speed of convergence

- **LS:** $O(n^2)$ algorithm requires $O(nE)$ msgs
 - may have oscillations
- **DV:** convergence time varies
 - may be routing loops
 - count-to-infinity problem

robustness: what happens if router malfunctions?

LS:

- node can advertise incorrect *link* cost
- each node computes only its own table

DV:

- DV node can advertise incorrect *path* cost
- each node's table used by others
 - error propagate thru network

Hierarchical Routing

- **Routing protocols within a single network (an autonomous system) are determined by some network administrator**
 - Provide routing information for hosts *within the network* (i.e. Intra-AS routing)
 - For example, a company or ISP would use some routing protocol for their network
- **Routers within an AS also need a way to route packets to hosts on other networks**
 - Route packets to **gateway router** that will direct packets to the next network (AS)
 - All packets destined for another network are routed to the gateway router
 - What if there are multiple gateways?
To which gateway should packets get routed?
 - Handled by **inter-AS routing protocols (BGP)**
 - Routers get information from both intra and inter-AS routing protocols and maintain info in forwarding table

